

Anders Bjorholm Dahl

Vedrana Andersen Dahl

Advanced Image Analysis

SELECTED TOPICS

DTU Compute, April 22, 2020

Contents

Preface 5

1 *Introduction* 7

PART I FEATURE-BASED IMAGE ANALYSIS

2 *Scale-space* 23

3 *Feature-based segmentation* 29

4 *Feature-based registration* 35

PART II IMAGE ANALYSIS WITH GEOMETRIC PRIORS

5 *Markov random fields* 43

6 *Deformable models* 53

7 *Using geometric priors for volumetric segmentation* 59

PART III IMAGE ANALYSIS WITH NEURAL NETWORKS

8 *Feed forward neural network* 67

9 *MNIST classification* 75

10 *Convolutional neural networks* 79

PART IV MINI PROJECTS

11 *Texture analysis* 87

12 *Optical flow* 91

13 *Layered surfaces* 95

14 *Spectral segmentation and normalized cuts* 101

15 *Probabilistic Chan-Vese* 107

16 *Orientation analysis* 109

- 17 *CNN for segmentation* 115
18 *Superresolution from line scans* 117

Preface

THIS LECTURE NOTE is a collection of topics for the students taking the course 02506 Advanced Image Analysis at Technical University of Denmark. The note provides background material for the course together with practical guidelines and advice for carrying out tasks in image analysis. The topics are selected to represent problems that you typically meet as an engineer that specialize in image analysis.

Image analysis is a rapidly growing field of research with a wealth of methods constantly being developed and published. This note is not intended to give a complete overview of the field. Instead, we focus on general principles for image analysis with the aim of giving students the skill-set needed for exploring new methods. General principles relate to identifying relevant image analysis problems, finding suitable methods for quantification, implementing image analysis algorithms and verifying their performance. This requires programming skills and the ability to translate a mathematical description into an efficient functioning program.

Image analysis methods published in scientific articles can be challenging to implement as a computer program, since they are described using mathematical notation, which may vary between articles. In some cases the notation can be very different from the code you need to write. One aim with this note is to guide the implementation of image analysis algorithms from descriptions in articles to the functioning programs. This is done through examples, practical tips, and advice on designing useful test to ensure that the obtained implementation gives the expected output.

There are several papers and book chapters that describe the methods to be implemented during the course. These are integral parts of the course curriculum, and should be read when working with the examples in this note.

The structure of the note is as follows. First comes a general introduction to a few central aspects relevant for image analysis along with the first introductory exercise, which has the purpose of refreshing basic image analysis concepts. The main text includes three compul-

sory exercises. Towards the end of the note we provide a number of examples for the final exercise in form of a mini-project.

During the course you may be implementing methods and algorithms that are already integrated in existing software libraries. These commercial or public implementations might be better than what you can achieve given the time available for the exercises. The reason to redo what other people have already done is to gain insight and understanding of how image analysis methods work, and to give you the skill-set needed for implementing or modifying advanced methods where there might not be an available implementation. It can however be a good idea to use existing implementations for evaluating your implementation.

1 Introduction

WE REFER TO IMAGES as regularly sampled signals in 2D or 3D space that typically represent a measurement, e.g. a measure of light intensity. In image analysis, we use the image to obtain some information about the signal we have measured. There are aspects to consider when working with images regarding what they represent and how they were created, that we will discuss here.

We start with the mathematical notation of images. One way of representing an image using mathematical notation is as a function $I(x, y)$ with $I : \Omega \subset \mathbb{R}^2 \rightarrow \mathbb{R}$, which means that each coordinate (x, y) in the image domain Ω a scalar value $I(x, y)$ is assigned. Typically the image is sampled at integer values, e.g. running from 1, such that $x \in \{1, \dots, X\}$ and $y \in \{1, \dots, Y\}$. There can be some variation between scientific articles on the notation of an image, e.g. that it is implicitly assumed that the image lives in 2D space, and the notation would simply be a symbol like I .

A 3D image is typically termed a volume, but again we can model it as a function $I : \Omega \subset \mathbb{R}^3 \rightarrow \mathbb{R}$. Volumetric images are often reconstructed from projection data obtained using a scanner, e.g. a CT or an MRI scanner. Again the image is represented as a function that maps to a scalar value, but here from a 3D coordinate (x, y, z) . In a volumetric image, the three spatial dimensions encode intensity information similar to a 2D image. This means that if we apply operators on a volumetric image, we would use a 3D operator, e.g. an averaging filter in three dimensions. In some cases the sampling is anisotropic, which is typically seen in e.g. medical CT images, and this can influence the applied analysis methods.

Spectral images have multiple measures in each pixel (sometimes represented as a vector) that encode the recorded spectral bands. A common example are the RGB images where $I : \Omega \subset \mathbb{R}^2 \rightarrow \mathbb{R}^3$ encodes the red, green, and blue bands. If more spectral bands are recorded, we are typically talking about multispectral or hyper-spectral images where $I : \Omega \subset \mathbb{R}^2 \rightarrow \mathbb{R}^n$ normally with $n > 3$. For 2D spectral images we would often apply operators in the spectral bands independently. Using

the smoothing example from before, but now in a RGB image, would be a 2D smoothing in the R-band, G-band, and B-band respectively.

Another common image-related representation is a movie. A movie is a set of consecutive images also called frames sampled over time, which we can model as $I(x, y, t)$ where $I : \Omega \subset \mathbb{R}^3 \rightarrow \mathbb{R}$ for a gray scale movie or $I : \Omega \subset \mathbb{R}^3 \rightarrow \mathbb{R}^3$ for an RGB movie. You can also have a multispectral movie ($I : \Omega \subset \mathbb{R}^3 \rightarrow \mathbb{R}^n$ with $I(x, y, t)$) or a volumetric movie ($I : \Omega \subset \mathbb{R}^4 \rightarrow \mathbb{R}$ with $I(x, y, z, t)$). For movies we would typically expect small changes between frames, and this can be utilized in the analysis.

1.1 Introductory exercise

This exercise is aimed at refreshing or introducing concepts from basic image analysis curriculum and other related subjects. It contains some topics that will be useful at a later stage in the course. You are expected to carry out the first exercises, whereas the last exercises are not mandatory, but you are welcome to read or solve those exercises also.

1.1.1 Image convolution

Image convolution is a central tool in image analysis, and in this exercise you will investigate some properties of image convolution with a Gaussian kernel and its derivatives.

For two continuous functions convolution is defined as

$$(f * g)(x) = \int_{-\infty}^{\infty} f(x - \tau)g(\tau)d\tau. \quad (1.1)$$

Convolution is commutative, but we usually distinguish between the signal and the kernel, and we say that the signal f is convolved with the kernel g .

For a discrete sampled signal we get

$$(f * g)(x) = \sum_{i=-l}^l f(x - i)g(i). \quad (1.2)$$

A convolution with a square kernel in 2D is given by

$$(f * g)(x, y) = \sum_{i=-l}^l \sum_{j=-l}^l f(x - i, y - j)g(i, j). \quad (1.3)$$

In image analysis, Gaussian kernel is often used for image smoothing by filtering. The 1D Gaussian is defined by

$$g(x; t) = \frac{1}{\sqrt{2t\pi}} e^{-x^2/(2t)}, \quad (1.4)$$

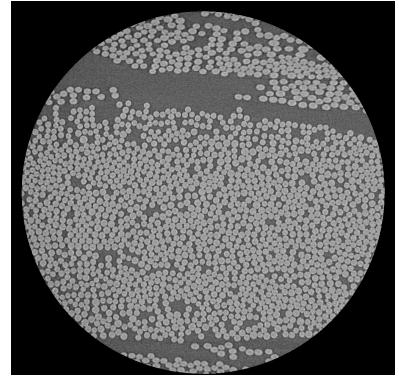


Figure 1.1: Slice of a CT image of glass fibers viewed orthogonal to the fiber direction.

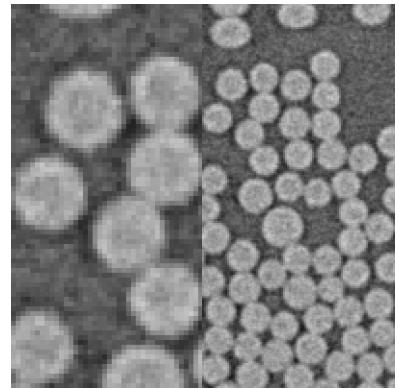


Figure 1.2: Two zoomed in images from image shown in Figure 1.1.

where $t = \sigma^2$ is the variance of the Gaussian normal distribution. The 2D isotropic Gaussian is given by

$$g(x, y; t) = \frac{1}{2t\pi} e^{-(x^2+y^2)/(2t)} . \quad (1.5)$$

The Gaussian is separable, which means that we can convolve the image using two orthogonal 1D Gaussians, and obtain the same result as when convolving with 2D Gaussian of the same variance. This can speed up convolutions significantly, especially for large convolution kernels.

Another property of the Gaussian convolution is the so-called *semi-group structure*, which means that we get the same convolution using a single large Gaussian as we get using several small ones

$$g(x, y; t_1 + t_2) * I(x, y) = g(x, y; t_1) * g(x, y; t_2) * I(x, y) , \quad (1.6)$$

where I is an image. On the right part of equation, the order of convolution does not matter, as convolution is associative.

When computing various features for image $I(x, y)$, we often need to know a local change in intensity values for all positions (x, y) . This can be achieved by taking the spatial derivative of the image. Since the image is a discretely sampled signal, we can only compute an approximation of the derivative. For example, we can take the difference between neighboring pixels. On the other side, we often want to smooth the image in connection with taking the derivative.

However, it turns out that if we want to convolve the image with a Gaussian and then take the derivative, we can instead convolve the image with the derivative of the Gaussian

$$\frac{\partial}{\partial x} (I * g) = \frac{\partial I}{\partial x} * g = I * \frac{\partial g}{\partial x} . \quad (1.7)$$

Since we can compute the derivative of the Gaussian analytically, we get an efficient and elegant approach to computing a smoothed image derivative.

The analytic 1D Gaussian derivative is given by

$$\frac{d}{dx} g(x) = \frac{x}{\sigma^3 \sqrt{2\pi}} e^{-x^2/(2\sigma^2)} . \quad (1.8)$$

The semi-group structure also holds for image derivatives, such that we get

$$\frac{\partial}{\partial x} g(x, y; t_1 + t_2) * I(x, y) = \frac{\partial}{\partial x} g(x, y; t_1) * (g(x, y; t_2) * I(x, y)) , \quad (1.9)$$

which implies that we can convolve with a large Gaussian derivative kernel or we can convolve with a smaller Gaussian and a smaller Gaussian derivative and get the same result.

Data For this exercise you will use an X-ray CT image of fibres `fibres_xcth.png`, shown in Figure 1.1 and Figure 1.2.

Tasks

1. Experimentally verify the separability of the Gaussian convolution kernel. Do this by convolving a test image with 2D kernel, and convolving the same image with two orthogonal 1D kernels. Subtract the result and verify that the difference is very small.
2. Investigate the difference between the derivative of the image convolved by a Gaussian and the image convolved with the derivative of the Gaussian as described in Eq. 1.7. Note that you can compute the derivative of the image by convolving with the kernel $k = [0.5, 0, -0.5]$.
3. Test if a single large convolution with a Gaussian of $t = 20$ is equal to ten convolutions with a Gaussian of $t = 2$.
4. Test if convolution with a large Gaussian derivative

$$I * \frac{\partial g(x, y; 20)}{\partial x},$$

is equal to convolving with a Gaussian with $t = 10$ and a Gaussian derivative with $t = 10$

$$I * g(x, y; 10) * \frac{\partial g(x, y; 10)}{\partial x}.$$

1.1.2 Computing length of segmentation boundary

Segmentation is one of the basic image analysis tasks, and we will also cover a few segmentation methods in the course. For an outcome of a segmentation, as for example shown in Figure 1.3, it may be important to measure some quality of the result. One relevant measurement is a length of the segmentation boundary.

Assume that segmentation is represented by an image $S(x, y)$ which takes n discrete values, i.e. $S : \Omega \rightarrow \{1, 2, \dots, n\}$, where n is the number of segments. We define the length of the segmentation boundary as

$$L(S) = \sum_{(x,y) \sim (x',y')} d(S(x, y), S(x', y')) ,$$

where $(x, y) \sim (x', y')$ indicates two neighboring pixel locations, and d is a discrete metric

$$d(a, b) = \begin{cases} 0 & \text{if } a = b \\ 1 & \text{otherwise} \end{cases}$$

which in this case operates on pixel intensities. In other words, $L(S)$ counts all occurrences of two neighboring pixels having different labels.

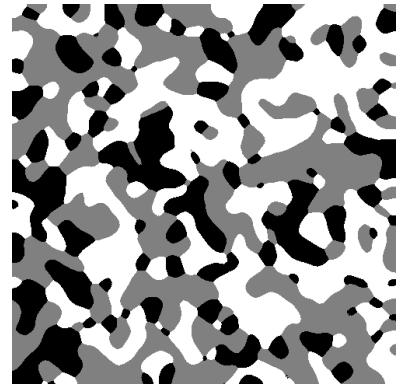


Figure 1.3: Image of a segmented fuel cell with three phases. Black represents air, grey is cathode, and white is anode.

Tasks

1. Compute the length of the segmentation boundary for provided segmentation images of a fuel cell, where one is shown in Figure 1.3. For efficient and compact implementation, avoid loops and use vectorization provided by Matlab or numpy.
2. Collect your code in a function which takes segmentation as an input, and returns the length of the segmentation boundary as an output. Your function will be useful when we will be working with Markov random fields later in the course.

Data In this exercise you should use the volume slice `fuel_cell_1.tif`, `fuel_cell_2.tif`, and `fuel_cell_3.tif` that you can find on Campusnet.

1.1.3 Curve smoothing

A segmentation boundary may be explicitly represented using a sequence of points connected by line segments, which typically delineates an object in the image. Assume that N -times-2 matrix \mathbf{X} contains x and y coordinates of N points which define a closed curve, a so-called snake¹.

To impose smoothness to this representation, we will need to smooth the snake. This can be achieved in a simple way by displacing every snake point towards the average of its two neighbors, possibly iteratively. Point displacement can be seen as a result of filtering the snake with a kernel $\lambda \begin{bmatrix} 1 & -2 & 1 \end{bmatrix}$, where λ is a parameter controlling the magnitude of the displacement. For efficiency, we want to implement the snake-smoothing step as

$$\mathbf{X}^{\text{new}} = (\mathbf{I} + \lambda \mathbf{L}) \mathbf{X} \quad (1.10)$$

where \mathbf{L} is a N -times- N matrix with elements 1, -2, and 1 in every row such that -2 is on the main diagonal, and 1 on its left and right (also circularly in the first and the last row), and zeros elsewhere.

Confirm that $\lambda = 0.5$ displaces every snake point exactly to the average of its neighbors. Try smoothing one of the provided contours, also shown in Figure 1.4. We have included both original and noisy curves.

Maybe you noticed two important limitations of our simple approach. First, for larger values of λ the curve will start oscillating, but using a small λ requires many iterations of the smoothing step for a noticeable result. Second, smoothing leads to the shrinkage of the curve.

Stability issues can be avoided by evaluating the displacement on the new snake \mathbf{X}^{new} . In other words, we can use an implicit (backwards

¹ Michael Kass, Andrew Witkin, and Demetri Terzopoulos. Snakes: Active contour models. *International Journal of Computer Vision*, 1(4):321–331, 1988

Euler) approach. Instead of Equation 1.10 where $\mathbf{X}^{\text{new}} = \mathbf{X} + \lambda \mathbf{L} \mathbf{X}$ we use $\mathbf{X}^{\text{new}} = \mathbf{X} + \lambda \mathbf{L} \mathbf{X}^{\text{new}}$ leading to

$$\mathbf{X}^{\text{new}} = (\mathbf{I} - \lambda \mathbf{L})^{-1} \mathbf{X}. \quad (1.11)$$

We can now choose an arbitrary large λ and obtain the desired smoothing in just one step. The price to pay is matrix inversion, but for many applications, this needs to be computed only once.

Shrinkage is caused by the kernel which minimizes curve length. Instead, we can use a kernel which minimizes the curvature, or even better, we can weight the elasticity (length minimizing) and rigidity (curvature minimizing) term. The kernel with the two contributions is

$$\alpha \begin{bmatrix} 0 & 1 & -2 & 1 & 0 \end{bmatrix} + \beta \begin{bmatrix} -1 & 4 & -6 & 4 & -1 \end{bmatrix}$$

as derived in ², section 3.2.4, with α and β weighting the two terms.

Tasks

1. Implement curve smoothing as in Equation 1.10 and test it for various values of λ . Try using smoothing iteratively to achieve a visible result for small λ .
2. Implement curve smoothing as in Equation 1.11 (implicit smoothing) and test it for various values of λ . Do you need an iterative approach of this smoothing?
3. Implement implicit curve smoothing but with the extended kernel. This means that your implementation instead of $\lambda \mathbf{L}$ uses a matrix that combines elasticity and rigidity, and has two parameters instead of λ . Note also that this matrix is a (sparse) circulant matrix. Test smoothing with various values of α and β . What do you achieve when choosing a large β and small α ?
4. Implement a function which returns a smoothing matrix needed for implicit smoothing with the extended kernel. You will be using this when working with deformable models later in the course.

Data In this exercise you should use the curves given as text files containing point coordinates `dino.txt`, `dino_noisy.txt`, `hand.txt`, and `hand_noisy.txt`.

1.1.4 Optional: Total variation

In many image analysis applications, such as image denoising and image segmentation, we are interested in producing a result which has

² Chenyang Xu, Dzung L Pham, and Jerry L Prince. Image segmentation using deformable models. *Handbook of medical imaging*, 2:129–174, 2000a

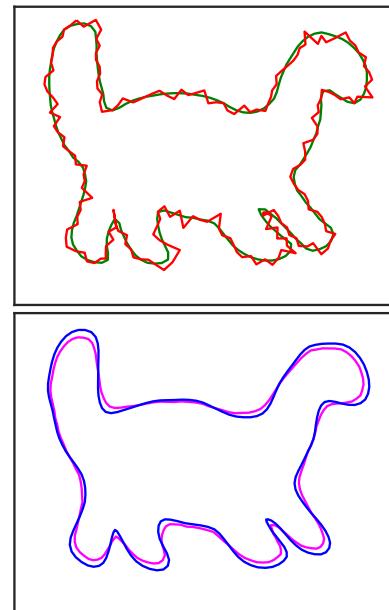


Figure 1.4: Top image shows the dinosaur curve in green, while red shows the curve with added noise. Bottom image shows two smoothing results with different α and β weights.

a quality that we loosely call *smoothness*. A common way of estimating smoothness is by considering a total variation defined for an image I as

$$V(I) = \sum_{x \sim x'} |I(x) - I(x')|,$$

where $x \sim x'$ indicates two neighboring pixel locations. We expect smooth images to have a low total variation.

Implement a function which computes the total variation of a 2D grayscale image and test it on the image shown in Figure 1.1 and Figure 1.2. Use Gaussian smoothing to remove some of the noise from the image, and confirm that the smoothed image has a smaller total variation.

Data In this exercise you should use the volume slice `fibres_xcth.png` that you can find on Campusnet.

1.1.5 Optional: Unwrapping image

A solution to image analysis problem may involve geometric transformations. When working with spherical or tubular objects, we sometimes want to represent an image in polar coordinate system. Implement a function which performs such *image unwrapping* using a desired angular and radial resolution. Use your function to unwrap one of the slices from the dental data set, an example is shown in Figure 1.5 and 1.6. Unwrapping will be useful when we will be working with deformable models later in the course.

Data In this exercise you should use one of the central slices from the dental folder that you can find on Campusnet.

1.1.6 Optional: Working with volumetric image

To give you a taste of working with 3D images, we have prepared a small data set containing slices from an X-ray CT scan. By convention in X-ray imaging, dense structures (having a high X-ray attenuation) are shown bright compared to less dense structures. Furthermore, the direction given by image slices is most often denoted z . The volume you are given contains a metal (very bright) object. Show orthogonal cross sections of the object, see Figure 1.7. Can you determine an optimal threshold for segmenting the object from the background?

Optionally, show a volumetric 3D rendering of the thresholded object using any available software. An example is shown in Figure 1.8. If you are using MATLAB, check a function `isosurface`.

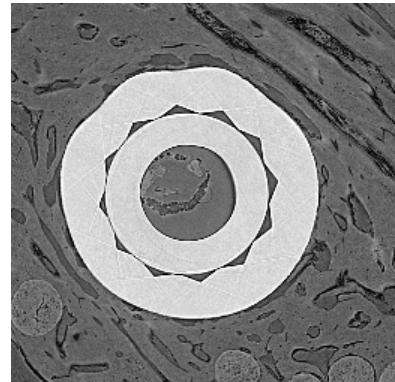


Figure 1.5: Image of a dental implant that should be unwrapped.

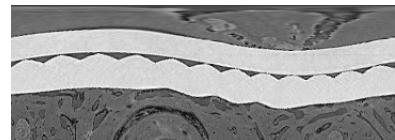


Figure 1.6: Unwrapped image of a dental implant.

Data In this exercise you should use the volumetric image stored as individual slices in the folder called dental that you can find on Campusnet.

1.1.7 Optional: PCA of multispectral image

Principal component analysis (PCA) is a linear transform of multivariate data that maps data points to an orthogonal basis according to maximum variance. A basic introduction in PCA is given in ³, and here we will apply it on a multispectral image.

We provided an image acquired with the VideometerLab, which is a multispectral imaging device, that uses coloured LED's to illuminate a material, in this case samples in a petri dish. This gives an 18 channel image $I : \Omega \subset \mathbb{R}^2 \rightarrow \mathbb{R}^{18}$ where each channel corresponds to a wavelength, and channels cover the range from 410 nm to 955 nm, i.e. the visible and near-infrared spectrum. The image depicts vegetables on a dish and is shown in false colours in Figure 1.9.

The aim of this exercise is to carry out PCA and visualise the principal components as images. PCA can be done by eigenvalue decomposition of a data covariance matrix. In our analysis we view each pixel as an observation, so we rearrange I into a N -by-18 data matrix \mathbf{X} . Each row of \mathbf{X} represents one pixel (observation), with the successive columns corresponding to wavelengths (variables).

Data covariance matrix \mathbf{C} is defined as

$$\mathbf{C}_{ij} = \frac{1}{N-1} \sum_{n=1}^N (\mathbf{x}_{ni} - \mu_i)(\mathbf{x}_{nj} - \mu_j), \quad (1.12)$$

where $i, j \in \{1, \dots, 18\}$, and μ is a 18 dimensional empirical mean vector computed for each variable.

Covariance matrix \mathbf{C} can be computed as a matrix product

$$\mathbf{C} = \frac{1}{N-1} \bar{\mathbf{X}}^T \bar{\mathbf{X}}, \quad (1.13)$$

where $\bar{\mathbf{X}} = \mathbf{X} - \mathbf{1}_{n \times 1} \mu^T$ is a zero-mean matrix obtained by independently centering each row of \mathbf{X} around its mean value. Convince yourself that is correct.

Principal components are given by the eigenvectors of \mathbf{C} , e.i. vectors such that $\mathbf{C}\mathbf{v}_i = \lambda_i \mathbf{v}_i$. Eigenvector corresponding to the largest eigenvalue gives the direction of the largest variance in data, the eigenvector corresponding to the second largest eigenvalue is the direction of the largest variance orthogonal to the first principal direction, etc. The projections of the data points onto principal directions $\hat{\mathbf{X}}\mathbf{v}_i$ can be rearranged back into image grid, and viewed as images.

If \mathbf{V} is a matrix containing eigenvectors in its columns, all principal

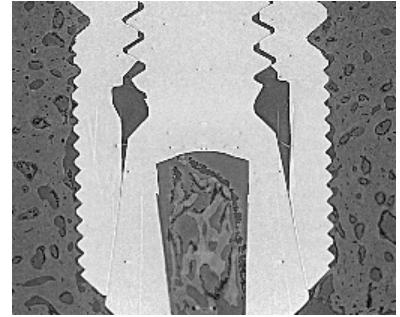


Figure 1.7: A longitudinal slice (an xz -plane) of the volumetric image of a dental implant.



Volume Viewer

Figure 1.8: 3D rendering of the thresholded volumetric image of a dental implant.



Figure 1.9: False colour image obtained from an 18 band VideometerLab image.

³ Lindsay I Smith. A tutorial on principal components analysis. Technical report, 2002

components can be computed as

$$\mathbf{Q} = \bar{\mathbf{X}}\mathbf{V}. \quad (1.14)$$

Data In this exercise you should use the images in the folder called `mixed_green` which contains png-images. You find the file on Campusnet.

Suggested approach The following steps takes you through computing the principal components.

1. Write a script to read in the images and display them. Convince yourself that there is a difference between the spectral bands. Make sure to change the data type to float or double.
2. Rearrange the image into a matrix \mathbf{X} as described above with one pixel in each row. Compute the column-wise mean μ and subtract this from \mathbf{X} to get the zero mean $\bar{\mathbf{X}}$.
3. Compute the covariance matrix \mathbf{C} .
4. Compute the eigenvectors \mathbf{V} and eigenvalues λ .
5. Compute the principal component loadings \mathbf{Q} .
6. Rearrange \mathbf{Q} into images and display the result.

You can compare your implementation to an already implemented PCA function in e.g. MATLAB or Python.

1.1.8 Optional: Bacterial growth from movie frames

Image data with a temporal component can be stored in the form of a movie. The purpose of this exercise is to read in image frames from a movie and analyze them. The movie contains microscopic images of listeria bacteria growing in a petri-dish acquired at equal time steps. An example frame is shown in Figure 1.10. Your task is to make a small program that visualize bacterial growth by counting the cells.

The image quality is however not very good due to the low resolution and compression artifacts, making it difficult to separate the individual bacteria. So, we make a rough assumption that the number of pixels covered by bacteria is proportional to the number of bacteria. The task is therefore to make a plot of the number of pixels covered by bacteria as a function of time.

Data In this exercise you should use the movie `listeria_movie.mp4` that you can find on Campusnet.



Figure 1.10: Example of microscopic image of listeria bacteria in a petri dish.

Suggested approach You can for example first read in one representative frame from the movie and build an cell segmentation method. A simple threshold is not sufficient, but with a few processing steps the cells become distinguishable from the background. You can try the following steps:

1. Convert the image I to a grey scale image G .
2. Compute the gradient magnitude $M = \sqrt{(\partial G / \partial x)^2 + (\partial G / \partial y)^2}$ using an appropriate filter.
3. Smooth M using a Gaussian filter.
4. If the parameters have been chosen appropriately, the pixels covering bacteria can now be segmented by thresholding.

When you have made a functioning segmentation model, you can apply this to all images in the movie using the following steps for each image in the movie:

1. Apply the segmentation and sum the bacteria pixels.
2. Store this number in an array.
3. Plot the number of pixels as a function of time.

The obtained curve has a characteristic shape. Can you recognize the function that could describe this shape?

1.1 Information for 02506

The purpose of this exercise is to refresh a few concepts from image analysis.

Requirement to complete the exercise We will be giving a feedback on this exercises during the next exercise session. To make this efficient, you should prepare, such that you can quickly show us your results. This should take less than five minutes. You should give an overview of what you accomplished, and focus on elements that you would like to get feedback on.

We suggest that you prepare a script (e.g. a Jupyter Notebook or MATLAB script) which covers the exercise. Make sure that it runs fast and clearly shows your results.

We will use the same feedback format for other exercises in the course. If there are parts of your code that take long time to run, you should save the results in advance, and load the results for the feedback session.

Exam questions related to this exercise Later in the course, the questions for the oral exam will be given at the end of each exercises. But for this exercise, we have no specific questions. However, the tools used here are either expected to be known before the course or are the concepts used in other exercises. The introductory exercise may therefore be discussed during the examination.

Part I

Feature-based image analysis

ANALYZING IMAGES using feature-based representations is central to many applications. We have chosen three topics that we will cover. First we will work with scale-space for detecting image features independently of scale. Specifically, we will focus on scale-space blob detection. Then we will investigate features as a basis for pixel classification as a method for segmentation. Finally, we will work with feature-based image registration.

Information for 02506

This part covers week 2 to 4 in the course, with a chapter for each week. For the first week you should read:

- Week 2: Lindeberg (1996): Scale-space: A framework for handling image structures at multiple scales

The reading material is uploaded to DTU Inside (Campusnet).

2 Scale-space

METHODS FROM SCALE-SPACE allow scale invariant detection of image structures. This means that we can find features like blobs (binary large objects), corners, ridges, edges, and other structures at different scale. When we talk about image features like corners and edges, it is not corners or edges of the physical depicted objects, but corners and edges in the image intensities. To visualize this, you can think of a 2D image as a landscape, with pixel intensities corresponding to height measurements at regularly placed positions. In this landscape, an edge is a line with high abruptly changes. A corner will be a height-change point where two (more or less) orthogonal edges meet, and other types of features can be described in the same way.

Using a feature-based image representation is convenient, because we break up the image into manageable parts that are more descriptive than the individual pixels. Scale invariance, which means that we characterize (make a mathematical description) the same feature shown at different scale in two images, is also very convenient. In e.g. computer vision where images of the same object are often captured from difference distance, it is typically a desired property to be able to measure the features independent of its scale. But it also allows us to measure image structures that are different in size for example from microscope images, as we will be working with here.

Here we will base our work on the article of Lindeberg¹ that gives an introduction to scale-space theory. Scale-space representation has made the basis for a range of image analysis methods and is extensively used in computer vision. In the exercise you will implement scale-space blob detection and use it for detecting and measuring the size of fibres that are imaged using X-ray CT.

The computation of scale-space is done by representing image features at all scales at once and detect features based on criteria that is independent of the scale. Here we will be working with the Gaussian scale-space, and the analysis is in practice done by smoothing the image using a Gaussian filter. In Lindeberg² the scale-space representation is defined for a general N -dimensional signal $f : \mathbb{R}^N \rightarrow \mathbb{R}$. Here we

¹ Tony Lindeberg, Scale-space: A framework for handling image structures at multiple scales. 1996

² Tony Lindeberg, Scale-space: A framework for handling image structures at multiple scales. 1996

will work with a 2D image $I : \mathbb{R}^2 \rightarrow \mathbb{R}$. For 2D image, its Gaussian scale-space representation is $L : \mathbb{R}^2 \times \mathbb{R}_+ \rightarrow \mathbb{R}$, which in practice becomes a 3D object, with the two spatial image dimensions (x, y) and the scale in the third dimension. Since scale is obtained by smoothing with a Gaussian, the variable determining the degree of smoothing is the variance t . Also the standard deviation $\sigma = \sqrt{t}$ is used in the article, but here we have simplified the notation and use only the variance t .

The Gaussian scale-space L is defined for N -dimensional signals by

$$L(x; t) = \int_{\xi \in \mathbb{R}^N} f(x - \xi) g(\xi; t) d\xi \quad (2.1)$$

with $g : \mathbb{R}^N \times \mathbb{R}_+ \rightarrow \mathbb{R}$ being the N -dimensional Gaussian kernel

$$g(x; t) = \frac{1}{(2\pi t)^{N/2}} e^{-(x_1^2 + \dots + x_N^2)/(2t)}. \quad (2.2)$$

In practice we will work with the Gaussian scale-space for 2D images on a discrete set of pixels. Therefore, we can write the Gaussian scale-space (ignoring boundary issues) as

$$L(x, y; t) = \sum_{-\gamma}^{\gamma} \sum_{-\delta}^{\delta} I(x - \gamma, y - \delta) g(\delta, \gamma; t) \quad (2.3)$$

where $g : \mathbb{R}^2 \times \mathbb{R}_+ \rightarrow \mathbb{R}$ is the 2D Gaussian kernel

$$g(x, y; t) = \frac{1}{2\pi t} e^{-(x^2 + y^2)/(2t)}. \quad (2.4)$$

Computing the scale-space is done at a discrete set of steps, and we have the start condition with $t = 0$ defined as $L(x, y; 0) = I(x, y)$.

For feature detection, we need to compute the derivatives of a scale-space representation. Note that this is conveniently achieved by convolving an image with a kernel that is a derivative of a Gaussian. Blob detection uses second order derivatives, more precisely the Laplacian of a Gaussian $\nabla^2 L = L_{xx} + L_{yy}$ which gives a high response where there is a blob in the image. To detect blobs we need to find local maxima and minima of the Laplacian.

What we still need to do is to ensure that we can detect blobs across different scales. The image in scale-space representation is increasingly smoothed, and with increasing scale t , pixels will change their intensity value towards the average value of the image. Therefore, the absolute values of derivatives will become smaller when increasing t . For blob detection, this means that the magnitude of the local maxima and minima in the scale-space of the Laplacian $\nabla^2 L$ will decrease and this smoothing must be compensated. The compensation factors for different features are given in Lindeberg³ and for the blob feature it is t such that the scale normalized Laplacian of Gaussian is $t\nabla^2 L$.

³ Tony Lindeberg, Scale-space: A framework for handling image structures at multiple scales. 1996

2.1 Exercise 2 – part I, Scale-space blob detection

In this part of the exercise you will implement scale-space blob detection with the purpose of detecting and measuring glass fibres from images of a glass fibre composite. An image example is given in Figure 2.1, that shows a polished surface of a glass fibre composite sample, where individual fibres can be seen. Since these fibres are relatively circular we will model them as circles. This means that we must find their position (center coordinate) and diameter, and for this we will use the scale-space blob detection. After having computed the fibres parameters, we will carry a statistical analysis of the results.

2.1.1 Computing Gaussian and its second order derivative

We will approach this analysis in steps that lead to the final algorithm. First we will use synthetic data to develop and test our algorithm, and after that we will carry out the analysis on the real images. Since we focus on blob detection, we must have a Gaussian kernel and its second order derivative. Since the Gaussian is separable, we can employ 1D filters for our analysis. The 1D Gaussian is given by

$$g(x) = \frac{1}{\sqrt{2\pi t}} e^{-\frac{x^2}{2t}}. \quad (2.5)$$

Suggested procedure

- Derive (analytically) the second order derivative of the Gaussian

$$\frac{d^2g}{dx^2}.$$

- Implement a function that takes the variance t as input and outputs a filter kernel of g and d^2g/dx^2 . You should use a filter kernel of at least $3t$. Why?
- Try the function on the synthetic test image `test_blob_uniform.png`.

2.1.2 Detecting blobs on one scale

Blobs can be found as spatial maxima (dark blobs) or minima (bright blobs) of the scale-space Laplacian

$$\nabla^2 L = L_{xx} + L_{yy}. \quad (2.6)$$

Suggested procedure

- Compute the Laplacian at one scale using the synthetic test image `test_blob_uniform.png`.

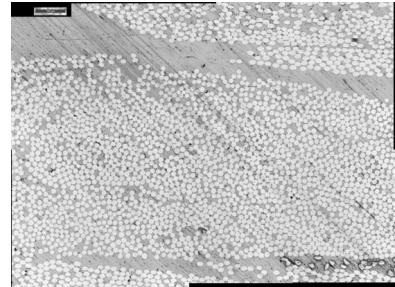


Figure 2.1: Example of fibre image acquired using an optical microscope.

2. Build a function that detects the coordinates of maxima and minima in the Laplacian image (detect blobs).
3. Plot the center coordinates and circles outlining the detected blobs. The radius of the circles should be $\sqrt{2t}$.
4. Try varying t such that the blobs in `test_blob_uniform.png` are exactly outlined.

2.1.3 Detecting blobs on multiple scales

To find blobs at multiple scales, we must use the scale-space representation. This can conveniently be done by representing $\nabla^2 L$ as a 3D array (volumetric image).

Suggested procedure

1. Decide on scales at which the Laplacian must be computed. You could make it equal steps in the size of the blobs ($\sqrt{2t}$).
2. Compute the scale normalized scale-space Laplacian $t\nabla^2 L$ for the test image `test_blob_uniform.png`.
3. Find coordinates and scales of maxima and minima in this scale-space and plot the detected blobs on top of the image. What are the detected scales?
4. Detect blobs in the test image `test_blob_varying.png`.

2.1.4 Detecting blobs in real data

We will now continue with the real images of fibers. The fibre data is obtained using different scanning methods including scanning electron microscopy (`SEM.png`), optical microscopy (`Optical.png`), synchrotron X-ray CT (`CT_synchrotron.png`), and three resolutions of laboratory X-ray CT (`CT_lab_high_res.png`, `CT_lab_med_res.png`, `CT_lab_low_res.png`). The CT data is a single slice very close to the top, so we assume the data to be from the same part of the sample, and this allows us to directly compare the fibers. We will do this comparison in next exercise, but in this we will compute the fiber location and their diameter. In Figure 2.2 you can see a visualization of the fibre data from the high resolution X-ray CT scan.

We start by testing the blob-detection on this real data.

Suggested procedure

1. Run your blob-detection function from above on a cut-out example one of the images. It is important that you tune your parameters to get the best possible results.

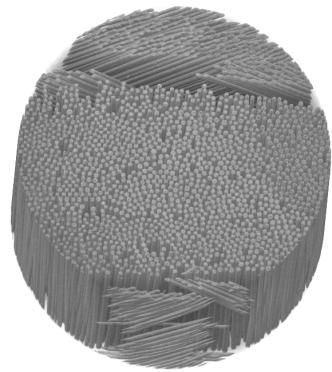


Figure 2.2: Visualization the 3D fibers scanned with the high resolution X-ray CT-scanner.

2.1.5 Localize blobs

It is difficult to detect blobs in the scale-space Laplacian, such that all fibers are found. To overcome this, we will detect the fibers as maxima in a Gaussian smoothed image. Since the fibers are almost the same size, we can use a single scale of the Gaussian to detect the fiber centers.

Suggested procedure

1. Smooth an image of fibers with a Gaussian and visualize the result.
2. Find locations of maxima in this image and plot the positions on top of the original image.
3. Compute the scale-space Laplacian for the image.
4. Find the scale of each fibre as the minimum over scales at the fiber locations.
5. Plot circles according to the found scale on top of the original image.
6. Detect fibers in all six fiber images. Save the locations and diameters.

In the next exercise, where you will work with image matching based on SIFT features, you will use the results obtained in this exercise. So, next time it will be possible to continue working on the parts that you did not finish here.

2.1 Information for 02506

This is the first part of the exercise in Part I, and will be continued the following two weeks. Please prepare a short (5 minutes presentation) of your code and results.

Exam questions related to this exercise

- Explain scale-space blob detection. This includes basic theory of scale-space blob-detection, explanation of the algorithm you made, and discussion of results.

3 Feature-based segmentation

SEGMENTATION OF AN IMAGE is done by partitioning the image. This can be described by a function $g(x, y) \rightarrow \ell$ that for each position (x, y) in the image I maps to a label $\ell \in 1, \dots, L$. Here we will do segmentation by assigning a label to each pixel in the image.

The simplest pixel labeling is obtained by classifying each pixel according to its intensity. In a gray-scale image, this will typically be one or more threshold values that separates the intensities into groups. For many image segmentation problems, this is not sufficient for obtaining a suitable segmentation as illustrated by the bone example in Figure 5.5. Here the fine structures cannot be separated by the pixel intensity alone, but we must utilize information in the image patterns such as size and shape of the depicted structures.

Instead of using the intensity in one pixel, we can use a neighborhood around a pixel to obtain our segmentation. Hereby, we capture the local appearance of the image, and we will try to use that for labeling each pixel in the image. The local appearance is also known as image texture.

3.1 Supervised feature-based segmentation

We will approach the segmentation as a supervised labeling problem, where we learn the parameters of our segmentation model from training data. The training data is obtained from labeled training images, which e.g. can be done manually. Now the segmentation problem is to transfer the labels from this image to an unknown test image of the same type, meaning an image with similar appearance.

3.1.1 Image features

The local appearance of the image can be computed as image features, and in this exercise we will compute two image features. The first feature is obtained by computing a set of image derivatives by convolving with Gaussians and its derivatives. The second is obtained as

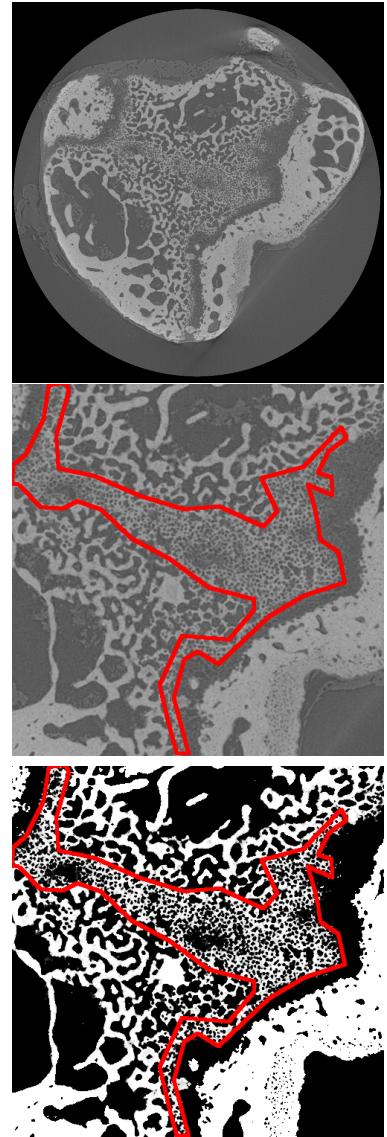


Figure 3.1: Slice of a CT-scan of a trabecular mouse bone. Top image shows the full image, middle is a zoom on the central part with the fine trabecular structure outlined by the red line, and bottom shows the same with a threshold.

patches centered on a pixel. There are many other features that we could choose, e.g. SIFT¹, SURF² or HOG³ features, but here we choose relatively simple features that are easy to compute.

Features from a Gaussian and its derivatives Similar to the previous exercises, we can smooth and compute image derivatives by convolving with a Gaussian kernel and its derivatives. The area over which we want to compute the Gaussian is determined by the scale of the Gaussian kernel.

The Gaussian features are obtained as a stack of Gaussian derivatives. Let us use the same notation of Gaussian derivatives as we did in the scale-space exercise, such that we have

$$L = I * g(t) ,$$

where L is an image I convolved with a Gaussian

$$g = \frac{1}{\sigma\sqrt{2\pi}} e^{\frac{-x^2}{2\sigma^2}}$$

at scale t , where $t = \sigma^2$. The images obtained from convolution with Gaussian derivatives are given by

$$L_x = I * \frac{\partial g}{\partial x} , \quad L_y = I * \frac{\partial g}{\partial y} , \quad L_{xx} = I * \frac{\partial^2 g}{\partial x^2} , \quad L_{xy} = I * \frac{\partial^2 g}{\partial x \partial y} ,$$

etc. In this exercise, we will compute the higher order derivatives until the fourth order, which will result in a 15-dimensional descriptor that captures the local appearance of the image. The descriptor is formed by stacking the Gaussian convolved images, such that each pixel position has an associated 15-dimensional feature vector. That is, we have a $r \times c \times 15$ feature image

$$\begin{aligned} F = [& L, L_x, L_y, L_{xx}, L_{xy}, L_{yy}, L_{xxx}, L_{xxy}, L_{xyy}, L_{yyy}, \\ & L_{xxxx}, L_{xxxxy}, L_{xxyy}, L_{xyyy}, L_{yyyy}] . \end{aligned}$$

Patch-based features Another way of computing image features is by extracting small patches centered on a pixel and concatenating the pixels into a feature vector. If we e.g. have a 9×9 image patch, this will result in an 81 dimensional feature vector for each pixel position.

3.1.2 Probabilistic clustering-based segmentation

The basic idea in the segmentation model that we will use, is that parts of the image that have similar appearance should have the same label. Since the features encode the local appearance of the image, it means that we can give features that are similar the same label.

¹ David G Lowe. Distinctive image features from scale-invariant keypoints. *International journal of computer vision*, 60(2):91–110, 2004

² Herbert Bay, Tinne Tuytelaars, and Luc Van Gool. Surf: Speeded up robust features. In *European conference on computer vision*, pages 404–417. Springer, 2006

³ Navneet Dalal and Bill Triggs. Histograms of oriented gradients for human detection. In *2005 IEEE computer society conference on computer vision and pattern recognition (CVPR'05)*, volume 1, pages 886–893. IEEE, 2005

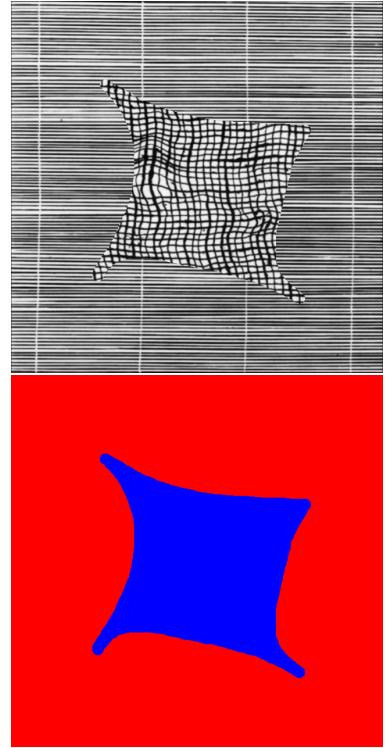


Figure 3.2: Top shows a training image with ground truth labeling. Red is one label and blue is another. Typically, this will be two scalar values, e.g. {0, 1}.

Using an already labeled image, we can learn the desired labels of the features. In practice, we typically have labels given as a separate image of the same size as the training image. This is shown in Figure 3.2 for one of the two-label images that you will work with in this exercise. Since the training image and the label image are of the same size, for each feature from the training image we can look up the label at a corresponding position in the label image.

By computing image features in both a training and a test image, we can transfer the labels from training image to the test image. This is done by using a similarity measure of the image features, and here we will use Euclidean distance

$$d(f_p, f_q) = \sqrt{\sum_{i=1}^n (f_p(i) - f_q(i))^2},$$

where f_p and f_q are two n -dimensional feature vectors.

A direct approach for labeling the image would be to compute the Euclidean distance from each feature vector in the test image to each feature vectors in the training image (or a random subset of the features in the training image). By selecting the label of one or more of the nearest feature vectors, we could obtain a labeling. If we select more than one feature, we can e.g. label according to majority vote. This would be using k -nearest neighbor classifier, which might not always be robust, because outliers could introduce noise.

A more robust approach is using k -means clustering, where each cluster center makes up a representative feature vector. This is sometimes referred to as a dictionary-based approach, where each cluster center is referred to as a *visual word*, and the collection of cluster centers make up a dictionary. Each of the clusters is made up of a number of feature vectors from the training image. This allows us to compute the probability of a given label for each feature cluster

$$p_C(\ell = \lambda) = \frac{1}{m} \sum_{f \in C} \delta(\ell(f) - \lambda), \quad (3.1)$$

where C is a feature cluster with m elements f , $\ell(f)$ is the label of feature element f , and

$$\delta(x) = \begin{cases} 1 & \text{if } x = 0 \\ 0 & \text{otherwise} \end{cases}. \quad (3.2)$$

3.2 Exercise

You should implement a probabilistic clustering-based segmentation using Gaussian features. If time allows, you should also implement the same using image patches.

There is a set of images available for training and testing named `train_{A,B,C}_image.png` and `test_{A,B,C}_image.png` and `ground_truth.png`, which is the label image. Furthermore, there is a set of images of a bone, which only has labels for training but not for test.

3.2.1 Gaussian features

The function for computing the Gaussian feature image is available as the functions `get_gauss_feat_im.m` for MATLAB and `get_gauss_feat_im` function in the `feature_based_segmentation.py` file for Python.

Suggested procedure

1. Read in the images and display them.
2. Compute the feature image.
3. Transform the feature image into a matrix where each column is a feature vector.

3.2.2 Clustering for building the dictionary

You should use k -means clustering for building the dictionary. Using all feature vectors would be too time consuming, and it is sufficient to select a random subset of features. You should select both features and image labels from the same pixel positions.

Suggested procedure

1. Select a random subset of feature vectors with corresponding labels (you can use random permutation). If you choose e.g. 5000-10000 vectors, it should be sufficient for clustering.
2. Cluster the feature vectors into a number of clusters. You can choose e.g. 100-1000 clusters.
3. Make a $2 \times n$ matrix to store label probabilities, where n is the number of cluster centers. Compute the probability of a cluster belonging to each of the two labels using Eq. 3.1 and 3.2, and store the probabilities in the matrix.

3.2.3 Computing probability image and segmenting

You should now extract image features from the test image. For each of the features you should find the nearest cluster center and assign the label probability to that cluster to the pixel position of the feature. This will result in a probability image of size $r \times c \times 2$, where each pixel has a probability of belonging to one of the two labels.

Suggested procedure

1. Compute a feature image from the test image.
2. Use a nearest neighbor algorithm (`knnsearch` in MATLAB or `Nearest Neighbors` from Scikit Learn in Python) and find the nearest cluster.
3. Build a probability image based on the nearest cluster for each pixel.
4. Obtain a segmentation by selecting the most probable label in each pixel. (You can obtain an improved segmentation by smoothing the probability image).

3.2.4 Segmentation with patch-based features

Do the same as above using image patches instead of Gaussian features. This is a little more difficult because of boundary effects. You can extract image patches using the `im2col` function in MATLAB, and we have provided an `im2col` function found in the `feature_based_segmentation.py` file for Python, that has the same functionality.

3.1 Information for 02506

This is the second part of the exercise in Part I, and will be continued next week. Please prepare a short (5 minutes presentation) of your code and results.

Exam questions related to this exercise

- Explain feature-based segmentation using k -means clustering (dictionary-based segmentation).

4 Feature-based registration

IMAGE REGISTRATION is the process of transforming one image (the moving image) to another (the target image). We will do this by matching image features. Here we will use SIFT features that is described in detail in ¹.

The concept of image features from interest points is essential in image analysis. The basic idea is to identify salient positions in an image (unique key-points) and use the surrounding image context to describe the image locally in the form of a descriptor vector. This principle has been extensively used for solving many problems including object recognition, image retrieval (image search), image stitching, geometric reconstruction, etc. Solutions to many of these problems have however been improved by machine learning methods and especially convolutional neural networks have replaced feature-based analysis in a range of applications.

Interest point features are often referred to as hand-crafted features whereas neural networks learn the image features. Since the end-to-end optimization employed in neural networks often results in an improved performance, much work now focuses on neural network methods. There are, however, still problems where interest point features will be a good choice, e.g. for matching images (feature-based image registration), i.e. finding point-wise correspondence between images. Some deep learning based alternatives to hand-crafted features have been suggested², showing superior performance. But due to the simplicity of hand-crafted features, where their design make them easy to compute without the time consuming learning part, we will here work with interest point features.

Scale invariant feature transform (SIFT) described in³ is a widely used interest point feature. Here we focus on SIFT. But bear in mind that many other interest point features have similar properties⁴.

We will use SIFT for feature-based image registration. The problem we address is that we are given two images that depict the same object and we want to find an image transformation that aligns the two images. This allows us to compare structures found in one image

¹ David G Lowe. Distinctive image features from scale-invariant keypoints. *International journal of computer vision*, 60(2):91–110, 2004

² Kwang Moo Yi, Eduard Trulls, Vincent Lepetit, and Pascal Fua. Lift: Learned invariant feature transform. In *European Conference on Computer Vision*, pages 467–483. Springer, 2016

³ David G Lowe. Distinctive image features from scale-invariant keypoints. *International journal of computer vision*, 60(2):91–110, 2004

⁴ Tinne Tuytelaars, Krystian Mikolajczyk, et al. Local invariant feature detectors: a survey. *Foundations and trends® in computer graphics and vision*, 3(3):177–280, 2008

with structures found in the other image. In some cases the difference between images can be modeled by a rotation \mathbf{R} , translation \mathbf{t} and scale s . This is e.g. the case in microscopy or CT, where pixels have a fixed physical size. A more general affine transformation can be computed using a homography, but in cases where the imaging system can only transform the image by a rotation, translation and scale, these will be the appropriate parameters to compute. Here we will work with sets of corresponding 2D points, where the correspondence is found by matching SIFT features.

Fitting two sets of 2D points in least squares sense. For two 2D point sets \mathbf{p} and \mathbf{q} with corresponding elements \mathbf{p}_i and \mathbf{q}_i (2×1 column vectors), where $i = 1, \dots, n$, we are interested in finding a rotation \mathbf{R} , a translation \mathbf{t} and a scale s that minimizes the squared distance between the two point sets. The transformation from one point set to the other is given by

$$\mathbf{q}_i = s\mathbf{R}\mathbf{p}_i + \mathbf{t}. \quad (4.1)$$

The scale s can be found e.g. by computing the ratio between average distance to the centroid of each point set

$$s = \frac{\sum_{i=1}^n \|\mathbf{q}_i - \mu_q\|}{\sum_{i=1}^n \|\mathbf{p}_i - \mu_p\|}, \quad (4.2)$$

where $\mu_p = \frac{1}{n} \sum_{i=1}^n \mathbf{p}_i$ and $\mu_q = \frac{1}{n} \sum_{i=1}^n \mathbf{q}_i$ are the centroids of \mathbf{p}_i and \mathbf{q}_i respectively.

Now we want to find the rotation and translation that minimize

$$\sum_{i=1}^n (\mathbf{q}_i - s\mathbf{R}\mathbf{p}_i - \mathbf{t})^2.$$

One way of least-squares fitting 2D point sets involves 2-by-2 covariance matrix (normalization is not needed)

$$\mathbf{C} = \sum_{i=1}^n (\mathbf{q}_i - \mu_q)(\mathbf{p}_i - \mu_p)^T,$$

and its singular value decomposition

$$\mathbf{U}\Sigma\mathbf{V}^T = \mathbf{C}.$$

Here, \mathbf{U} and \mathbf{V} are the left and right singular matrices respectively and Σ is a diagonal 2-by-2 matrix containing the singular values. From this we obtain the rotation

$$\hat{\mathbf{R}} = \mathbf{U}\mathbf{V}^T. \quad (4.3)$$

In rare cases this computation can result in a reflection instead of a rotation. If the determinant of $\det(\hat{\mathbf{R}}) = 1$ it is a rotation and if

$\det(\hat{\mathbf{R}}) = -1$ it is a reflection. This computation will typically only result in a reflection e.g. if there are only two points for determining the rotation. Therefore, we can compute the rotation taking this into account by

$$\mathbf{R} = \hat{\mathbf{R}}\mathbf{D}, \quad (4.4)$$

where

$$\mathbf{D} = \begin{bmatrix} 1 & 0 \\ 0 & \det(\hat{\mathbf{R}}) \end{bmatrix}. \quad (4.5)$$

Finally, we find the translation as the average vector from points in \mathbf{q} to the rotated points in \mathbf{p}

$$\mathbf{t} = \frac{1}{n} \sum_{i=1}^n (\mathbf{q}_i - s\mathbf{R}\mathbf{p}_i) = \mu_q - s\mathbf{R}\mu_p. \quad (4.6)$$

See e.g. Arun et al.⁵ which covers a more general 3D case.

4.1 Exercise on feature-based registration

This exercise aims at finding correspondence between images by matching SIFT features and computing the transformation, i.e. the rotation, translation and scale between the two images. For computing the SIFT features you can use `vlFeat` for MATLAB or `OpenCV` for Python. In the extra exercise you can use the transformation obtained here to compare the fiber diameters that you got from using blob detection for fibers. But for now, the purpose is to compute the transformation.

4.1.1 Rotation, translation and scale

You should implement a function that takes two point sets as input and returns the rotation, translation and scale. To ensure that you have the correct implementation, you can make a set of random 2D points that you rotate and translate with known parameters. By applying the transformation to the points, you can verify that you obtain the correct transformation parameters.

4.1.2 Compute and match SIFT

Here you should compute SIFT features in two images and match them using Euclidean distance between their descriptor vectors. You should use the criterion by Lowe where a correct match is found if the fraction between the closest feature vector and the second closest feature vector is less than a certain value, e.g. 0.6. There are functionality for matching features in both `vlFeat` and `OpenCV` that you can use. You can also

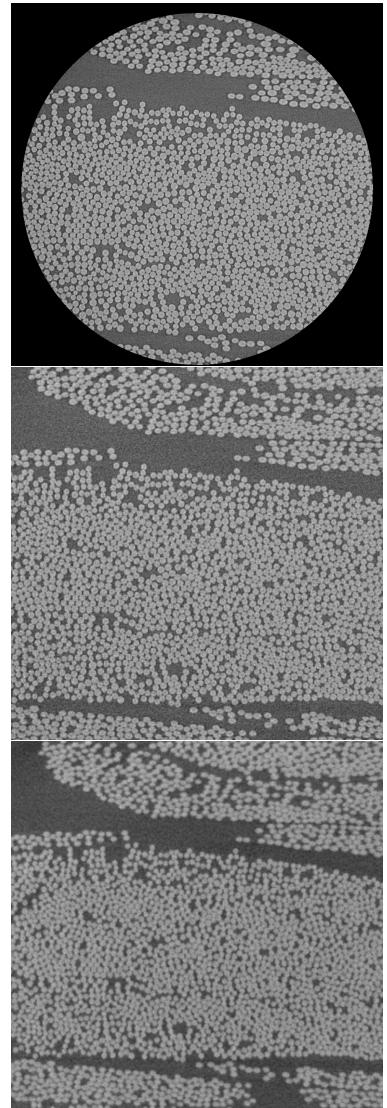


Figure 4.1: Example of an image of the same fiber sample acquired using a CT scanner at three resolutions.

⁵ K Somani Arun, Thomas S Huang, and Steven D Blostein. Least-squares fitting of two 3-D point sets. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, (5):698–700, 1987

implement your own matching function where you take e.g. scale and rotation into the matching criterion.

To make sure that you compute and match the SIFT features, you can make two corresponding images by rotating and scaling one image and matching it to the original. After that, you can try to match the CT-images of fibers at the three resolutions shown in Figure 4.1. Visualize the matching feature points by drawing lines between them e.g. as shown in Figure 4.2. This allows you to visually evaluate the matching.

4.1.3 Transform the matched features

Now you should combine the function for computing the transformation parameters and the SIFT feature matching. The matching will not be perfect and you will most likely see some wrongly matched features. The larger the difference in appearance or scale of the image that is being matched, the more wrongly matched features can be expected. If the majority of the correspondences are correct then the least squares fit will give a relatively good result.

Since we are computing the transformation by a least squares fit the outliers will affect the result to some extent. Outliers can however be removed relatively easily. You can compute the Euclidean distance between the two point sets after you have aligned them. There you will see that most of the distances are relatively small. And if you remove matching points with a distance larger than a certain threshold, you can repeat the computation of the transformation and obtain higher precision in the matching. You should implement a function that makes this two step computation of the transformation and choose a good criterion for a threshold.

Illustrate your transformed feature points by plotting the two point sets on top of each other in the image e.g. as illustrated in Figure 4.3. You should be able to see a difference in the precision of the matching after removing the outliers.

4.1 Information for 02506

Exam questions related to this exercise

- Explain the principle of feature-based image registration and how SIFT can be used for feature matching.

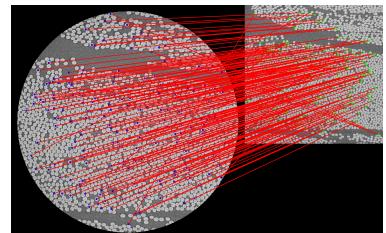


Figure 4.2: Matching SIFT features illustrated by red lines.

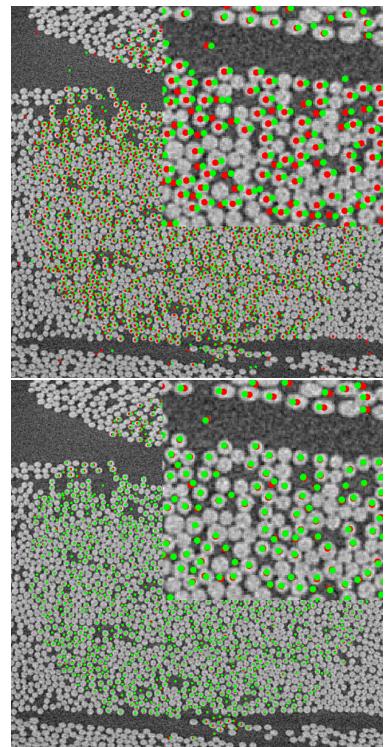


Figure 4.3: Matching features shown in red and green (zoom in upper right corner). Top is after computing least squares of all matching features and bottom is a recomputed match after removing outliers.

Part II

Image analysis with geometric priors

IN THE CONTEXT OF IMAGE ANALYSIS, the term *prior knowledge* refers to all the information about problem that is available in addition to the image data. There are numerous ways of incorporating priors when solving image analysis problems, and here we look into Markov random fields which are used to model local contextual information, and two models for handling a distinctive geometry: parametric deformable curves and layered surfaces.

Information for 02506

The second part of the course, image analysis with geometric priors, is covered in three weeks of the course. Counting from the start of semester, these are weeks 5, 6, and 7.

- Week 5: Markov random fields (MRF) modeling and optimization using graph cuts. In addition to lecture notes, the optional reading material is a book by Stan Li, which is freely available for download at DTU Findit. Exercises in MRF is on bone segmentation.
- Week 6: Mumford-Shah functional, Chan-Vese algorithm and snakes. In addition to lecture notes, the optional reading material is an article by Chan and Vese, and a book chapter on snakes. For exercises in deformable models you will track a simple object in a sequence of images.
- Week 7: Layered surface detection. In addition to lecture notes, the reading material is an article by Li et al. For exercise you will detect layers in bone data. Additionally, you can combine layered surface detection with deformable model and segment nerves from volumetric data.

This was changed to:

- Week 7: Applicability of the introduced geometric models for solving image segmentation problems. There will be no additional reading material introduced this week. Exercise involves a nerve segmentation task. You will be using and adapting functions developed in week 5 and 6. *The deliverable* is a one-pager on the nerve segmentation.

5 Markov random fields

MARKOV RANDOM FIELDS (MRF) is a probabilistic framework that can be used for various image analysis tasks, where contextual information needs to be considered. All MRF formulations involve labeling, i.e. assigning labels to some image entity, typically assigning labels to pixels. MRF are characterized by the Markov property, i.e. that the probability of a (pixel) label is only dependent on the local neighborhood (of a pixel).

In exercises on MRF, we will use MRF model for image segmentation. A segmentation can be solved by assigning a discrete label to each pixel in the image according to the pixel intensity. Often, we would like the segmentation to be smooth, meaning that the probability of a pixel label being different from label of neighboring is low. Provided an image, we aim at finding a label configuration that maximizes the *a posterior* (MAP) probability which is a combination of a likelihood (data) term and the term modelling a smoothness prior.

One characteristics of MRF is that the probability of the MRF configuration is an exponential of the negative configuration energy. Maximizing the posterior probability is therefore equivalent to minimizing the posterior energy of the configuration f given by

$$E(f) = U(f|d) = U(d|f) + U(f) \quad (5.1)$$

or, in terms of clique potentials

$$E(f) = \sum_{\{i\} \in \mathcal{C}_1} V_1(f_i) + \sum_{\{i, i'\} \in \mathcal{C}_2} V_2(f_i, f_{i'})$$

where \mathcal{C}_1 is the set of one-cliques, V_1 is a one-clique potential used for modeling the likelihood term, \mathcal{C}_2 is the set of two-cliques, and V_2 is a two-clique potential used for modeling the prior term.

As discussed in the book by Li¹, Chapter 1, Introduction, first paragraph, the main concerns of the MRF framework are how to define an objective function, i.e. clique potentials (modelling part), and how to find the optimal solution for a given objective function (optimization part).

¹ Stan Z Li. *Markov random field modeling in image analysis*. Springer Science & Business Media, 2009

5.1 Example: Gender determination (optional, written for students wanting an easy introduction to MRF)

We start with the extremely small 1D example with the aim of introducing MRF terminology, demonstrating the modelling possibilities provided by MRF, and the use of the data term and likelihood. A student comfortable with these MRF concepts may skip the example and proceed with the exercises.

Imagine entering a bar and observing 6 persons standing along the counter. You estimate persons heights (in cm) and record this data as

$$d = \begin{bmatrix} 179 & 174 & 182 & 162 & 175 & 165 \end{bmatrix}.$$

You want to estimate the persons gender, i.e. you want to assign either a label M or F to each person by combining a data term and your knowledge of the contextual information. You decide to pose the problem as a MRF with the neighbourhood given by the first neighbor (person to the left and person to the right).

We first consider the likelihood (data) term. You know that the average male height is 181 cm, and the average female height is 165 cm, and that height for each gender may be described as following a normal distribution where you assume the standard deviation for both genders being the same. For this reason you define the likelihood terms as one clique potentials

$$V_1(f_i) = (\mu(f_i) - d_i)^2$$

where d_i is the height of person i , f_i is a label assigned to person i , and μ is defined as $\mu(M) = 181$, $\mu(F) = 165$. The likelihood energy of a configuration $f = [f_1 \dots f_6]$ is the sum of all one-clique potentials

$$U(d|f) = \sum_{i=1}^6 V_1(f_i).$$

To find the configuration which minimizes the likelihood energy you can consider the one-clique potentials for all i and both labels

$$\begin{array}{rccccc} (\mu(M) - d_i)^2 & : & 4 & 49 & 1 & 361 & 36 & 256 \\ (\mu(F) - d_i)^2 & : & 196 & 81 & 289 & 9 & 100 & 0 \end{array}$$

Obviously, the minimal likelihood energy is obtained if we choose a label which minimizes the cost for each i , resulting in a labeling

$$f^D = \begin{bmatrix} M & M & M & F & M & F \end{bmatrix}, \quad (5.2)$$

and giving $U(d|f^D) = 99$. Another thing to notice is that additional cost for deviating from this labeling varies, depending on which label we change. For example, it costs additional 352 to label the forth person

as male, while it costs only additional 32 to label the second person as female.

Now you want to incorporate the contextual (prior) information about the gender of the people standing along the bar counter. Your experience is that people tend to group by gender, and that a configuration with a man standing next to a woman occurs less frequently. For this reason, you decide to incorporate a cost β which penalizes a less-frequent configuration. For prior energy you define 2-clique potentials as

$$V_2(f_i, f_{i+1}) = \begin{cases} 0 & \text{if } f_i = f_{i+1} \\ \beta & \text{otherwise} \end{cases}$$

The prior energy is the sum of all 2-clique potentials for all 2-cliques (all pairs of neighbors) in a configuration.

$$U(f) = \sum_{i=1}^5 V_2(f_i, f_{i+1})$$

Obviously, this prior energy is minimal for a configuration with all labels being equal, while spacialy alternating labels yield maximal prior energy of 5β . The prior energy for the configuration f^D which minimizes the likelihood energy (5.2) is $U(f^D) = 3\beta$.

Last modelling choice involves setting a suitable parameter β . This choice depends on your confidence in the prior, compared to the data. You choose to use $\beta = 100$. According to (5.1), the posterior energy of configuration f^D is

$$U(f^D|d) = U(d|f^D) + U(f^D) = 99 + 300 = 399.$$

The question is, can we find another configuration which yields a smaller posterior energy? And finally, which configuration minimizes posterior energy?

For our small problem, we can simply try different configurations, two are worth considering

$$f^P = [M \ M \ M \ M \ M \ F] ,$$

$$f^O = [M \ M \ M \ F \ F \ F] .$$

Relatively easy we can confirm that configuration f^O with $U(f^O|d) = 163 + 100 = 263$ is an optimal configuration.

Note again that the prior knowledge encodes our assumptions of how the solution is supposed to look like, and it also influences the result such that what we find is what we expect to find. Our experience (prior knowledge) about people standing in a bar might have motivated another prior energy which encourages configurations where males

and females stand next to each other. For example a prior

$$V_2(f_i, f_{i+1}) = \begin{cases} \beta & \text{if } f_i = f_{i+1} \\ 0 & \text{otherwise} \end{cases}$$

This prior would yield in another optimal configuration.

Note also the distinction between modeling (setting up the problem by defining a likelihood term and a prior term) and optimization (finding the configuration which minimizes the posterior energy).

5.2 Exercise: MRF modelling

In this exercise we define an objective function for segmenting a noisy image, similar to the problem in Li Section 3.2.2. Here we will compute the energy of different configurations to confirm that minimizing an objective function leads towards the desired solution. The model we use is very similar to the model used for gender determination. In this exercise we use synthetic data (i.e. we produce our input data by adding noise to a ground truth image) shown if Figure 5.1. This allows us to evaluate the quality of our objective function. In the text the input image is denoted D and ground truth segmentation S_{GT} where elements of S_{GT} are from the set $\{1, 2, 3\}$ corresponding to the darkest, medium gray, and brightest class.

Write a function that given D and a segmentation, for example S_{GT} , produces a histogram of the pixel intensities and histograms of the pixel intensities divided into segmentation classes. An example is shown in Figure 5.2. What does this histogram say about the chances of obtaining a reasonable segmentation of D using a method which considers only pixel intensities, for example thresholding?

Now we pose image segmentation as a MRF. Sites are pixels, labels are from $\{1, 2, 3\}$, and we choose a first-order neighborhood (four closest pixels). As in the previous example, we define the one-clique potentials for the likelihood energy as the squared distance from the class mean

$$V_1(f_i) = \alpha (\mu(f_i) - d_i)^2$$

where d_i are intensities of the (noisy) image, f_i are pixel labelings given by the configuration, and μ is estimated from the histogram and set to $\mu(1) = 70$, $\mu(2) = 130$, $\mu(3) = 190$. In this example, we will weight the data term with the parameter α (unlike the previous example, where we weight the prior term). You may use $\alpha = 0.0005$. The likelihood energy is defined similar to before

$$U(d|f) = \sum_i V_1(f_i),$$

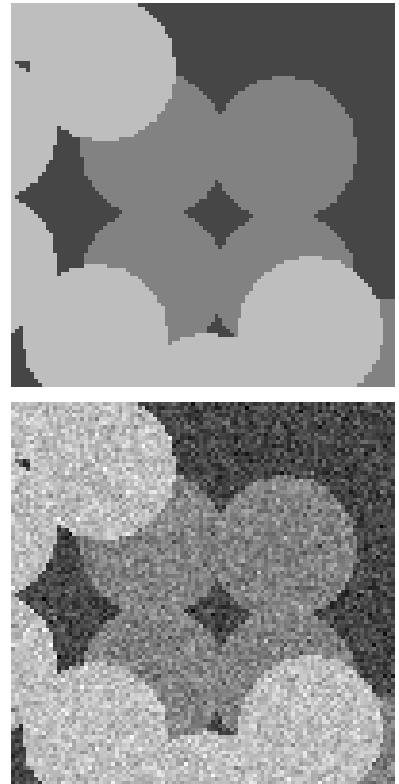


Figure 5.1: A ground truth (the desired segmentation should resemble ground truth) and a noisy image (input data).

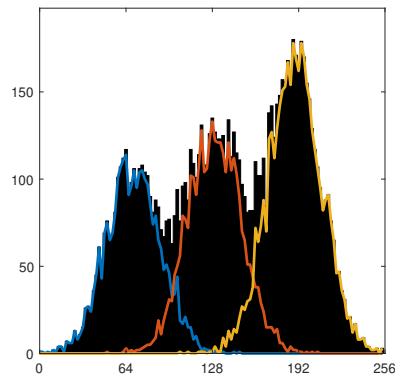


Figure 5.2: Histograms.

where summation covers all image pixels. Similarly to the previous example, we define 2-clique potentials for discrete labels which penalizes neighbouring labels being different

$$V_2(f_i, f_{i'}) = \begin{cases} 0 & f_i = f_{i'} \\ 1 & \text{otherwise} \end{cases},$$

and prior energy

$$U(f) = \sum_{i \sim j} V_2(f_i, f_j)$$

where summation runs over all pairs of neighbouring pixels. The posterior energy is now given by (5.1).

We want to check that our optimal function leads to the desired result. That is, we want to make sure that posterior energy gets smaller when we approach the desired result. Therefore we want to compute the likelihood, prior and posterior for some reasonable segmentations (MRF configurations). For the purpose of this testing, produce at least two segmentations of the noisy image D . The first segmentation S_T is obtained by thresholding D at intensity levels 100 and 160 (valleys of the histogram). The second segmentation S_M is computed by median filtering S_T using an appropriate kernel. You are welcome to produce additional configurations, e.g. by applying a Gaussian filter to D prior to thresholding, or by using morphological operations. For all candidate configurations you should take a look at the intensity histograms of three classes, similarly as for the S_{GT} earlier.

Write a function which given an image D , a configuration S and the MRF parameters μ (intensities for segmentation classes) and α (weighting of the data term), returns computed likelihood, prior and posterior energy. Use your function for computing energies of different configurations, and also the ground truth S_{GT} . If we consider only the likelihood, which configuration is the most probable? If we consider only the prior energy, which configuration is the most probable? What if we consider the posterior energy? Would you expect that minimizing the optimal energy leads to a good segmentation? If not, try adjusting α and improve the posterior.

Tasks

1. Implement a function which produces histograms, as explained in the text.
2. Implement a function which computes segmentation energies, as explained in the text.
3. Produce a number of configurations. Apply your two functions to all configurations, and answer the questions from the text.

5.3 Exercise: Iterative optimization for MRF (optional, for students wanting additional challenges)

The interactions modelled by MRF prior make optimization (finding an optimal configuration) of the MRF very difficult. Standard MRF optimization methods may be very slow, but efficient graph cut algorithms can be used for a subset of problems. To gain a better understanding of MRF you may first implement an standard MRF optimization called iterated conditional modes (ICM), briefly sketched in Li Section 3.2.2. and elaborated in Section 9.3.1. Later, in the exercises following this one, you will be given graph cut implementation which you will use for optimization. You may also chose to first solve the problem using graph cuts, and then return to this exercise.

The general principle of ICM is the following. Every pixel contributes to the overall energy only locally, so for each pixel we can find a label that locally minimizes the energy (i.e. maximizes the conditional probability given all other labels). The iterative process continues until convergence. In our case, for a pixel i we need to compute

$$V(f_i|d_i, f_{\mathcal{N}_i}) = \alpha (\mu(f_i) - d_i)^2 + \#\{f_i \neq f_{i'} | i' \in \mathcal{N}_i\}$$

(see Li Equation (9.15) from Section 9.3.1) for three possible labels $f_i \in \{1, 2, 3\}$ and choose the label yielding the lowest value. The symbol $\#$ denotes the number of neighbors of i which have a label different from f_i .

To implement ICM, write a function which takes as input a segmentation S , the data term D (the noisy image), and MRF parameters μ and α . The function should output conditional local potential, i.e. values $V(f_i|d_i, f_{\mathcal{N}_i})$ for all pixels and all labels. For our purpose the output should have three layers, each layer as big as the image, such that the k -th layer gives a pixel-wise local energy for label k . You can use your function to iteratively improve the configuration by labeling each pixel with the locally optimal label.

The convergence of ICM is guaranteed only for serial updating (updating labels one after another). To see why, we will first try parallel update (updating all labels at once), where we in each iteration at once overwrite all labels of the current configuration with locally optimal labels. Start for example with S_T and run for 10 or 20 iterations. What do you observe? Does the algorithm converge?

Instead of a fully serial update, we can in parallel update a set of labels where no two sites are neighbors. In our case, this can be obtained by dividing all pixels in two sets using a checkerboard pattern. Why can we update such sets in parallel, and why use a checkerboard pattern in our case? Modify the algorithm such that you in each iteration compute conditional local potentials (using the function you wrote) and update

half of the pixels according to checkerboard pattern, then compute local potentials again and update the other half of the pixels. Does the algorithm converge?

Compute the posterior energy for the configuration obtained using ICM, and compare with the energies for configurations given by ground truth and all the segmentation configurations. Did you come closer to the optimal configuration? Try also starting with a random initialization of the labels. Do you obtain the same result regardless of the initialization? Try reducing and increasing the smoothness by changing the weighting of likelihood and prior. For example, make α 10 times bigger or 10 times smaller. What do you observe?

Optionally, you can try implementing another popular and well-known optimization algorithm. The Gibbs sampling algorithm (Li Section 7.1.6) is a randomized sampling algorithm for finding the optimal configuration of the MRF. It is based on changing the labeling f_i with a probability which is proportional to the probability of the labelings f_i . The sketch of the Gibbs sampling algorithm is as follows. Initialized based on the maximum likelihood. Iterate for a number of times. In each iteration compute the local probability of each label every pixel. In our case this is easily obtained from the output of your function by taking an exponential of negative energy and normalizing probabilities to sum to 1. For each pixel, divide the interval $[0, 1]$ according to probabilities, then choose a random number from $[0, 1]$ and determine which subpart it belongs to – this indicates the label which should be assigned for the pixel. The Gibbs sampler might be further improved using simulated annealing (Li Section 10.1). An easy way of implementing simulated annealing is to multiply conditional local potentials with a increasing value, for example iteration number.

5.4 Example: Graph cuts for MRF

A binary (two label) MRF problem with submodular second order energy (loosely speaking an energy favoring smoothness) can be exactly solved by finding a minimum $s-t$ cut of a graph constructed from the energy function^{2,3,4}. A minimum $s-t$ graph cut can be found e.g. using the Ford and Fulkerson algorithm, or an efficient freely available graph cut implementation by Boykov and Kolmogorov. A multiple-label discrete MRF problem can also utilize graph cuts via iteratively solving multiple two-label graph cuts, e.g. by using α expansion.

In the following exercises we are using graph cuts to optimize discrete MRF. MATLAB users may use the provided code, in particular the `GraphCutMex` function. This is a slightly modified version of the older Boykov implementation, the newest version can be found at <http://pub.ist.ac.at/~vnk/software.html>. Python users may use `PyMaxflow`

² Yuri Boykov, Olga Veksler, and Ramin Zabih. Fast approximate energy minimization via graph cuts. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 23(11):1222–1239, 2001

³ V Kolmogorov and R Zabih. What energy functions can be minimized via graph cuts? *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(2):147–159, 2004

⁴ Y Boykov and V Kolmogorov. An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(9):1124–1137, 2004

package documented at <http://pmneila.github.io/PyMaxflow/index.html>.

To begin with, we look back at the small example with gender labeling. Recall that the heights (in cm) of 6 persons are

$$d = \begin{bmatrix} 179 & 174 & 182 & 162 & 175 & 165 \end{bmatrix}$$

and we want to estimate the persons gender. For likelihood we use squared distance from the means $\mu(M) = 181$, $\mu(F) = 165$. For the prior we use $\beta = 100$ as a penalty for neighbouring labels being different.

We want to $s-t$ graph corresponding to this problem. The solution for this is not unique. The focus is often on constructing a graph with fewest edges, an approach suggested in Li book Section 10.4.2. However, you might prefer drawing a more intuitive graph despite a higher number of edges. This approach is sketched in Figure 5.3. Terminal edges (linking to source and sink) are used for the likelihood energy terms, while internal edges model the prior energy terms. Confirm that a cost of an $s-t$ cut in this graph equals to the posterior energy of the corresponding configuration.

To be able to compute the optimal configuration using the MATLAB GraphCut function, we need to create two matrices which contain edge weights to be passed to the function. The matrix containing terminal weights and the matrix containing weights between internal nodes for gender assignment example are shown in Figure 5.4. Python wrapper has slightly different manner of passing graph weights to the function, as explained in the package documentation.

Find the minimum $s-t$ cut by calling `[Scut, flow] = GraphCutMex(N, Et, Ei)`, with first inputs being the number of internal nodes in the graph, followed by the two weight matrices. What is given in the outputs? Which configuration is optimal? Change $\beta = 10$ and solve again. Which configuration is optimal now? Try also $\beta = 1000$.

Tasks

1. Download and test the provided software for graph cuts. For further help, we have provided small scripts which show a way of achieving this.

5.5 Exercise: Binary segmentation using MRF

Take a look at the bone image `V12_10X_x502.png`. The image is a slice from a CT scan of a mouse tibia. You can visually distinguish air (very dark), bone (very bright) and cartilage (dark). The task in this exercise is to segment the image in two segments: air and bone. Cartilage should

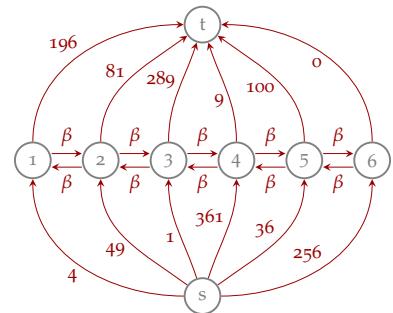


Figure 5.3: A sketch of an $s-t$ graph for a gender labeling problem.

$$E_{\text{terminal}} = \begin{bmatrix} 1 & 4 & 196 \\ 2 & 49 & 81 \\ 3 & 1 & 289 \\ 4 & 361 & 9 \\ 5 & 36 & 100 \\ 6 & 256 & 0 \end{bmatrix}$$

$$E_{\text{internal}} = \begin{bmatrix} 1 & 2 & \beta & \beta \\ 2 & 3 & \beta & \beta \\ 3 & 4 & \beta & \beta \\ 4 & 5 & \beta & \beta \\ 5 & 6 & \beta & \beta \end{bmatrix}$$

Figure 5.4: Representing an $s-t$ graph using two matrices, one containing weights of terminal edges and one matrix for internal edges.

be segmented together with air. The model we use is still the same as in the previous exercises, with the likelihood as the sum of squared distances, and the prior penalizing neighbouring labels being different. The bone image is of type `uint16`, and should be converted into double precision before any computation. You may also want to divide image intensities with $2^{16} - 1$, as this might simplify the weighting between the likelihood and the prior term.

Produce the histogram of the image to determine the mean intensities of the air and bone classes. Formulate the likelihood energy. Construct the matrices containing edge weights of the corresponding graph. Choose a parameter β and compute the optimal configuration using the `GraphCutMex` function. Adjust β to obtain a visually pleasing segmentation with reduced noise in air and bone. Produce a figure showing histogram of the entire image, and on top of that the intensity histograms of the air and bone classes, similar to how it was done in the modelling exercise.

Tasks

1. Segment bone image of circles into two classes. Check how changing β affects the segmentation.

5.6 Exercise: Multilabel segmentation usign MRF (optional)

Multilabel segmentation is obtained using an iterative α expansion algorithm. A MATLAB function `multilabel_MRF` implements α expansion. Read the help text of the function for explanation on input and output variables. Python users may try the `maxflow.fastmin` which is a part of `maxflow` package.

We will first verify the quality of the solution provided by the α expansion algorithm by returning to the segmentation of the synthetic image. How does the energy of the graph cut solution compare to the energies of the configurations found previously? Try changing β to see how it affects the result.

Use the α expansion algorithm to segment the bone image into air, cartilage and bone class. The challenge here is to distinguish between air and cartilage. You should aim at producing a visually pleasing result with cartilage as solid as possible (without noisy pixels segmented as air) and air as clean as possible (without noisy pixels segmented as cartilage). A good result can be obtained by tweaking two parameters: the mean value for the cartilage class and the smoothness weight β . As means for air and bone you can use the values estimated from the histogram, and you may assume the same standard deviation for all three classes (so there is no need for additional weighting of

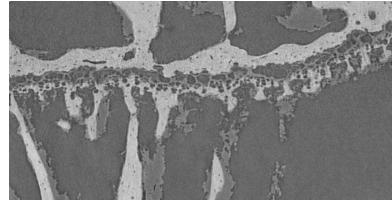


Figure 5.5: A bone image.

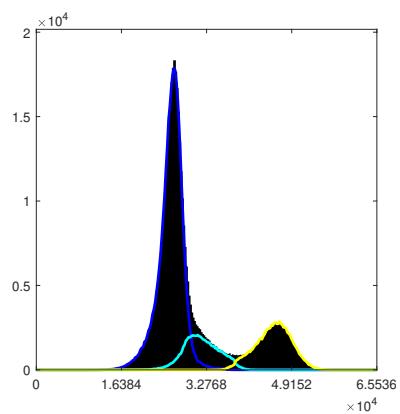
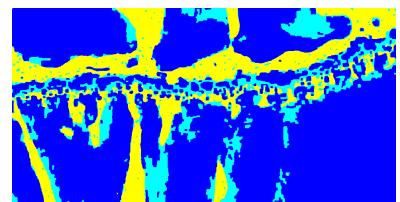
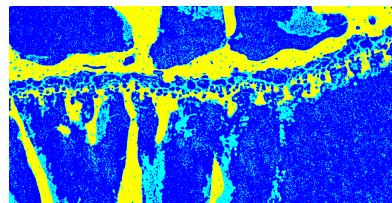


Figure 5.6: A segmentation of bone using maximum likelihood (top) and maximum posterior (middle). Histograms show intensity distributions for the three segmented classes.

the likelihood terms). When adjusting a mean value for cartilage, choose no smoothing ($\beta = 0$) and try to obtain a reasonable (but noisy) segmentation. Then increase β to remove the noise.

After you have tuned the parameters and obtained a nice segmentation, try segmenting the other bone image (`V8_10X_x502.png`). Can we use the same parameters? Why?

5.1 Information for 02506

Exam question related to this exercise

- Explain the use of MRF for image segmentation. Hints: Explain the concepts of likelihood and prior. Explain the difference between modelling and optimization. Explain how graph cuts are used to for MRF optimization.

6 Deformable models

TYPICALLY, IMAGE SEGMENTATION INVOLVES a combination of two terms: one dealing with the image data and the other describing a desirable segmentation. For example, we have used Markov random fields to impose smoothness on the segmentation. Using deformable models for image segmentation is another strategy which combines two contributions: the first originating from the image, and the second imposing smoothness.

The basic principle of deformable models is to perform image segmentation by evolving a curve in an image. The curve moves under the influence of *external forces*, which are computed from the image data, and *internal forces* which have to do with the curve itself. Deformable models are generally classified as either *parametric* (also called explicit) or *implicit* (in the context of image segmentation also called geometric), depending on the method used for representing the curve, see Figure 6.1. Despite this fundamental difference in curve representation, the underlying principles of both methods are the same¹.

In this exercise we use parametric curve representation, often called a *snake*², $C(s) = (x(s), y(s))$ where parameter $s \in [0, 1]$ is an arclength. In a discrete setting this reduces to a sequence of points, and parameter s becomes a discrete index $s = \{1, \dots, N\}$ indicating ordering of the points. We consider an image where the task is to separate the foreground from the background, and we use subscripts F and B for the corresponding image entities, such as for the image domain Ω consisting of Ω_F and Ω_B . At the same time, a curve C divides the image into inside and outside region, and for those regions we use subscripts in and out.

Deformable models are guided by the segmentation energy E , which should be defined such that the desired segmentation has a minimal energy. A segmentation is obtained by iteratively moving the curve to minimize the energy, and the most challenging part of the approach is deriving energy-minimizing curve deformation forces $F = -\nabla E$. To allow deformation, the curve is made dynamic (time-dependable), and

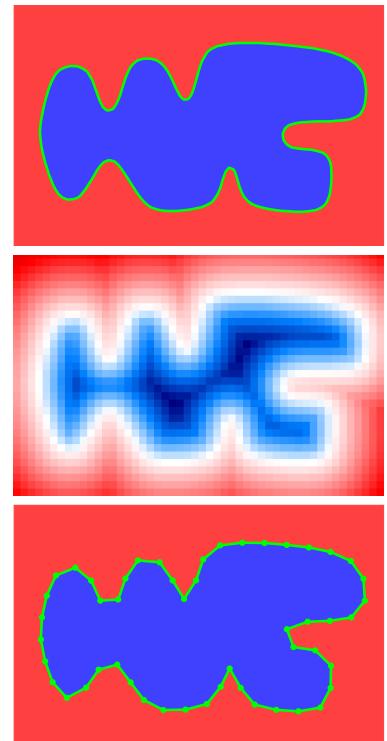


Figure 6.1: A curve (top) and its two discrete representations: implicit (middle) and parametric (bottom).

¹ Chenyang Xu, Anthony Yezzi Jr, and Jerry L Prince. On the relationship between parametric and geometric active contours. In *The Asilomar Conference on Signals, Systems, and Computers*, volume 1, pages 483–489. IEEE, 2000b.

² Michael Kass, Andrew Witkin, and Demetri Terzopoulos. Snakes: Active contour models. *International Journal of Computer Vision*, 1(4):321–331, 1988.

its change in time, often denoted *evolution*, is given by

$$\frac{\partial C}{\partial t} = F(C).$$

We use E_{ext} to denote external energy, which is a contribution to the segmentation energy determined by image data. We use E_{int} for internal energy, which has to do with the curve itself. Correspondingly, deformation forces on the curve are divided into external and internal F_{ext} and F_{int} .

This exercise is inspired by the Chan-Vese algorithm³, a deformable model for image segmentation which minimizes a piecewise-constant Mumford-Shah functional. Chan-Vese uses an implicit (level-set) curve representation and a two-step optimization. We will use the solution of the Chan-Vese approach, but will combine it with a parametric curve representation. For this reason our model uses an external energy similar to Chan-Vese algorithm, and an internal energy similar to snakes. The detailed explanation on how external forces are derived from external energy can be found in the Chan-Vese article. How internal forces are derived is explained in Chapter 3 of the Handbook of Medical Imaging, Image Segmentation Using Deformable Models⁴, subsection 3.2.1 and 3.2.4. Recall that you already implemented curve smoothing as one of the introductory exercises during the first week of the course.

6.1 External energy, Chan-Vese

An external energy closely related to the two-phase piecewise constant Mumford-Shah model is

$$E_{\text{ext}} = \int_{\Omega_{\text{in}}} (I - m_{\text{in}})^2 d\omega + \int_{\Omega_{\text{out}}} (I - m_{\text{out}})^2 d\omega$$

where I is an image intensity as a function of the pixel position, while m_{in} and m_{out} are mean intensities of the inside and the outside region. This energy seeks the best (in a squared-error sense) piecewise constant approximation of I . An evolution that will deform a curve toward an energy minimum is derived as

$$F_{\text{ext}} = (m_{\text{in}} - m_{\text{out}}) (2I - m_{\text{in}} - m_{\text{out}}) N. \quad (6.1)$$

where N denotes an outward unit normal.

In other words, curve deforms in the normal direction, and for every point on the curve we only need to evaluate the (signed) length of the displacement. We will denote the scalar components of the force as $f_{\text{ext}} = (m_{\text{in}} - m_{\text{out}}) (2I - m_{\text{in}} - m_{\text{out}})$. Note that this can be written as $f_{\text{ext}} = (m_{\text{in}} - m_{\text{out}}) \left(I - \frac{1}{2}(m_{\text{in}} + m_{\text{out}}) \right)$, i.e. the signed length of displacement is proportional to the difference between the image intensities under the curve and the mean of m_{in} and m_{out} .

³ Tony F Chan and Luminita A Vese. Active contours without edges. *IEEE Transactions on image processing*, 10(2):266–277, 2001

⁴ Chenyang Xu, Dzung L Pham, and Jerry L Prince. Image segmentation using deformable models. *Handbook of medical imaging*, 2:129–174, 2000a

6.2 Internal forces, snakes

The internal energy is determined solely by the shape of the curve. In the classical snakes formulation internal forces discourage stretching and bending of the curve

$$E_{\text{int}} = \frac{1}{2} \int \alpha \left| \frac{\partial C}{\partial s} \right|^2 + \beta \left| \frac{\partial^2 C}{\partial s^2} \right|^2 ds,$$

with weights α and β controlling the elasticity (first-order derivative) and the rigidity (second-order derivative) term. Corresponding deformation forces are

$$F_{\text{int}} = \frac{\partial}{\partial s} \left(\alpha \frac{\partial C}{\partial s} \right) + \frac{\partial^2}{\partial s^2} \left(\beta \frac{\partial^2 C}{\partial s^2} \right). \quad (6.2)$$

Those regulatory forces are the key to success of deformable models, as they provide robustness to noise.

Since our snake is discrete, the derivatives should be approximated by finite differences. The regularization now corresponds to filtering (smoothing) the curve with filters for the first and the second derivative, (i.e. a filter $[1 -2 1]$ and a filter $[-1 4 -6 4 -1]$). Those contributions, weighted by parameters α and β are now used to regularize the curve. In efficient implementation this is done by a matrix multiplication, and for better stability we use a backward Euler scheme. For slightly more detail, you can revise the introductory exercise on curve smoothing 1.1.3.

6.3 Final model

For a snake consisting of n points and represented using an $n \times 2$ matrix C , a final discrete update step is, adapted from Handbook of Medical Imaging, Eq. (3.22),

$$C^t = B_{\text{int}} \left(C^{t-1} + \tau \text{diag}(f_{\text{ext}}) N^{t-1} \right). \quad (6.3)$$

In this expression τ is the time step for displacement, while B_{int} is the $n \times n$ matrix used for regularizing the curve and taking the role of the internal forces. Curve normals are represented as $n \times 2$ matrix N , and pointwise displacement is obtained by multiplying N with a $n \times n$ diagonal matrix containing the displacement lengths (this is in principle a row-wise multiplication).

6.4 Exercise: Segmentation and tracking

We will use a deformable model to segment and track a simple organism in a sequence of images. You are provided with two image sequences:

crawling amoeba⁵ and water bear⁶. The same code can be used for both sequences, with only a minor adjustment in a pre-processing step.

Tasks Steps for solving the problem are listed below, with the hints for MATLAB and python users.

1. Read in and inspect the movie data. In MATLAB you may use `VideoReader`. You may save the image sequence as a multi-dimensional array, or as a movie object using `im2frame` conversion. In python you may use function `get_reader` from `imageio` package.
2. Process movie frames. For our segmentation method to work, movie frames need to be transformed in grayscale images with a significant difference in intensities of the foreground and a background. For the movie showing the crawling amoeba (which is white on a dark background), it is enough to convert movie frames to grayscale. Transforming intensities to doubles between 0 and 1 is advisable, as it might prevent issues in subsequent processing. For the movie of the echinicsus, we want to utilize the fact that foreground is yellow while background is blue. A example of suitable transformation is $(2b - (r + g) + 2)/4$, with r, g, b being color channels (with values between 0 and 1).
3. Choose a starting frame and initialize a snake so that it roughly delineates the foreground object. You may define a circular snake with points $(x_0 + r \cos \alpha, y_0 + r \sin \alpha)$, where (x_0, y_0) is a circle center, r is a radius and angular parameter α takes n values from $[0, 2\pi]$. See Figure 6.2 for example, but use approximately 100 points along the curve.
4. Compute mean intensities inside and outside the snake. In MATLAB you can use `poly2mask` function. In python use `polygon2mask` from package `skimage.draw` starting with version 0.16.
5. Compute the magnitude of the snake displacement given by Eq. (6.1). That is, for each snake point, compute the scalar value giving the (signed) length of the deformation in the normal direction. This depends on image data under the snake and estimated mean intensities, as shown in Figure 6.3. A simple approach evaluates the image intensities under the snake by rounding the coordinates of the snake points. A more advanced approach involves interpolating the image at the positions of snake points for example using bilinear interpolation, which is in MATLAB implemented in function `interp2`, and is in python available under the same name in `scipy.interpolate` package.

⁵ The video of crawling amoeba is from Essential Cell Biology, 3rd Edition Alberts, Bray, Hopkin, Johnson, Lewis, Raff, Roberts, & Walter, <https://www.dnatube.com/video/4163/Crawling-Amoeba>

⁶ The video of water bear is from Olympus microscopy resources, <https://www.olympus-lifescience.com/ru/microscope-resource/moviegallery/pondscum/tardigrada/echiniscus>

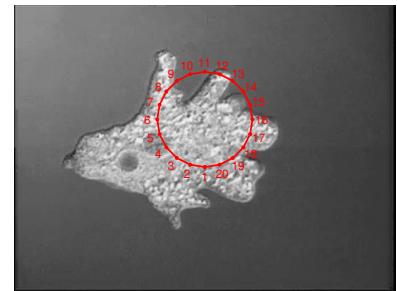


Figure 6.2: The first frame of a crawling amoeba and a circular a 20-point snake.

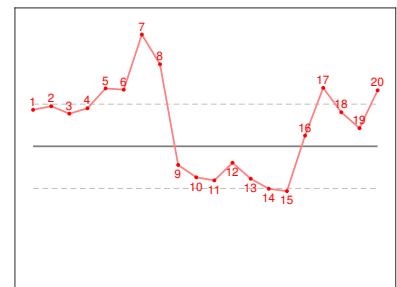


Figure 6.3: Red curve shows image intensities along the snake in Figure 6.2. Dashed gray lines indicate m_{in} and m_{out} , while gray line indicates the mean of m_{in} and m_{out} . Signed length of the curve displacement is given by the difference between the red curve and the gray line.

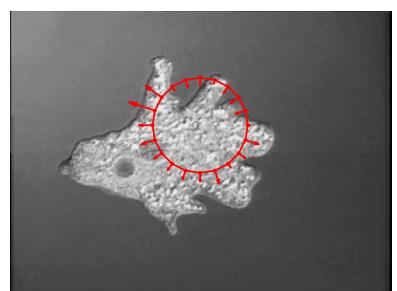


Figure 6.4: Force on the curve indicated by arrows. Displacement is in the normal direction and the length of the displacement is given by the values shown in Figure 6.3.

6. Write a function which takes snake points \mathbf{C} as an input and returns snake normals \mathbf{N} . A normal to point c_i can be approximated by a unit vector orthogonal to $c_{i+1} - c_{i-1}$. (Alternatively, and slightly better, you may average the normals of two line segments meeting at c_i .) Displace the snake. Estimate a reasonable value for the size of the update step by visualizing the displacement. You should later fine-tune this value so that the segmentation runs sufficiently fast, but without introducing exaggerated oscillations. This step corresponds to computing the expression in the parentheses in the Eq. (6.3).
7. Write a function which given α , β and n constructs a regularization matrix \mathbf{B}_{int} . Your code from introductory exercise could be used. Apply regularization to a snake. Estimate a reasonable values for the regularization parameters α and β by visualizing the effect of regularization. You should later fine-tune these values to obtain a segmentation with the boundary which is both smooth and sufficiently detailed. This step corresponds to matrix multiplication on the right hand side of the Eq. (6.3).
8. The quality of the curve representation may deteriorate during evolution, especially if you use a large time step θ and/or weak regularization, i.e. small α and β . To allow faster evolution without curve deterioration, you may choose to apply a number of substeps (implemented as subfunctions) which ensure the quality of the snake:
 - Distribute points equidistantly along the snake. This can be obtained using 1D interpolation (function `interp1`).
 - Constrain snake to image domain.
 - Apply heuristics for removing crossings from the snake. For example, if you detect self-intersection, identify two curve segments separated by the intersection and reverse the ordering of the smallest segment.
- We provide a few functions for improving the quality of the snake, in MATLAB and in python.
9. Repeat steps 4–8 until a desirable segmentation is achieved. Note that the regularization matrix only depends on regularization parameters and a number of snake points. This is constant when the size of the snake and the regularization are fixed, which is a typical case. It is therefore sufficient to precompute \mathbf{B}_{int} prior to looping. Figure 6.5 shows our 20-point snake during evolution.
10. Read in the next frame of the movie, and use the results of the previous frame as an initialization. Evolve the curve a few times by repeating steps 4–8.

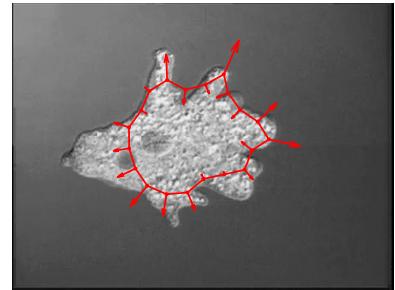


Figure 6.5: The curve and the forces after 20 iterations.

11. Process additional frames of the image sequence.

6.1 Information for 02506

Exam question related to this exercise

- Explain the use of deformable models for image segmentation.
Hints: Explain the piecewise-constant Mumford-Shah functional and how it is minimized using Chan-Vese algorithm? Explain the difference between the external and the internal forces. Explain curve representations used for deformable models, and their advantages/ disadvantages.

7 Using geometric priors for volumetric segmentation

IN THIS CHAPTER the practical aspects of volumetric segmentation using geometric priors. This is demonstrated on a problem of segmenting myelinated axons from the volumetric data containing scans of the human posterior interosseous nerve.

There are various strategies for solving this problem. Here we demonstrate the use of two approaches we already presented: Markov random fields introduced in Chapter 5 and deformable models introduced in Chapter 6. Furthermore, an approach from Chapter 13, which is a suggestion for mini-project, can be used for solving nerve segmentation problem.

You will be working on the small part of the large data set. For better understanding of the goals of image analysis, we provide background information on the large study involving a complete data set.

7.1 X-ray tomography of human peripheral nerves

Nerve disorders caused by trauma or disease can have serious consequences for the affected people. One aspect of understanding nerve disorders involves a knowledge of the structure of peripheral nerves and their subcomponents. This is typically obtained through microscopy.

Imaging using conventional light and electron microscopical techniques only allows a two-dimensional visualization of tissues such as peripheral nerves. With recent advances in synchrotron imaging techniques, we can now also obtain detailed three-dimensional images of tissue. This in turn allows extraction of 3D morphological information.

For a larger study¹, biopsies of the posterior interosseous nerve at wrist levels were taken from otherwise healthy subjects and from subjects with type 1 and 2 diabetes. The aim of the study was to determine whether diabetes influences the radius, trajectory and organization of myelinated axons in human peripheral nerves.

Biopsies were stained in osmium (a heavy metal used for staining

¹ Dahlin L B, Rix K R, Dahl V A, Dahl A B, Jensen J N, Cloetens P, Pacureanu A, Mohseni S, Thomsen N O B, and Bech M. X-ray phase contrast zoom tomography to visualize human diabetic peripheral nerves. *Nature Methods* (in submission)

lipids) which provides contrast to the image, and embedded in Epon (epoxy resin) for stability. The samples were then imaged using X-ray phase contrast zoom tomography at the European Synchrotron Radiation Facility (ESRF, Grenoble, France) with an isotropic voxel size of 130nm. In the obtained volumetric data, the nerve fibers are aligned with the z direction, and the stained myelin sheaths around axons (see Figure 7.1 for a schematic drawing of a nerve) appear circular in the x-y slices through the volume, as shown in Figure 7.2.

Complete data-set contains more than 10 samples, each resulting in a volume of a size $2048 \times 2048 \times 2048$ voxels. For the exercise, we extracted a small region from one volume, as indicated in Figure 7.2. Furthermore, extracted volume has been downsized by a factor 2, which yields a volume of size $350 \times 350 \times 1024$. The extracted volume is saved as a stacked tiff image `nerves_part.tiff`.

7.2 Segmentation of myelinated nerves

A good contrast between stained myelin sheaths and the background makes it possible to clearly distinguish individual nerve cells in the volume. For this reason, a reasonable segmentation strategy would utilize dark appearance of the myelin. Segmentation may be improved by incorporating a prior knowledge about the directionality of the nerves.

Furthermore, circular appearance of myelin sheaths allows a segmentation of a single nerve cell by aligning a closed curve with the periphery of the myelin. For this, the circle can be manually initialized around the nerve cell, and automatically moved to the boundary of the myelin. For segmenting a whole nerve, the curves are automatically propagated trough the volume, such that the surface moves only slightly between the slices. In every slice the surface is to be attracted to the boundary of the myelin layer.

Try segmenting nerves using Markov random fields and deformable models. In Figure 7.3, 7.4, 7.5 and 7.6 we show results of image analysis performed for the original study. However, those results are obtained on a full-resolution volumes, and using a combination of deformable models(Chapter 6) and layered surfaces(Chapter 13). Your results might therefore be of poorer quality.

For this open assignment, we provide some tips, but you are free to investigate other approaches.

- For MRF segmentation you might consider further downsizing the data or using only a subset of slices.
- When using MRF segmentation you can process the volume slice-by-slice. This corresponds to a situation where MRF-modelled smooth-

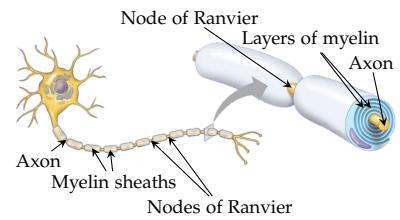


Figure 7.1: A schematic drawing of a nerve showing an axon, myelin sheaths and nodes of Ranvier.

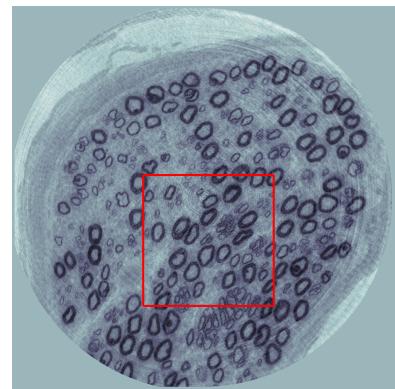


Figure 7.2: One slice from the volume showing peripheral nerves, and a region which was extracted for the exercise.

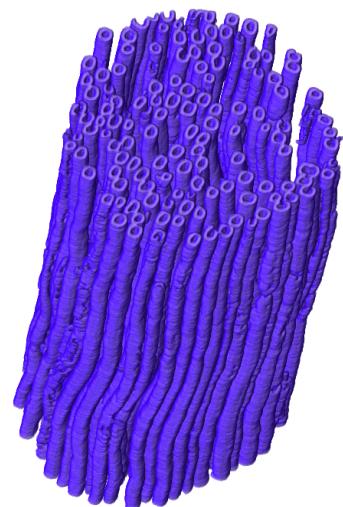


Figure 7.3: 3D visualization of segmented myelinated nerves.

ness prior has a parameter β for two neighbouring pixels in x and y direction, while the change of labels for neighbours in z direction is not penalized. However, note that nerves are elongated in the direction orthogonal to image slices. For this reason, it would be more appropriate to set MRF-modelled smoothness especially high along the z direction, and this calls for a full 3D implementation of the MRF segmentation.

- When using deformable models note that assumption of Chan-Vese about a object of different intensity than the background does not apply for nerves. This is because a nerve consists of a dark myelin and a bright axon. Instead of allowing for the automatic estimation of the parameters m_{in} and m_{out} by averaging, it might be better to fix those parameters using the values estimated from the images. Alternatively, values m_{in} and m_{out} may be estimated from a thin band inside and outside the curve.
- The robustness of the deformable models might be improved by moving the curve towards the point where the change of intensity in the normal direction is high. This can be implemented by unwrapping the image (similar to exercise 1.1.5) following the curve normals. The gradient in normal direction can then be computed for the unwrapped image, and curve moved to the point where gradient is high.
- A node of Ranvier (see Figure 7.1 for schematic drawing of a nerve) can be seen on a few of the nerves in the volume to be analysed, as shown in Figure 7.6. Nodes are of a high interest for understanding nerve disorders. However, those might be challenging to capture due to the lack of myelin.
- Visualizing results in 3D usually provides a useful information on the segmentation results. Visualization options provided by MATLAB and python are good. Still, for large datasets, and advanced visualization you may want to use a specialized software, and we suggest trying ParaView. A few notes on 3D visualization using ParaView be found in [ParaView notes](#).

Tasks

1. Consider following microstructural measurements which can be extracted from volumetric data:

Nerve density count: Number of axons per area of nerve-fibre cross-section. Measured in number per area.

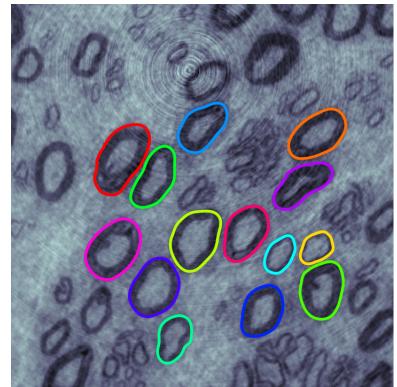


Figure 7.4: Axons segmented using deformable curves visualized on a single slice.

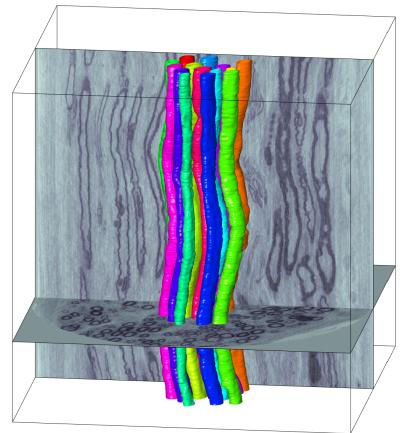


Figure 7.5: 3D visualization of axons segmented using deformable curves.

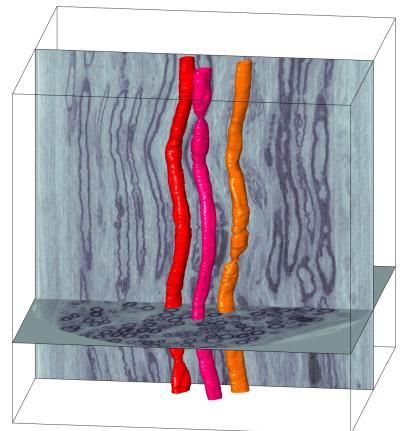


Figure 7.6: Three of the axons with visible node of Ranvier.

Myelin density: A fraction of nerve cross-section corresponding to myelin. Expressed as dimensionless fraction (a number between 0 and 1), or a percentage.

Average nerve area: Average area of nerves, broken down into average axon area, and average myelin area. This measure is very related to average nerve radius.

Average nerve radius: Average radius of nerves, broken down into average axon radius, and average myelin thickness. This measure is very related to average nerve area.

2. Perform a binary segmentation of the data using MRF. Which microstructural measurement can you extract from your MRF segmentation?
3. Perform a volumetric segmentation of few nerves using deformable models. Which microstructural measurements can you extract from your segmentation?

7.1 Information for 02506

A deliverable. We will have a small deliverable related to this exercise, in form of a one-page presentation of your work. You may chose to present MRF results, deformable surfaces, or both. We suggest that you choose an illustration of your best result, and explain how it was obtained. You are welcome to briefly describe your segmentation strategy and suggest possible improvements. Your one-pagers will be collected in a single document and shared with other students. You should therefore:

- Write your name on the one-pager.
- Upload the one-pager to DTU CampusNet as a pdf document in A4 format.
- Refrain from including content which you do not want to be shared.

The deliverable is not going to be used for assessment, and will have no influence on the exam. If you are prevented in producing a deliverable, or do not want it to be shared with other students, just write a comment in CampusNet.

Exam questions related to this exercise

- Explain the problem of nerve segmentation, the strategy you chose for segmenting the nerves and the obtained results.

Part III

Image analysis with neural networks

NEURAL NETWORKS are very useful for a range of image analysis tasks including segmentation, detection, classification, etc. Neural networks are often easy to adapt to a specific problem and they allow approximating an unknown function f^* that based on some input \mathbf{x} can predict the output \mathbf{y} even without *a priori* knowing the relation between \mathbf{x} and \mathbf{y} . This is done by learning a set of parameters θ from a training set, i.e. of corresponding input values \mathbf{X} and predictions \mathbf{Y} . In image analysis problems the input will typically be an image or a part of an image, and the output is a scalar vector or an image.

A range of high-performance libraries for neural networks exists, that are very well suited for solving a number of problems also in image analysis. The aim here is however to understand the basic elements of neural networks and get experience with their functionality. This will be done by implementing a feed forward neural network, a Multilayer Perceptron (MLP). The first task is to separate simple point sets. This is not an image analysis task, but it is chosen as a simple and easy to visualize task that can help verifying that the implementation is correct. Furthermore, it will allow some experience with various model parameters. This implementation will later be applied to image classification and image segmentation.

Information for 02506

The third part of the course, image analysis with neural networks, is covered in three weeks of the course. Counting from the start of the semester, these are weeks 8, 9, and 10.

- Week 8: Feed forward neural network (multilayer perceptron (MLP)) will be introduced. The reading for this exercise is this chapter. You can find extensive background material in Chapter 6 in the book Deep Learning by Goodfellow et al., which you can find online <https://www.deeplearningbook.org/>, which gives more details and examples of MLP. During the exercise you will implement a feed forward neural network in MATLAB or Python. This time you will only work with simple 2D point sets to verify your implementation, but in the following exercise your network will be used on image data.
- Week 9: Regularization and optimization for neural networks. The background material is found in Chapters 7.4, 8.3, 8.4 and 8.5 in the book Deep Learning by Goodfellow et al. (<https://www.deeplearningbook.org/>). Deep learning is highly dependent on data for training the model param-

eters. Here data augmentation and model regularization is important. In the exercise you will implement some of these techniques to optimize the performance of your network. The exercise will be a competition on performance of your network, and there is a prize for the group with the best performing network.

- Week 10: Convolutional neural networks will be covered. The background reading material is found in Chapter 9 in the book Deep Learning by Goodfellow et al. (<https://www.deeplearningbook.org/>). In the exercise you will experiment with a convolutional neural network and compare its performance to the feed forward network from the two previous exercises.

8 Feed forward neural network

A **NEURAL NETWORK** is modeled as a directed graph as shown in Figure 8.1. The input layer is shown on the left, hidden layers are in the middle, and the output layer is to the right. This exercise is based on the description in the Deep Learning book¹ but also Chapter 5 in the book Pattern Recognition and Machine Learning² gives a good introduction to neural networks.

8.1 Concept of neural network

The fundamentals for the deep learning method are explained in the Deep Learning book (Goodfellow et al.), but here we will give a brief introduction to a simple feed forward network. You will later implement a more general version of a feed forward network, but first we will focus on the basic elements.

Conceptually we want to make a function f that takes a vector \mathbf{x} as input and predicts a vector \mathbf{y}

$$f(\mathbf{x}) = \mathbf{y}.$$

In image analysis, the input \mathbf{x} will be an image concatenated into a vector (just reshaping the image into a vector). The output depends on the problem to be solved. If we solve a classification problem, it will typically be a vector of class probabilities, e.g. if there are k classes, \mathbf{y} will be a k -dimensional vector, where each element in the vector has the value of the probability for belonging to one of the k classes. If we are doing image segmentation into L labels, then the output will be an $D \times L$ matrix where each element is a probability of the i -th pixel belonging to the j -th label where $i = 1, \dots, D$ and $j = 1, \dots, L$.

There are some steps in deep learning that requires design choices, which are referred to as hyper-parameters of the model (opposed to the edge weights that are the model parameters). This includes choices such as number of layers, number of elements in these layers and decisions regarding regularization parameters.

¹ Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT press Cambridge, 2016

² CM Bishop. *Pattern recognition and machine learning*, 2006

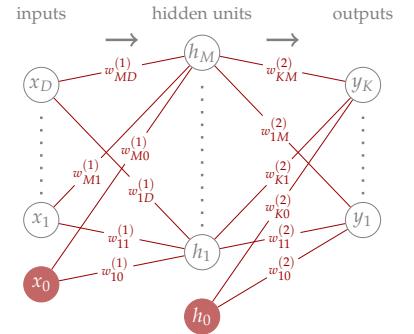


Figure 8.1: Example of a neural network with an input layer, one hidden layer, and an output layer. This is termed a one layer network, since it has one hidden layer. Typically, there will be many hidden layers.

We will start by explaining the *forward model* of the feed forward neural network. Through the forward pass we get the model prediction, which solves our image analysis problem. But in order to get a model, we must tune the parameters of the model. This is done by changing the edge weights such that the model can correctly predict the values of a training set. To obtain the model parameters we use the *backpropagation algorithm*.

8.2 Forward model

For simplicity, we start by describing the network shown in Figure 8.2. This network will solve a classification problem. The network takes a two dimensional input vector and outputs a probability distribution over two classes. The network contains an input layer of three nodes (neurons), where two are the input variable (x_1, x_2) (independent variables) and $x_0 = 1$ is the bias node. The hidden layer contains four nodes including three nodes connected to the input layer (h_1, h_2, h_3) and the bias node $h_0 = 1$, and the output layer contains the two predicted values (y_1, y_2) (dependent variables). The weights connecting the nodes are termed $w_{ij}^{(l)}$ connecting node j from layer $l - 1$ to node i in layer l .

The values of the nodes in the hidden layers are computed by first computing a weighted linear combination z_i of the node values and the edge weights followed by an activation function, and we use the max function $a(z_i) = \max(z_i, 0)$, which in deep learning is called the rectified linear units function (ReLU) to obtain h_i . We have

$$z_i = \sum_{d=0}^D w_{id}^{(1)} x_d , \quad (8.1)$$

$$h_i = a(z_i) = \max\{0, z_i\} , \quad (8.2)$$

$$\hat{y}_j = \sum_{m=0}^M w_{jm}^{(2)} h_m . \quad (8.3)$$

Since the output of the network should be used for classification, we want to interpret the output values as probabilities, i.e. each element in the output y_j can be seen as the probability of the observation belonging to the j -th class. The reason for obtaining probabilities as output is related to the use of backpropagation for parameter optimization. Therefore, we must transform the values of \hat{y}_j , and this is done through normalization. For normalization, we use the softmax function and get the values

$$y_j = \frac{e^{\hat{y}_j}}{\sum_{k=1}^K e^{\hat{y}_k}} . \quad (8.4)$$

This makes the output values sum to one. For classification, the output class j^* with the highest value is chosen. The values of y can however

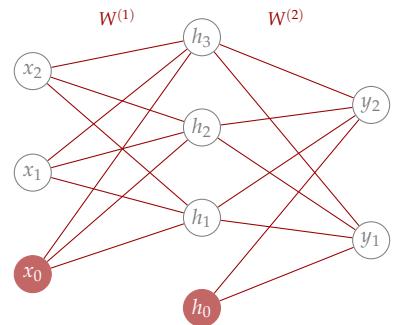


Figure 8.2: Simple three layer neural network.

also be used directly for visualization. This is done by computing the output values in a regular grid and displaying their value as an image.

8.3 Backpropagation

We now want to adjust the weights of the network to minimize the loss over a training set of inputs and the associated target values. This is done through the backpropagation algorithm which uses an iterative optimization method called stochastic gradient descent. Just as in gradient descent the loss is minimized by taking steps proportional to the negative gradient. We compute the gradient from the loss, which is derived from the predicted value y and the target t , which is given in the training set. We use the cross entropy loss function

$$L = - \sum_{k=1}^K t_k \log y_k , \quad (8.5)$$

where t_k is the target value of the prediction where

$$t_k = \begin{cases} 1 & \text{if class label is } k \\ 0 & \text{otherwise} \end{cases} . \quad (8.6)$$

The cross entropy loss gives zero if the prediction and target are the same and otherwise a value larger than zero.

When computing the gradient we can compare the predictions and targets for all data in our training set. However, we do not consider the loss function for all inputs and targets at once. Instead, we consider one input (or a smaller sample of inputs) in a random order. The derivation below computes update for one iteration of the stochastic gradient descent.

We will evaluate derivatives for a certain (fixed) input in order to determine how change in each $w_{ij}^{(l)}$ affects L , and then we will use an update

$$w_{ij}^{(l)\text{new}} = w_{ij}^{(l)} - \eta \frac{\partial L}{\partial w_{ij}^{(l)}} ,$$

where η is a user-chosen learning rate.

Change in each $w_{ij}^{(l)}$ contributes to the change in L only through $z_i^{(l)}$ and using the chain rule the derivative may be separated into two elements

$$\frac{\partial L}{\partial w_{ij}^{(l)}} = \frac{\partial L}{\partial z_i^{(l)}} \frac{\partial z_i^{(l)}}{\partial w_{ij}^{(l)}} .$$

Since $z_i^{(l)}$ is a linear function of $w_{ij}^{(l)}$ the second (easy) derivative evaluates to $h_j^{(l-1)}$ (or, in the case of the first layer, input values x_j). This is

valid regardless of the non-linear activation. The first (more difficult) derivative needs to be evaluated for each $z_i^{(l)}$ and we denote those values by $\delta_i^{(l)}$, such that we have

$$\frac{\partial L}{\partial w_{ij}^{(l)}} = \delta_i^{(l)} h_j^{(l-1)}. \quad (8.7)$$

Notice the simplicity: the update for weight $w_{ij}^{(l)}$ is a product of two values, the first value depends only on the *to*-node and the second value depends only on the *from*-node.

We still need to evaluate $\delta_i^{(l)}$, i.e. establish how a change of $z_i^{(l)}$ affects L . This depends only on what happens in the layers further down the pipeline, and on the choice of the non-linear activation used on $z_i^{(l)}$. We distinguish between the last layer (where we use the softmax function) and the internal layers.

For the last layer we express L as a function of $z_k^{(l^*)}$

$$\begin{aligned} L &= -\sum_k t_k \ln \frac{\exp z_k^{(l^*)}}{\sum_j \exp z_j^{(l^*)}} = \text{using the properties of ln and the distributive rule} \\ &= -\sum_k t_k z_k^{(l^*)} + \left[\sum_k t_k \ln \sum_j \exp z_j^{(l^*)} \right]_{=1} \text{equal for all } k. \end{aligned}$$

The derivative of L with respect to $z_i^{(l^*)}$ is therefore

$$\delta_i^{(l^*)} = -t_i + \frac{1}{\sum_j \exp z_j^{(l^*)}} \exp z_i^{(l^*)} = y_i - t_i. \quad (8.8)$$

Now consider the internal layers. The change in $z_i^{(l)}$ may change all z_k^{l+1} , and any of these changes may affect L . The chain rule gives

$$\frac{\partial L}{\partial z_i^{(l)}} = \sum_k \frac{\frac{\partial L}{\partial z_k^{(l+1)}}}{\frac{\partial z_k^{(l+1)}}{\partial z_i^{(l)}}} \text{we have} \quad \text{we need}$$

The first set of derivatives are $\delta_k^{(l+1)}$ for the layer further down the pipeline. We already evaluated those for the last layer in (8.8), and this is why we compute the update backwards through the network. The only remaining is to determine how the change of $z_i^{(l)}$ affects $z_k^{(l+1)}$. From definition (8.1) we see that $z_k^{(l+1)}$ is a linear function of $a(z_i^{(l)})$ which gives

$$\frac{\partial z_k^{(l+1)}}{\partial z_i^{(l)}} = w_{ki}^{(l+1)} a'(z_i^{(l)}),$$

where a' denotes the derivative of the activation function, which for ReLU function takes a value zero for arguments smaller than zero, and

one otherwise. This is easy to determine by assessing whether $h_i^{(l)}$ is zero or larger. The final expression for internal layers is

$$\delta_i^{(l)} = a'(z_i^{(l)}) \sum_k w_{ki}^{(l+1)} \delta_k^{(l+1)}. \quad (8.9)$$

8.4 Implementation

You are now ready to implement your neural network.

Data preprocessing The recommended preprocessing is to center the data to have mean of zero. For features of different scale, it is advised to normalize the scale along each feature.

Initialization Weights should be initialized with small numbers. Initializing with zeros is not a good idea, as it introduces no asymmetry between neurons. The current recommendation in the case of neural networks with ReLU neurons is to draw the weight from the Gaussian distribution with standard deviation of $\sqrt{\frac{2}{n}}$, where n is the number of inputs to the neuron.

Batch optimization A variant of stochastic gradient descend splits the training data into smaller subsets (minibatches) and updates weight according to the average of the gradients for the minibatch. One cycle through the entire training dataset is called a training epoch.

Matrix multiplications The layered structure makes it efficient to evaluate and train neural networks using matrix vector operations. The forward propagation trough layers is evaluated using a product

$$\mathbf{h}^{(l)} = a(\mathbf{W}^{(l)} \mathbf{h}^{(l-1)}), \quad (8.10)$$

where vectors $\mathbf{h}^{(l)}$ and $\mathbf{h}^{(l-1)}$ contain values $h_i^{(l)}$ and $h_i^{(l-1)}$ (without a bias), $\mathbf{W}^{(l)}$ is a matrix containing weights and a is activation. Vectorized expression (8.10) is also valid when passing multiple inputs trough the network: a $2 \times m$ input results with a $2 \times m$ output. Likewise, during training, values $\delta_i^{(l)}$ can be represented as vectors, such that computation (8.9) utilizes a matrix multiplication, while (8.9) becomes an outer product of two vectors. When working with batches (8.9) is a matrix-matrix multiplication.

8.4.1 Setting up the problem

You should implement the network shown in Figure 8.2 using the description given above. This network contains a single hidden layer and takes a two dimensional input and classifies the input to a two dimensional output. Before you get there you should have some test data for

running your method. We have provided a function for MATLAB called `make_data.m` and a script for running it called `data_example.m` and in Python it is all in the file `make_data.py`. This will create test point-sets similar to the ones shown in Figure 8.3. In this experiment you will use the same points for training and testing. This first experiment is to ensure that you build the neural network in a correct way, and later you will test the performance of your method on an independent validation dataset.

8.4.2 Simple three layer network

You should start with a three layer network with a single hidden layer with five neurons, i.e. four neurons that are connected to the input layer and a bias variable. It is a good idea to start making a hand drawing of the network you should implement will the nodes, edges and notation for the parts of the network.

Suggested procedure

1. Start by implementing the forward propagation step with random weights. Make sure that it works as expected. It is important to verify the obtained results which is easy when working with 2D data and can be done by sampling points on a regular grid. This can be displayed as an image where each pixel takes the label number.
2. Make a script that computes the classification results on a regular grid. You can look at the results from your forward propagation with random numbers.
3. Implement the backpropagation algorithm. Note that both the forward and backward propagations can be implemented efficiently using matrix operations. Display the loss as you iterate to ensure that your algorithm is converging.
4. Test your implementation on the generated data. What should the learning rate be? How is performance affected by noise? Does it converge to the same result each time you run it?

8.4.3 Variable number of layers and hidden units

In the exercise above you have obtained a neural network with a fixed architecture, i.e. the number of neurons and hidden layer. The architecture is however central in modeling with neural networks, and therefore you should make the number of layers and the number of neurons in each of the hidden layers a part of your input choice. You will also be needing this flexibility in the later exercises for classifying the MNIST handwritten digits.

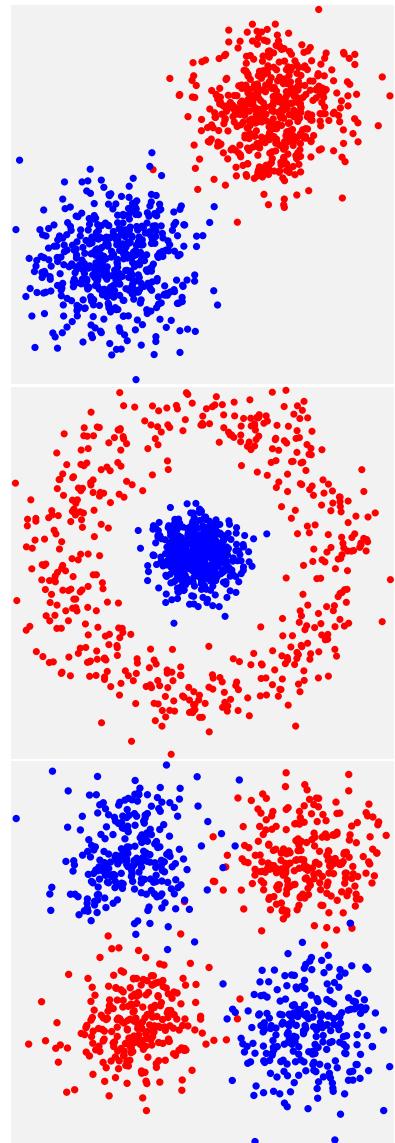


Figure 8.3: Scatter plots of point sets to test neural network

Suggested procedure

1. Implement a neural network keeping the single hidden layer and with variable number of hidden units.
2. Implement a neural network with variable number of hidden layers but with variable number of hidden units. Now you can play around with your model and from here it is easy to modify it to include other elements like other activation functions.
3. Again test your implementation on the given data.

The expected output is shown in Figure 8.4. The coloring of the pixels in the background is obtained by running all the coordinates of the pixels in an through the neural network, and coloring the result in dark or bright gray depending on the classification result.

8.1 Information for 02506

Exam question related to this exercise

- Explain the components of a feed forward neural network including input, hidden units, output, etc. Explain the back-propagation algorithm and show using a simple example how it is computed. This could include the loss function and how weights are updated using partial derivatives.

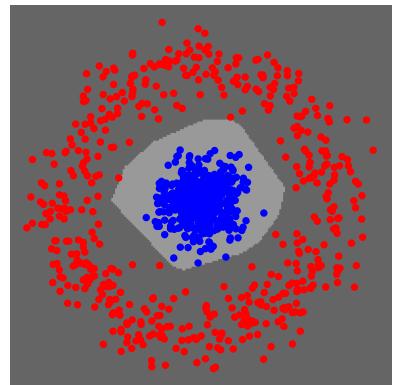


Figure 8.4: Result of training on the input data shown in colored points, and the test result is shown in the pixel colors in the background.

9 MNIST classification

IN THIS EXERCISE you should build a neural network for classifying the MNIST image data. Performance will be measured in number of misclassified images, and the goal is to obtain the result with the least number of misclassified images.

The MNIST dataset contains images of handwritten digits of 28×28 pixels as shown in Figure 9.1. Ground truth class labels are given together with the images. The ground truth is 10 dimensional vectors with 1 in the dimension representing the class containing the digit and zeros elsewhere. MNIST contains 60000 images for training your network. If you use all 60000 images for training your network, you might overfit your model, and to choose when to stop training you can split the data into a training and validation set. You can for example use 50000 images for training the network and reserve 10000 for validating it. By classifying the validation images, you can measure if you have overfitted your model, which is seen by a drop in classification performance of the validation data. In addition there are 10000 images for testing, but these should only be used for evaluating the performance of your networks after they have been trained. Figure 9.2 and 9.3 show a classification performance example and examples of correctly and misclassified digits.

The rules for the competition are:

1. Implement your own neural network for classifying the MNIST images.
2. Train the neural network on a part of the training data (e.g. 50000 images) and validate it on another part (e.g. the remaining 10000 images).
3. When you are satisfied with the obtained result – upload the trained network together with Matlab or Python code for running it.
4. Hand in a description of your network and a guide on how to run your code

0	1	6	3	6	4	6	6	3	1
6	7	1	1	4	2	3	4	8	2
1	8	3	5	0	1	1	4	9	4
6	1	9	2	2	7	9	9	1	1
9	4	6	3	0	4	7	4	2	8
9	0	1	1	7	3	0	3	4	1
1	6	2	4	2	9	5	1	5	1
0	7	2	2	8	1	4	6	0	8
7	7	1	7	2	3	7	8	6	1
8	4	7	0	7	3	0	3	4	2

Figure 9.1: Example of the MNIST images.

94.14% success										
classification	0	1	2	3	4	5	6	7	8	
0	965	0	8	0	1	8	12	2	4	9
1	0	1113	0	0	2	1	3	14	1	6
2	1	3	951	17	3	5	3	21	10	1
3	2	2	13	944	1	20	1	5	18	12
4	0	0	9	0	930	5	9	8	7	28
5	4	2	3	22	1	817	7	0	16	4
6	5	3	7	2	12	12	920	0	8	1
7	1	2	14	12	3	4	1	952	11	10
8	2	10	23	9	4	15	2	3	891	7
9	0	0	4	4	25	5	0	23	8	931
	0	1	2	3	4	5	6	7	8	9
	target									

Figure 9.2: Table showing a classification performance example of a classification of the MNIST handwritten dataset.

classification									
classification	0	1	2	3	4	5	6	7	8
0	00	02	4	55	67	88	99		
1	11	11	1	1	1	1	1	1	1
2	0	12	23	34	45	56	67	78	89
3	1	22	33	44	55	66	77	88	99
4	22	33	44	55	66	77	88	99	99
5	10	27	34	41	52	63	74	85	96
6	01	12	23	34	45	56	67	78	89
7	00	13	24	35	46	57	68	79	80
8	11	22	33	44	55	66	77	88	99
9	01	12	23	34	45	56	67	78	89
	0	1	2	3	4	5	6	7	8
	target								

Figure 9.3: Examples of classified handwritten digits and misclassified digits.

5. Be a fair player and do not use the MNIST test data for training your network (can be found on the internet, but do not use it!).
6. Do an effort in making the code efficient.

9.1 *Modifications of your network for MNIST*

The classification will be using a fully connected feed forward neural network, similar to the one you implemented last exercise. But in contrast to last exercise, where data was points in two dimensions, you now have images. We treat these as vectors, so even though MNIST images are only 28×28 pixels the vector representation is 784 dimensions. Therefore, the network should take in 784 dimensional vectors and return a 10 dimensional vector for classifying the digits 0 to 9. You can use the book on deep learning by Goodfellow et al.¹ to get ideas for this exercise.

Obtaining a strong classifier requires many iterations of the backpropagation algorithm using 50000 training images. Therefore, it is a good idea to utilize vectorized code in your implementation. This can be done by computing the gradients for subsets of the training data using minibatches. You can have a minibatch as a matrix and compute the forward propagation and gradients using matrix operations. By averaging the gradients obtained from the minibatch, the backpropagation can be carried out in the same way as you would do when training with one sample at a time. Due to the averaging, the obtained gradients are less affected by noise and it is typically possible to have higher learning rates.

A part of optimizing the neural network is by changing its architecture. Therefore, it is recommended that you implement your network such that you can change the number of layers and the number of neurons in each layer.

1. Implement a fully connected neural network for MNIST classification. The data should be normalized and centered, i.e. vectors of unit length using the 2-norm and with zero mean. Besides vectorizing the code it is also worth considering the data type. Single is faster than double, so you should consider if you want faster computations, at the cost of lower precision. You can experimentally evaluate if the high precision is necessary.
2. Train the network and plot the training and validation error for each iteration (epoch). The dataset could be split into 50000 images for training and 10000 for validation.

¹ Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT press Cambridge, 2016

9.2 Optimizing the neural network

A large number of techniques for optimizing the performance of the neural network has been proposed. Neural networks are typically initialized with random numbers, but the performance depends on the choice of the initialization strategy. You can therefore optimize the network by experimenting with different initialization strategies.

Optimization with stochastic gradient decent can be slow, but updating the gradients using momentum can accelerate the learning rate. Momentum is obtained by computing the gradients as a weighted combination of the previous gradient and the new gradient. Hereby, the gradient is computed as a moving average with exponential decay. Another way of ensuring convergence of the gradient decent is by adapting the learning rate. Here you can adapt the learning rate to the individual gradient estimates.

1. Implement one or more optimization strategies and document how it affects the obtained result.

9.3 Regularization methods

It is important that the neural network generalizes well such that it can classify new unseen data. Since neural networks often have many parameters it is easy to overfit the model, especially on small datasets where a very low training error can be obtained, but the validation error will be high. One way to overcome the problem with training a neural network on small datasets is through dataset augmentation, where fake data is fabricated by small modifications of the input data. This can be done by small permutations or by adding small amounts of noise. Hereby a much larger dataset can be obtained, which can help the training.

Instead of adding noise to augment the training data, small amounts of noise may be added to the hidden units in the network. You can add random noise in each minibatch iteration. Noise can also be added to the output targets for obtaining better performance.

Dropout is another method for regularizing the neural network. Here a random selection of neurons are set to zero during each minibatch iteration leaving out these neurons in that iteration. Setting the neurons to zero resembles having a number of different neural networks and is inspired by ensemble methods.

1. Try experimenting with regularization methods. You can also get inspired by architectures that other people have had success with.

9.1 Information for 02506

Exam questions related to this exercise Explain how a fully connected neural network can be used for image classification. This includes explanation about the input, functionality of the network, and the backpropagation algorithm. You can also discuss

- Optimization of a neural network. This could e.g. be momentum or adaptive learning rate.
- Exemplify methods for regularizing a neural network. This could e.g. be data augmentation, regularization by adding noise, or dropout.

10 Convolutional neural networks

CONVOLUTIONAL NEURAL NETWORKS (CNNs) have many aspect in common with multilayer perceptrons (MLPs – fully connected feed forward neural networks) such as being a directed acyclic graph with weighted edges and non-linear activations. Instead of having unique weights for all connections, the CNNs share their weights, which means that only one edge weight is learned for many edges. This results in a significant reduction in number of model parameters, and therefore makes it possible to have much larger input data compared to MLP networks. Furthermore, the shared weights can efficiently be implemented as convolutions, which for images are 2D convolutions. The parameters that must to be learned are the weights of the convolutional kernels, which similar to MLPs, are learned using backpropagation.

Working with images in 2D makes it possible to apply additional operations to the hidden layers. This includes for example a pooling step typically combined with a down-scaling step. Max-pooling is an example of a widely used pooling method, where pixels are replaced by the local maximum in a local neighborhood. Max-pooling ensures that only the important features are kept, and makes the analysis robust to small translations. There is a number of other operations that can be applied, and an overview is given in chapter 9 in Goodfellow et al.¹.

Despite that CNNs have many aspects in common with MLPs, they are typically more complicated and therefore not as simple to implement. A large number of software frameworks are available, and training neural networks for many engineering applications will involve GPU processing. This is however implemented in a user friendly way in many of the software packages and highly sophisticated neural networks can be developed using high-level programming using e.g. Python that makes these frameworks easy to use.

10.1 Exercise

In this exercise we will work with existing software frameworks and we will use a pre-trained network, namely the VGG19². We will

¹ Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT press Cambridge, 2016

² Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014

use a network that is trained for the 1000 categories in the ImageNet challenge³, but instead of classifying photographs we will use it for classifying microscope images of histopathology tissue samples.

The data originates from the CAMELYON17 challenge⁴ for breast cancer diagnosis. In clinical practice, a pathologist will spend long time inspecting images of tissue that has led to much research in automating this inspection task, which is also the goal of the CAMELYON17 challenge. Here the aim is to detect and classify breast cancer metastasis from so-called whole-slide images of stained histological lymph node sections. In this exercise you will do the core part of this analysis, namely classifying samples to healthy or diseased tissue.

In Figure 10.1 some examples of the histopathology images are shown. These are 224 by 224 pixels color images cropped from whole-slide images, and you should note that the visual difference between normal and diseased tissue is small. The basis for the exercise is to classify these images using pre-trained convolutional neural networks.

Data is provided in file called `histo_images.mat` where the images are stored in a 4D array called `histo_images` and each image is labeled given in a 1D array called `labels` with 1 for normal tissue and 2 for diseased tissue. There are 992 images in all where 629 are from normal tissue and 363 are from diseased tissue.

Suggested procedure

1. Load in the data, visualize and get familiar with the format

10.2 Setting up the CNN framework

If you use MATLAB you should use MatConvNet⁵ and if you use Python you should use Keras with Tensorflow⁶. You should use the VGG19 model pre-trained for the ImageNet ILSVRC classification task.

Suggested procedure

1. Download the pre-trained VGG19 model for the ImageNet challenge.
2. Get MatConvNet or Keras to work on your computer.
3. Make a forward pass on your computer and investigate the output of the network including access to the hidden layers.

10.3 Classifying images

You should do a simple classification of the images using k -nearest neighbors, based on features extracted from the downloaded CNN. A part of this is to investigate the features extracted in the hidden layers in

³ <http://www.image-net.org/>

⁴ <https://camelyon17.grand-challenge.org/>

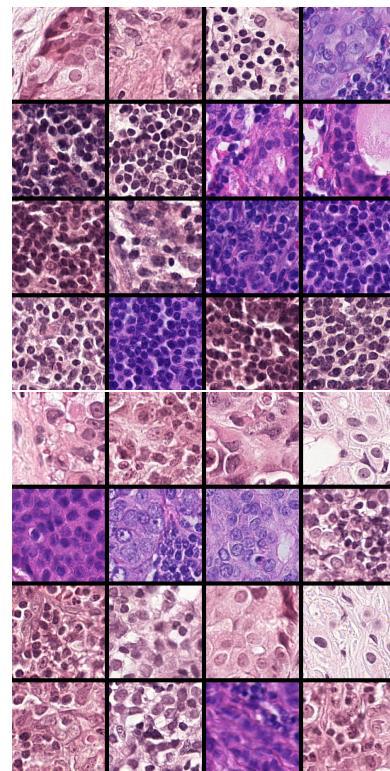


Figure 10.1: Example images of histological tissue samples from healthy tissue in the top and breast cancer metastasis in the bottom. These are from the CAMELYON17 challenge.

⁵ <http://www.vlfeat.org/matconvnet/>

⁶ <https://keras.io/>

the network. You should not expect improved performance by looking at the pooling or ReLU layers, since there is no information introduced in these layers. The features that you should use for classification might not be the last layer, since it is trained for a very specific task, but still the chosen layer should be at a deep layer in the network. Therefore, it is suggested that you look at the fully connected layers, and determine which is suitable for the classification task.

Suggested procedure

1. Extract and store features for the histopathology images provided for the exercise. Be careful to preprocess the images, such that they are normalized and fed into the model.
2. Do a k -nearest neighbor classification based on the extracted features by measuring the Euclidean distance between image features. You should take one out and measure the distance to 991 other images and classify to the most frequent occurring label among k neighbors. Investigate how k should be chosen.
3. Make an overview of classification performance with different choices of k and layer from the CNN. Which combination performs best?

10.1 Information for 02506

Exam questions related to this exercise

- Explain the principle of convolutional neural networks (CNNs) including shared weights, pooling, tensor convolution, etc.

Part IV

Mini projects

HERE, WE COVER a number of topics relevant for image analysis.

Information for 02506

In this exercise you are free to choose the image analysis case that you find most interesting. It is expected that you will find a relevant image analysis problem, implement a suitable analysis method, set up test data that verifies the correctness of your implementation, and report and present your results on a poster for the last day at the course.

You are welcome to come up with your own idea for a project, but you are also welcome to choose from the suggested projects in the following.

10.2 Information for 02506

Exam questions related to this exercise

- Shortly explain the problem, choice of method, the algorithm and your implementation, and the results you obtained.

11 Texture analysis

IMAGE TEXTURE is important for a range of image analysis problems like object classification and quality control. Also a number of image processing problems like denoising and inpainting are based on principles of texture analysis. Here you will solve a texture classification based on the Basic Image Features described in Crosier and Griffin¹ and an inpainting problem described in the paper by Efros and Leung².

11.1 Data

The data for the exercise is found in the file called `texture_data.zip` that contains images with a wide variety of textures for BIF characterization used in the first part and corrupted images to be used during the texture synthesis part of the exercise. Examples images are shown in Figure 11.1. You are also welcome to find your own data set for the exercise.

11.1.1 Basic Image Features

This section will provide a small summary of how BIFs are estimated. The purpose of BIF is to go from an image of high dimensionality to a lower dimensional vector representation of the image texture. This representation uses simple geometric image features and will enable differentiation between different textures. The following recipe is used to estimate BIF:

1. Convolve with six Gaussian filters to get scale-normalized filter responses $(s, s_x, s_y, s_{xx}, s_{yy}, s_{xy})$.
2. Calculate the flat, slope, blob ($2\times$), line ($2\times$) and saddle feature responses using the formulas from³ and $(s, s_x, s_y, s_{xx}, s_{yy}, s_{xy})$ from Step 1.
3. Classify each pixel as flat=0, slope=1, dark blob=2, white blob=3, dark line=4, white line=5, saddle=6, by finding the label index of the maximum feature responses of Step 2. Denote the resulting label image as \mathcal{L} .

¹ M. Crosier and L.D. Griffin. Using basic image features for texture classification. *International Journal of Computer Vision*, 88(3):447–460, 2010. ISSN 0920-5691. DOI: 10.1007/s11263-009-0315-0

² A.A. Efros and T.K. Leung. Texture synthesis by non-parametric sampling. In *Computer Vision*, 1999. *The Proceedings of the Seventh IEEE International Conference on*, volume 2, pages 1033–1038 vol.2, 1999

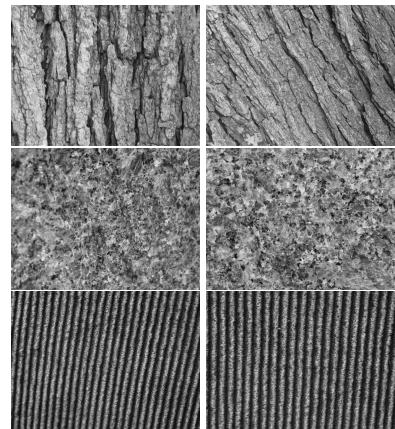


Figure 11.1: Examples of three classes of textured images.

³ M. Crosier and L.D. Griffin. Using basic image features for texture classification. *International Journal of Computer Vision*, 88(3):447–460, 2010. ISSN 0920-5691. DOI: 10.1007/s11263-009-0315-0

For a fixed scale, a seven bin histogram can now be formed by counting how many times a pixel is classified as one of the classes. This histogram is the BIF of an image for scale σ . Please note that if we discard the flat pixels, a six bin histogram is used per scale, however we are generally interested in a histogram that also models scale.

Four scales: How to get a histogram? When extending this BIF representation to multiple scales, the process of forming a histogram becomes increasingly complicated. If we choose four scales $\sigma = (1, 2, 4, 8)$ we need to run steps 1 - 3 four times. This will lead to a four channel label image $\mathcal{L}(\mathbf{x}; i)$, where \mathbf{x} is the position in the image and $i = 0, \dots, 3$ for the four scales. To get the texture characterization, we have to convert these four channels of the label image into a histogram. Each bin of this histogram counts how often a specific label configuration occurs across the four scales. If we ignore flat BIF regions, the pixels can be classified as one of the labels $\mathcal{L}(\mathbf{x}; i) \in \{1, \dots, 6\}$. First we want to translate this to a number between 0 and 1295 (i.e. $6^4 = 1296$ unique combinations). This is done in all pixels resulting in an image $\mathcal{B}(\mathbf{x})$ by converting the four BIF classes to one number

$$\mathcal{B}(\mathbf{x}) = \sum_{i=0}^3 (\mathcal{L}(\mathbf{x}; i) - 1) 6^i. \quad (11.1)$$

This means that the label combination $[1, 1, 1, 1]$ is a pixel classified as slope on all the scales, and similarly $[1, 3, 4, 6]$ is a pixel that is classified as a slope at the first scale, a blob on the second scale, a line on the third scale, and a saddle on the fourth scale.

11.2 Texture classification

In this exercise you will experiment with *Basic Image Features* (BIF) for texture description.

The BIF features are computed from the following equations:

Classify according to the largest of the features:

Flat: ϵs

Slope: $2\sqrt{s_x^2 + s_y^2}$

Blob: $\pm \lambda$

Line: $2^{-\frac{1}{2}}(\gamma \pm \lambda)$

Saddle: γ

where

$$\lambda = s_{xx} + s_{yy}$$

$$\gamma = \sqrt{(s_{xx} - s_{yy})^2 + 4s_{xy}^2}$$

Suggestions for experiments:

- Illustrate the BIF response in some images using color codes similar to how this is done in Crosier and Griffin⁴.
- Show the BIF histogram (set $\epsilon = 0, \sigma \in \{1, 2, 4, 8\}$) for an example image.
- Compare BIF histograms for a total of 30 images (6 texture classes, 5 images per class). Construct a 30×30 *confusion matrix* containing the histogram distances based on the L_1 -norm (sum of absolute difference). Show the histogram as an image and explain the pattern.

11.3 Texture synthesis: Task 2

In this exercise you will synthesize image texture using a method similar to the one presented in⁵. This method is based on fitting partly overlapping image patches to an image with holes by sampling (randomly) from a distribution of similar image patches. The distribution is approximated from the image itself by measuring distances to patches from the image itself.

Suggestions for experiments:

- Choose or construct a simple test example with repeated texture and a small hole and fill in the missing part.
- Choose a natural image and fill in a hole. Try varying number of patches, patch size, hole size, type of image, etc.
- Can the method be used for noise reduction? Experiment with e.g. salt and pepper noise.

⁴ M. Crosier and L.D. Griffin. Using basic image features for texture classification. *International Journal of Computer Vision*, 88(3):447–460, 2010. ISSN 0920-5691. DOI: 10.1007/s11263-009-0315-0

⁵ A.A. Efros and T.K. Leung. Texture synthesis by non-parametric sampling. In *Computer Vision*, 1999. *The Proceedings of the Seventh IEEE International Conference on*, volume 2, pages 1033–1038 vol.2, 1999

12 Optical flow

SMALL MOVEMENTS between two consecutive frames of an image series can be modeled as optical flow. The problem of optical flow is to determine local translations between two frames as a vector field such that the brightness constancy constraint

$$I(x, y, t) = I(x + \Delta x, y + \Delta y, t + \Delta t), \quad (12.1)$$

is fulfilled. Here, x and y are spatial coordinates and t is time. In this exercise you will implement methods for computing optical flow of the movement between two images.

A number of algorithms have been suggested for solving the flow problem, and a simple solution is to match patterns locally using block matching, where a window around a point in the first image is translated and compared to a window in the other image using e.g. sum of squared differences. This is however a time consuming task, so other methods based on computing differentials have been suggested. These include the Lucas-Kanade¹ and the Horn-Shunck methods² that you will work with in this exercise. This exercise is based on the book Computer Vision: Algorithms and Applications³, Chapter 8 (can be downloaded from <http://findit.dtu.dk/>). But you may also find relevant information from the internet and the two original papers.

In this exercise, different approaches for solving the optical flow problem are given, and it is suggested that you start by working with the basic elements of *Optical flow* and then try working with the Lucas-Kanade method or the Horn Shunck method and preferably both. These methods have a number of common elements, so when one is implemented it is relatively easy to implement the other.

¹ B D Lucas and T Kanade. An iterative image registration technique with an application to stereo vision. 1981

² Berthold KP Horn and Brian G Schunck. Determining optical flow. *Artificial intelligence*, 17(1-3):185–203, 1981

³ Richard Szeliski. *Computer vision: algorithms and applications*. Springer Science & Business Media, 2010

12.1 Optical flow

The assumption behind optical flow is that the movement between frames is small. Therefore, the optical flow can be computed by

$$\frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v = -\frac{\partial I}{\partial t}, \quad (12.2)$$

where u and v are displacements in the horizontal and vertical directions respectively. $[u, v]^T$ is the velocity of the flow field.

Task: Derive (12.2) from (12.1) by using a first order Taylor approximation.

In (12.2) two unknown parameters (u, v) must be computed, but in a single pixel there is just one equation, so the problem is underdetermined. If a small neighbourhood around a pixel is assumed to have the same displacement, it will be possible to solve (12.2) as a linear least squares problem, where we have $\mathbf{A}\mathbf{u} = \mathbf{b}$, where $\mathbf{u} = [u, v]^T$.

Task: Write up the elements of \mathbf{A} and \mathbf{b} for a 3×3 neighbourhood.

Two small images of 10×10 pixels called `composedIm_1.png` and `composedIm_2.png` are available in the `optical_flow_data.zip` file on Campusnet. They are made from an image by extracting two patches shifted by one pixel. Furthermore a 3×3 region of the same pattern is placed in the image, but shifted in the opposite direction as shown in Figure 12.1. You can use these two images to try simple experiments with optical flow.

Task: Compute the optical flow vector for a window of 3×3 pixels centered at $(r, c) = (2, 6)$ and $(r, c) = (5, 4)$, where r is the row and c is the column. You should use a simple pixel differences to compute the differential by using the central difference filters $[-1, 0, 1]$ and $[-1, 0, 1]^T$. You can also use $[-1, 1]$ and $[-1, 1]^T$, but then the derivative is computed between pixels. You can ignore this and compare the result to the central difference filter.

The least squares solution to $\mathbf{A}\mathbf{u} = \mathbf{b}$ is found by solving the minimization problem

$$\arg \min_{\mathbf{u}} \|\mathbf{A}\mathbf{u} - \mathbf{b}\|^2. \quad (12.3)$$

Taking the derivative with regards to \mathbf{u} and setting to zero yields

$$\mathbf{u} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{b}. \quad (12.4)$$

This allows us to precompute $\mathbf{A}^T \mathbf{A}$ and $\mathbf{A}^T \mathbf{b}$ as a number of sums over the image that can be obtained efficiently by filtering.

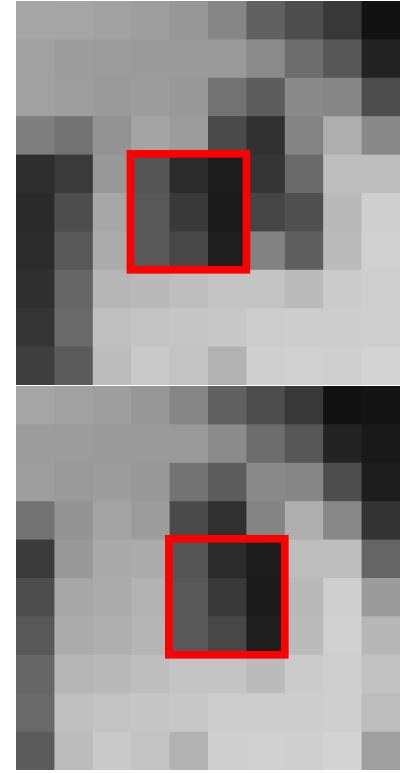


Figure 12.1: Two images of the same pattern shifted one pixel to the left, whereas the center part marked in the red box is shifted one pixel to the right.

Task: Write up the elements of $\mathbf{A}^T \mathbf{A}$ and $\mathbf{A}^T \mathbf{b}$.

Task: Precompute the input needed for $\mathbf{A}^T \mathbf{A}$ and $\mathbf{A}^T \mathbf{b}$ and implement a filtering scheme that sums the elements for the small test images `composedIm_1.png` and `composedIm_2.png`. Compute the flow vectors for the full images.

Task: Display the flow vectors on top of the images using the MATLAB function `quiver`.

12.2 Lucas-Kanade method

In this and the next part you should experiment with larger images. There are some benchmark images available from the Middlebury benchmark homepage <http://vision.middlebury.edu/flow/>. Some images from here have been uploaded to CampusNet that you can use for this exercise. But you are also welcome to experiment with your own images.

What you implemented in until now is a simple version of the Lukas-Kanade method. When you are working with larger images, there are however some practical aspects to consider. The linear system $\mathbf{A}\mathbf{u} = \mathbf{b}$ may be ill posed such that there are no vector \mathbf{u} that fulfills the equation. If this is the case, the 2×2 matrix $\mathbf{A}^T \mathbf{A}$ does not have an inverse.

Task: Find a way to check if there is a solution to the equation $\mathbf{A}\mathbf{u} = \mathbf{b}$, i.e. that $\mathbf{A}^T \mathbf{A}$ has an inverse. Implement that in your solution for computing the optical flow vector field.

Task: Compute and display the flow field in two larger images from e.g. the Middlebury dataset.

In the first part of this exercise you implemented the image differential using pixel differences. There are also other methods for computing image differentials like central differences using the filter $[-1, 0, 1]$ and its transpose. You may also the first order derivative of a Gaussian.

Task: Test one or more differential filters.

Instead of treating all pixels in the window equally, better results can be obtained by weighting the pixels using a weight matrix $\mathbf{W}\mathbf{A}\mathbf{u} = \mathbf{W}\mathbf{b}$ where \mathbf{W} is a diagonal weight matrix, resulting in the least squares solution $\mathbf{u} = (\mathbf{A}^T \mathbf{W}\mathbf{A})^{-1} \mathbf{A}^T \mathbf{W}\mathbf{b}$. This can efficiently be implemented using a weight filter e.g. a Gaussian.

Task: Implement a weighted sampling window as a filter operation. Display the result on test images using different filter size.

Task: Comment on performance and processing time for the Lucas-Kanade method.

12.2.1 Horn-Shunck method

A problem with the Lucas-Kanade method is that it operates locally, so in regions with no texture it is not possible to compute the flow vectors. This is solved in the Horn-Shunck method where the Laplacian over the vector field is minimized in addition to ensuring that the brightness is constant by minimizing

$$E = \int \int [(I_x u + I_y v + I_t) + \alpha^2 (\|\nabla u\|^2 \|\nabla v\|^2)] dx dy. \quad (12.5)$$

The energy E is minimized by iteratively updating the flow vectors using the update rules

$$u^{k+1} = \bar{u}^k - \frac{I_x(I_x \bar{u}^k + I_y \bar{v}^k + I_t)}{\alpha^2 + I_x^2 + I_y^2} \quad (12.6)$$

$$v^{k+1} = \bar{v}^k - \frac{I_y(I_x \bar{u}^k + I_y \bar{v}^k + I_t)}{\alpha^2 + I_x^2 + I_y^2} \quad (12.7)$$

where \bar{u}^k is the average flow over a window in the x -direction and

$$I_x = \frac{\partial I}{\partial x}, \quad I_y = \frac{\partial I}{\partial y}.$$

Task: Implement the Horn-Shunck method and test it on two larger images from e.g. the Middlebury dataset. Display the flow vectors.

The choice of the α parameter and the number of iterations influence the smoothness of the obtained result. Also, the choice of how the image is differentiated will affect the performance.

Task: Display results with varying parameter choices. Display results of a good and a bad choice of α . Do the same for number of iterations and size of averaging.

Task: Experiment with different choice of differentiation method.

Task: Comment on performance and processing time for the Horn-Shunck method.

13 Layered surfaces

A VOLUMETRIC SEGMENTATION PROBLEM can often be constrained in terms of topology. We may for example be interested in segmenting a roughly spherical object, tubular objects or surfaces. Such topological constraints strongly reduce the solution space, and may turn an otherwise challenging problem into a solvable problem, an example is shown in Figure 13.1.

In this mini-project we focus on optimal net surface detection via graph search originally suggested by Wu and Chen¹ for segmenting terrain-like surfaces. They construct a graph on a set of sample points from a volume, such that the roughness of possible solutions is constrained. The optimality of the solution is defined in terms of a volumetric cost function derived from the data. The approach has been extended for finding multiple interrelated layered terrain-like and tubular surfaces², which made it applicable for medical image segmentation and led to further extensions. One extension involves the volumetric cost function which determines an optimal placement of the surface. This was originally defined only in terms of on-surface appearance, and has been extended to incorporate appearance of the regions between surfaces³.

We will here review an algorithm for finding optimal layered surfaces, with the focus on the inputs and the outputs. For details on *how* this algorithm works, the reader is referred to article by Li et al. We also very briefly cover the principle of transforming the data into volumetric cost, which is the input to the layered surface detection algorithm.

13.1 Layered surface detection

In a discrete volume $x \in \{1, \dots, X\}$, $y \in \{1, \dots, Y\}$, $z \in \{1, \dots, Z\}$, a terrain-like surface s defined by $z = s(x, y)$ meets a smoothness constraint (Δ_x, Δ_y) if

$$|s(x, y) - s(x - 1, y)| \leq \Delta_x \quad \text{and} \quad |s(x, y) - s(x, y - 1)| \leq \Delta_y. \quad (13.1)$$

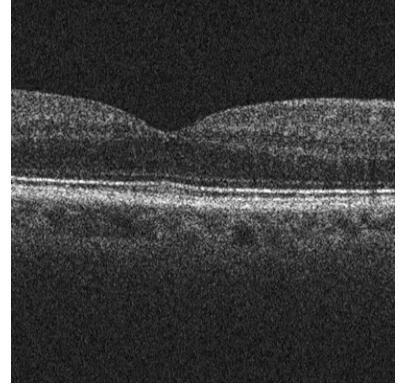


Figure 13.1: OCT (optical coherence tomography) image of retina. Quantifying the thickness of retinal layers is of clinical importance.

¹ Xiaodong Wu and Danny Z Chen. Optimal net surface problems with applications. In *Automata, Languages and Programming*, pages 1029–1042. Springer, 2002

² Kang Li, Xiaodong Wu, Danny Z Chen, and Milan Sonka. Optimal surface segmentation in volumetric images – a graph-theoretic approach. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 28(1):119–134, 2006

³ Mona Haeker, Xiaodong Wu, Michael Abràmoff, Randy Kardon, and Milan Sonka. Incorporation of regional information in optimal 3-d graph search with application for intraretinal layer segmentation of optical coherence tomography images. In *Information Processing in Medical Imaging*, pages 607–618. Springer, 2007

For a cost volume $c(x, y, z)$, an *on-surface* cost of s is defined as

$$C_{\text{on}}(s, c) = \sum_{x=1}^X \sum_{y=1}^Y c(x, y, s(x, y)). \quad (13.2)$$

The *optimal net surface problem* is concerned with finding a terrain-like surface with a minimum cost among all surfaces satisfying smoothness constraint.

The polynomial time solution presented in the work by Wu and Chan transforms the optimal net surface problem into that of finding a *minimum-cost closed set* in a node-weighted directed graph with nodes representing volume voxels. This is further transformed into a problem of finding a *minimum-cost s-t cut* in a related arc-weighted directed graph. Minimum-cost *s-t cut* can be solved in polynomial time and efficiently found using algorithm of Boykov and Kolmogorov ⁴, a well known tool for many image segmentation tasks. While the optimal net surface problem is ultimately solved using the minimum-cost *s-t cut* algorithm, it should be noted that the graph constructed for surface detection is rather different than when used for Markov random fields.

The extension to multiple surfaces developed in ⁵ may be exemplified by considering two terrain-like surfaces s_1 and s_2 . The surfaces are said to meet an overlap constraint $(\delta_{\text{low}}, \delta_{\text{high}})$ if

$$\delta_{\text{low}} \leq |s_2(x, y) - s_1(x, y)| \leq \delta_{\text{high}}. \quad (13.3)$$

Given two cost volumes c_1 and c_2 the total cost associated with surfaces s_1 and s_2 is

$$C_{\text{on}}(s_1, c_1) + C_{\text{on}}(s_2, c_2)$$

and the optimal surface detection will return a pair of surfaces with a minimum cost among all surfaces satisfying overlap and smoothness constraint. Depending on the problem at hand c_1 and c_2 may be different or identical, and likewise smoothness constraints may vary or be the same for the two surfaces.

Finally, two *layered* (i.e. non-intersecting and ordered) surfaces give rise to an *in-region* cost corresponding to the volume between two surfaces, and defined by

$$C_{\text{in}}(s_1, s_2, c_{1,2}) = \sum_{x=1}^X \sum_{y=1}^Y \sum_{z=s_1(x,y)+1}^{s_2(x,y)} c_{1,2}(x, y, z). \quad (13.4)$$

This cost, together with the cost for the region under the surface s_1 and the region over the surface s_2 , can be incorporated into the minimization problem ⁶. We use notation $c_{0,1}$ and $c_{2,3}$ for the cost volumes in connection with the boundary regions.

To summarize, for finding K cost-optimal layered surfaces we need to define

⁴ Y Boykov and V Kolmogorov. An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(9):1124–1137, 2004

⁵ Kang Li, Xiaodong Wu, Danny Z Chen, and Milan Sonka. Optimal surface segmentation in volumetric images – a graph-theoretic approach. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 28(1):119–134, 2006

⁶ Mona Haeker, Xiaodong Wu, Michael Abràmoff, Randy Kardon, and Milan Sonka. Incorporation of regional information in optimal 3-d graph search with application for intraretinal layer segmentation of optical coherence tomography images. In *Information Processing in Medical Imaging*, pages 607–618. Springer, 2007

- K on-surface cost volumes c_k , $k = 1, \dots, K$, and/or
- $K+1$ in-region cost volumes $c_{k,k+1}$, $k = 0, \dots, K$.

The set of feasible surfaces is given by

- K smoothness constraints (Δ_x^k, Δ_y^k) , $k = 1, \dots, K$ and
- $K-1$ overlap constraints $(\delta_{\text{low}}^{k,k+1}, \delta_{\text{high}}^{k,k+1})$, $k = 1, \dots, K-1$.

Layered surface detection has found an immediate use for detecting tubular surfaces. The main principle is the fact that a circle $x^2 + y^2 = \rho^2$ appears as a straight line $r = \rho$ when represented in polar (r, θ) coordinates. Detecting a tubular surface is achieved by representing the volumetric data in a cylindrical coordinate system (r, θ, z) with the longitudinal axis $r = 0$ roughly aligned with the center of the tube. We call this transformation *unwrapping* the volume, and we also say that the volume is sampled along the radial rays. Important practical parameters for unwrapping are the radial and the angular resolution. In the unwrapped representation, the tubular surface is terrain-like and can be defined as $r = s(\theta, z)$. When using layered surface detection for detection of tubular surfaces, additional constraints are added to ensure a smooth transition over $\theta = 0$.

13.2 Constructing cost volumes

The surfaces returned by the layered surface detection algorithm are optimal in terms of the volumetric cost. Therefore, to detect a surface we need to define a cost volume c_k which takes small values where a data $V(x, y, z)$ supports the surface k . This modelling step, crucial for the performance of the algorithm, is fully dependent on the data.

If the surface to be detected is characterized by a certain voxel intensity v_s , then the cost volume may be defined as $(V - v_s)^2$. More often, the surface divides two regions of different intensities, so cost volume needs to be defined in terms of change of intensity. When computing intensity changes for tubular surfaces, the best approach is to first unwrap the volume, and then compute the change in the r direction.

13.2.1 Examples

Figure 13.2 demonstrates the use of the layered surface detection.

13.3 Proposed mini-projects for layered surface detection

You may work with the layered surface detection in a number of ways. Contact the teacher to clarify the possibilities.

1. Implement the algorithm for layered surface detection. For finding

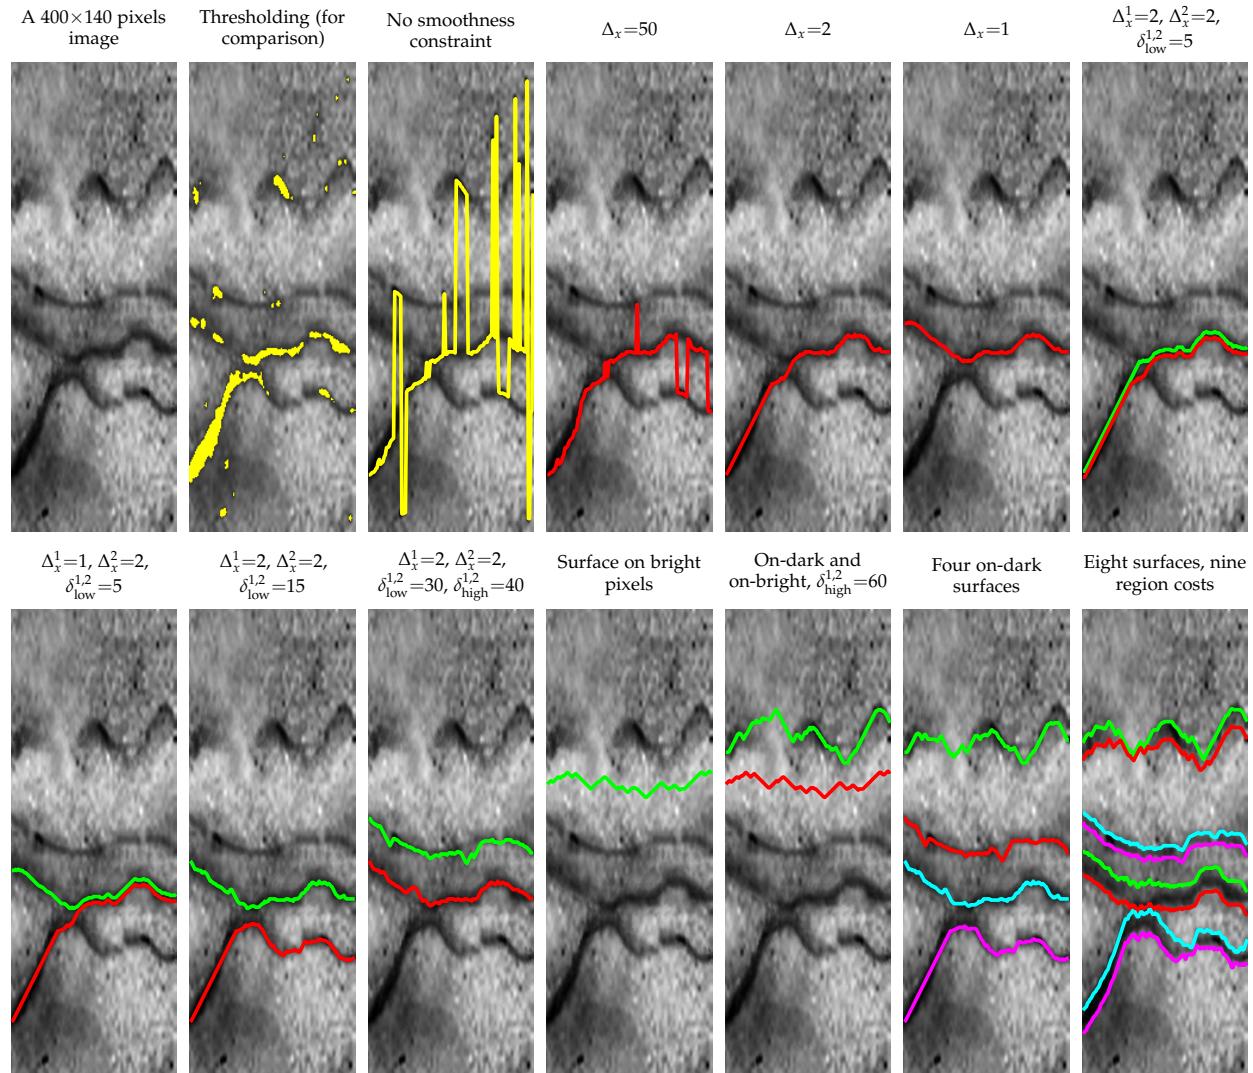


Figure 13.2: Output of layered surface detection. First three images serve to illustrate the problem). Images 4–6 show how changing smoothness constraint influences the result. Images 7–10 demonstrate the use of the overlap constraint. Images 11–12 demonstrate the use of different cost functions. Image 13 is a four-surface detection, while image 14 uses region costs.

- minimal $s-t$ cut use the same implementation as previously used for Markov random fields exercise. [Effort: medium.]
2. Adapt the layered surface detection algorithm so that it operates on the triangular mesh. [Effort: difficult.]

14 Spectral segmentation and normalized cuts

SPECTRAL CLUSTERING is an approach to data clustering problem, and it includes a number of related techniques. Spectral clustering is used in machine learning, computer vision and signal processing, with applications in processing speech spectrograms, DNA gene expression analysis, document retrieval and computation of Google page rank. The name *spectral* originates from the mathematical term *spectrum* (a set of the eigenvalues of a given matrix), and this is because spectral clustering utilizes eigenvalues and eigenvectors of the data similarity matrix.

Spectral clustering is one of the fundamental data clustering approaches, it is easy to implement and solve by standard linear algebra software, there is no assumptions on the nature of the clusters, and the techniques have been mathematically rigorously proved. Disadvantages include high computational and memory requirements of the direct implementation. For this reason the practical use of spectral clustering often involves computational simplifications and significant pre- or postprocessing.

Spectral approach has gained a great popularity for image segmentation following the seminal paper on normalized cuts by Shi and Malik¹. Recent uses of spectral approach is in the superpixel segmentation. Another intriguing example is the Copenhagen-based company Spektral (formerly known as CloudCutout) where spectral segmentation is a part of the successful green-screen removal product Figure 14.1.

Spectral methods can be applied to the data which is represented using a similarity matrix, and is often described in terms of graph partitioning. For a comprehensible coverage of (general) spectral clustering we recommend an excellent tutorial by von Luxburg². We will here briefly cover spectral clustering, and will then turn to its use in image segmentation.

Boiled down to four words, the essence of spectral clustering is: *Eigensolution gives graph partitioning*. To be able to understand spectral

¹ Jianbo Shi and Jitendra Malik. Normalized cuts and image segmentation. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 22(8):888–905, 2000



Figure 14.1: Spektral (formerly known as CloudCutout) green-screen product.

² Ulrike Von Luxburg. A tutorial on spectral clustering. *Statistics and computing*, 17(4):395–416, 2007

clustering, you need to be acquainted with the concept of graph cuts in order to describe the problem we want to solve. Furthermore, we need to represent a graph using an adjacency matrix and a closely related Laplacian matrix. Then we can see how eigensolution provides a solution to a graph cut problem.

14.1 Graph cuts, graph representations and eigensolutions

Recall that a graph consists of nodes and edges, and in general may be node-weighted, edge-weighted, directed or undirected. In context of spectral image segmentation, each pixel will correspond to a graph node, and pairs of pixels define graph edges – we will get back to image segmentation after covering the general case. For spectral clustering we work with edge-weighted undirected graphs which we represent using an adjacency matrix W with elements w_{ij} being the weight of the edge connecting the node i and a node j . Consider for example a graph in Figure 14.2 consisting of 12 nodes. This graph can be represented in terms of a 12×12 adjacency matrix, as illustrated in Figure 14.3.

A graph cut is a partitioning of a graph. The graph partitioning problems are concerned with finding a graph cut with the least cost. The simplest way of defining the cost of a cut is to consider all edges between the two partitions

$$\text{cut}(A, B) = \sum_{i \in A, j \in B} w_{ij},$$

but that might lead to unbalanced cuts. For this reason we might prefer using some other measure of the cut cost, which also consider the size of the partitions. Commonly used are normalized cuts

$$\text{Rcut}(A, B) = \frac{\text{cut}(A, B)}{|A|} + \frac{\text{cut}(A, B)}{|B|},$$

where we use the notation $|A|$ for a number of vertices in subset A , but one could also consider ratio cut and min-max cuts

$$\text{Ncut}(A, B) = \frac{\text{cut}(A, B)}{\text{vol}(A)} + \frac{\text{cut}(A, B)}{\text{vol}(B)},$$

$$\text{MMcut}(A, B) = \frac{\text{cut}(A, B)}{\text{cut}(A, A)} + \frac{\text{cut}(A, B)}{\text{cut}(B, B)},$$

where $\text{vol}(A)$ is weight of all edges associated with the subset, i.e. $\text{vol}(A) = \sum_{i \in A} d_i$ and $d_i = \sum_j w_{ij}$ is a degree of i th node. Figure 14.4 shows three graph cuts, while Figure 14.5 lists the costs associated with the three cuts given by different measures. You may confirm that the values are correct, and notice the different balancing properties of the cost measures.

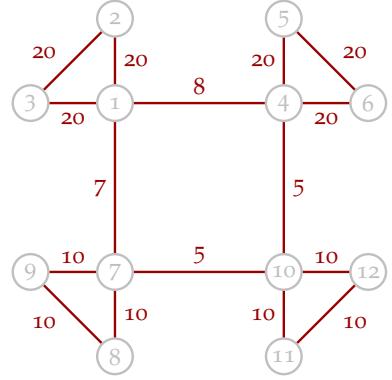


Figure 14.2: An example of an edge-weighted undirected graph.

	1	2	3	4	5	6	7	...
1	0	20	20	8	0	0	7	...
2	20	0	20	0	0	0	0	...
3	20	20	0	0	0	0	0	...
4	8	0	0	0	20	20	0	...
5								...
6								...
7								...
8								...
9								...
10								...
11								...
12								...

Figure 14.3: An adjacency matrix of a graph shown in Figure 14.2.

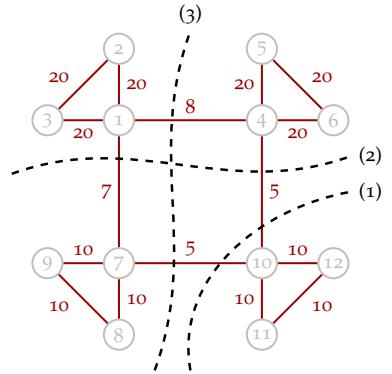


Figure 14.4: Three different partitioning of a graph.

	(1)	(2)	(3)
cut	10	12	13
Rcut	4.4	4.0	4.3
Ncut	0.1723	0.1293	0.1268
MMcut	0.3939	0.2784	0.2709

Figure 14.5: Values of the different cuts shown in Figure 14.4.

It turns out that the solution to the graph cut problem may be found by considering the eigensolution of the matrix closely related to the adjacency matrix – the graph Laplacian. Depending on the cost measure, the derivation will be slightly different, leading to the different (unnormalized and normalized) versions of the graph Laplacians. To normalize the Laplacian, we first define a diagonal degree matrix \mathbf{D} with diagonal elements being node degrees $d_i = \sum_j w_{ij}$.

We define unnormalized graph Laplacian

$$\mathbf{L} = \mathbf{D} - \mathbf{W},$$

and two normalized graph Laplacians

$$\mathbf{L}_{\text{sym}} = \mathbf{D}^{-1/2} \mathbf{L} \mathbf{D}^{-1/2} = \mathbf{I} - \mathbf{D}^{-1/2} \mathbf{W} \mathbf{D}^{-1/2},$$

$$\mathbf{L}_{\text{rw}} = \mathbf{D}^{-1} \mathbf{L} = \mathbf{I} - \mathbf{D}^{-1} \mathbf{W}.$$

These matrices are closely related to each other, and so are their spectra. All three matrices have real-valued eigenvalues with 0 being the smallest eigenvalue, see tutorial by von Luxburg for a more complete list of properties.

For our purposes, the most important are eigenvectors of \mathbf{L}_{rw} , i.e. vectors satisfying $\mathbf{L}_{\text{rw}} \mathbf{u} = \lambda \mathbf{u}$. The smallest eigenvectors (corresponding to smallest eigenvalues) yield solution relaxed versions of finding the normalized cut. The relaxed solution is subsequently transformed into an approximate discrete solution to the original problem, for example using k -means clustering. In particular, the second smallest eigenvector gives the partitioning of the data in two subsets – the very smallest eigenvector (corresponding to 0) is constant.

It can be shown that eigenvectors of \mathbf{L}_{rw} are generalized eigenvectors of \mathbf{L} and \mathbf{D} , see again von Luxburg's tutorial, Proposition (3) part 3. Therefore, to find the solution to normalized cut, one may also seek solution to generalized eigenproblem $\mathbf{L}\mathbf{u} = \lambda \mathbf{D}\mathbf{u}$. This is the formulation of normalized spectral clustering according to the original paper by Shi and Malik. In practice, solving a standard eigenproblem requires less computation.

We can utilize other matrix algebra identities to make computation of the spectra more accurate and/or efficient. Many eigensolvers are more accurate when working of symmetric matrices. A strategy exploiting matrix symmetry would involve computing eigenvectors of the (symmetric) \mathbf{L}_{sym} , and transforming those to eigenvectors of \mathbf{L}_{rw} by multiplying with $\mathbf{D}^{-1/2}$, see tutorial by von Luxburg, Proposition (3) part 2. Furthermore, if only a subset of eigenvectors is to be found, some eigensolvers are more efficient when finding eigenvectors corresponding to largest eigenvalues. This may be utilized by noticing that the smallest eigenvectors of $\mathbf{I} - \mathbf{A}$ are largest eigenvectors of \mathbf{A} .

14.2 Clustering 2D points

You should implement two functions. One function should take an affinity matrix and return eigenvectors corresponding to solution for normalized cuts. The second function should perform discretization of the eigenvectors using k -means.

You should first test your implementation on a small 2D point set. You can use points previously used for neural networks, but we also provide five point sets in a mat file `points_data.mat`. For each of the point sets you should:

- Visualize the point set, and identify the clusters (there are 2 clusters in set 3 and 5, and 3 clusters in set 1, 2 and 4).
- Construct the affinity matrix W . Use the fully connected graph and Gaussian similarity function (Luxburg, Section 2.2). You should initially estimate parameter σ so that it reflects the distance between the neighbouring points of the point cloud.
- Compute eigenvectors and clustering given by the normalized cut. Visualize the clustering.
- Determine ordering (permutation) of the points according to the clustering (so that points from the first cluster come first, followed by points in the second cluster, etc.)
- Visualize the values of the second eigenvector, first for unsorted points, then for sorted points.
- Visualize affinity matrix, and the affinity matrix for sorted points.
- Estimate the parameter σ which results in a meaningful clustering. You will need to change the parameter σ between point sets.

14.3 Image segmentation

Now we use spectral clustering on a pixels of a small image. Consider some of the provided test images. You might want to (drastically!) reduce the size of your images, to avoid memory problems.

You should:

- Construct the affinity matrix W using Equation (11) from article by Shi and Malik. Initially estimate parameters σ_I , σ_X and r . Instead of setting a radius r , for our small example you may use a fully connected graph (i.e. ignore if-otherwise condition of Equation (11) which sets affinity of distant points to 0).
- Visualize the spatial part of W , the brightness part of W and the final W .

- Compute eigenvectors and clustering given by the normalized cut.
- Visualize the values of the second eigenvector on the image grid.
- Visualize the segmentation results.
- Estimate the model parameters to obtain meaningful segmentation.

Start by the grayscale image. Use 2 clusters for plane and 5 clusters for vegetables. For the similarity (brightness, color) part of W treat an RGB value of each pixel as a vector to compute the Euclidian distance between the pair of pixels. Try also clustering the pixels using k-means clustering by treating RGB pixels values as vectors.

Consider adapting spectral methods to be able to handle larger images.

15 Probabilistic Chan-Vese

CHAN-VESE SEGMENTATION ALGORITHM alternates between two updates: update for mean intensities of the two regions given region boundaries, and update of the boundary between the two regions given mean intensities. The use of mean intensities implies that the two regions are distinguished by having different mean intensities. There are situations, where this is not the case, see Figure 15.1. Regions can be characterized by distributions of intensities or other features.

In this mini-project you will investigate generalizations of Chan-Vese algorithm, which allow for segmenting more challenging situations than with the original Chan-Vese. Some inspiration can be found in paper by Dahl and Dahl¹, which proposes a intensity-distribution approach Figure 15.2 and patch-distribution approach Figure 15.3. The approach is illustrated in Figure 15.2. Once the curve is initialized, instead of computing mean intensities for the inside and outside region, the distributions of intensities are collected for the inside and the outside region. For every pixel value we now have an information on how often it is in the inside and outside region, which can be translated into a probability that this pixel value is inside or outside. Computing such probabilities for all image pixels leads to probability image which can be used to deform the curve. Alternating, in a Chan-Vese manner, between computing probability image given the curve, and deforming the curve given probability image leads to segmentation.

For even more general case, instead of working with distributions of pixel intensities, distributions of image features may be used to compute the probability image. Figure 15.3 shows an example of using dictionary of image patches. For every patch from the dictionary we can compute the probability of it occurring in the inside or outside region.

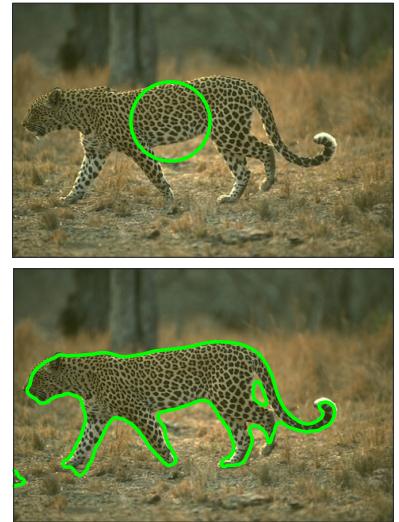


Figure 15.1: A deformable model used for segmenting forward-background image where two regions are characterized by different textures.

¹ Vedrana Andersen Dahl and Anders Bjorholm Dahl. A probabilistic framework for curve evolution. In *Scale Space and Variational Methods in Computer Vision*, pages 421–32. Springer, 2017

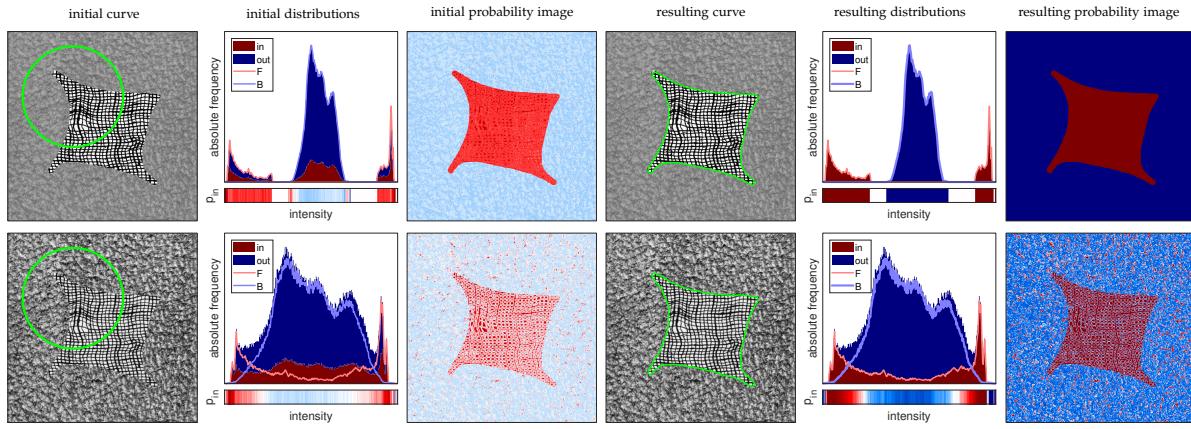


Figure 15.2: A probabilistic Chan-Vese approach when regions inside and outside are characterized by different distributions of pixel intensities. Top row shows easier problem of non-overlapping distributions, bottom row shows two overlapping distributions. Columns 1–3 show initialization, columns 4–6 show result after iterating.

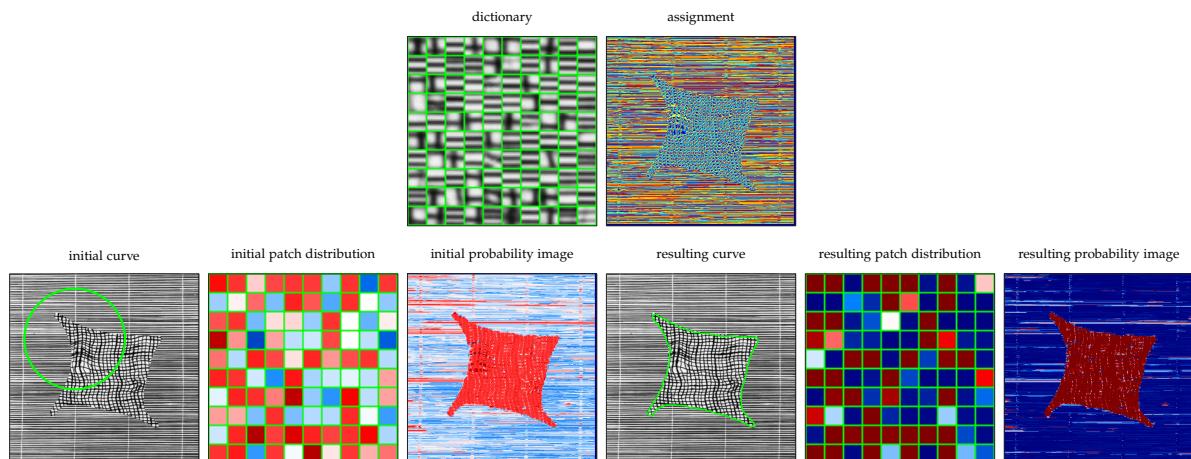


Figure 15.3: A probabilistic Chan-Vese approach when regions inside and outside are characterized by different texture. Top row shows used dictionary of image patches, and an assignment of image pixels to dictionary. In the second row, Images 1–3 show initialization and images 4–6 show result after iterating.

16 Orientation analysis

ANALYZING ORIENTATIONS OF IMAGES STRUCTURES is often needed if we want to visualize, quantify, or elsewhere utilize orientation information obtainable from images. A common tool for orientation analysis is a structure tensor.

In this mini-project you will be working with structure tensor and orientation analysis. You may decide to focus on computation of structure tensor, visualization of the orientation information, quantification of orientation, or some similar aspect.

16.1 Computing structure tensor

In the context of volumetric (3D) image analysis, a structure tensor is a 3-by-3 matrix which summarizes orientation in a certain neighborhood around a certain point.

For example, consider volumetric data showing a bundle of roughly parallel fibers. If we extract two cubes from this volume, mutually displaced along the predominant fiber orientation, the two cubes will have very similar intensities. For two cubes displaced along other orientations the cube intensities would be more different. For this reason, measuring the change of intensities between slightly displaced cubes may be used for determining predominant orientation of imaged structures.

It turns out that, given an initial point and a size of the cube, squared change of intensities may be expressed as $\mathbf{u}^T \mathbf{S} \mathbf{u}$, where \mathbf{u} is the direction of the displacement and \mathbf{S} is 3-by-3 symmetric positive semi-definite matrix – a structure tensor. Finding predominant orientation now amounts to finding \mathbf{u} which minimizes $\mathbf{u}^T \mathbf{S} \mathbf{u}$.

For a more formal derivation, consider a volumetric data V defined on the domain $\Omega \subset \mathbb{R}^3$, where $V(\mathbf{p})$ denotes voxel intensity at the point $\mathbf{p} = [p_x \ p_y \ p_z]^T$. Consider an arbitrary but fixed neighborhood N around a point \mathbf{p} , such that $N(\mathbf{p}) \subset \Omega$. We want to measure

$$D = \sum_{\mathbf{p}' \in N(\mathbf{p})} (V(\mathbf{p}' + \mathbf{u}) - V(\mathbf{p}'))^2 .$$

Assuming a small displacement we use first order Taylor expansion and arrive to

$$D = \sum_{\mathbf{p}' \in N(\mathbf{p})} ([V_x(\mathbf{p}') \ V_y(\mathbf{p}') \ V_z(\mathbf{p}')] \ \mathbf{u})^2$$

where we use notation $V_x = \frac{\partial V}{\partial x}$, and correspondingly for partial derivatives in y and z direction. Finally, exploiting commutativity of the inner product leads to

$$D = \mathbf{u}^T \sum_{\mathbf{p}' \in N(\mathbf{p})} \begin{bmatrix} V_x(\mathbf{p}') \\ V_y(\mathbf{p}') \\ V_z(\mathbf{p}') \end{bmatrix} [V_x(\mathbf{p}') \ V_y(\mathbf{p}') \ V_z(\mathbf{p}')] \ \mathbf{u}.$$

So, as earlier claimed, we arrived to expression $D = \mathbf{u}^T \mathbf{S} \mathbf{u}$, where \mathbf{S} is a 3-by-3 matrix – a structure tensor computed in a point \mathbf{p} and using a neighborhood N . That \mathbf{S} is symmetric and positive semi-definite follows directly from the construction of \mathbf{S} (note that $D \geq 0$).

Using a compact notation for gradient $\nabla V = [V_x \ V_y \ V_z]^T$, structure tensor is

$$\mathbf{S} = \sum \nabla V (\nabla V)^T.$$

Her we imply that structure tensor is computed for every voxel of the volume, i.e. structure tensor is a matrix-valued function over Ω . The summation is conducted over a neighborhood of each voxel, and result will be the same (up to the multiplicative factor) if summation is replaced by an averaging filter.

Two Gaussian filters are usually involved in computing structure tensor, see ¹ for detailed description of 2D case. The one Gaussian filter has to do with averaging orientation information in the neighborhood. This can be achieved using a convolution with a Gaussian K_ρ , where parameter ρ , called *integration scale*, reflects the size of the neighborhood. Now we have

$$\mathbf{S} = K_\rho * (\nabla V (\nabla V)^T).$$

The second Gaussian has to do with computing partial derivatives in gradient ∇V . To make differentiation less sensitive to noise we may convolve the volume with a Gaussian prior to computing derivatives. More efficiently, utilizing derivative theorem of convolution, partial derivatives can be computed by convolving with derivatives of Gaussian. We denote such gradient $\nabla_\sigma V$. The parameter σ is called *noise scale*. Expression with both Gaussians is

$$\mathbf{S} = K_\rho * (\nabla V_\sigma (\nabla V_\sigma)^T).$$

In summary, computing structure tensor for each voxel of a volume V involves three steps:

¹ Joachim Weickert. *Anisotropic diffusion in image processing*, volume 1. Teubner Stuttgart, 1998. URL <https://www.mia.uni-saarland.de/weickert/Papers/book.pdf>

1. Convolve V with derivatives of Gaussian with standard deviation σ to obtain V_x , V_y and V_z . For efficiency, use separability of Gaussian kernel.
2. Using element-wise multiplication compute six volumes V_x^2 , V_y^2 , V_z^2 , $V_x V_y$, $V_x V_z$ and $V_y V_z$.
3. Convolve each of the six volumes with the Gaussian kernel with standard deviation ρ . For efficiency, use separability of Gaussian kernel. The resulting volumes contain per-voxel elements s_{xx} , s_{yy} , s_{zz} , s_{xy} , s_{xz} and s_{yz} of the structure tensor

$$\mathbf{S} = \begin{bmatrix} s_{xx} & s_{xy} & s_{xz} \\ s_{xy} & s_{yy} & s_{yz} \\ s_{xz} & s_{yz} & s_{zz} \end{bmatrix} .$$

16.2 Computing orientations

Given structure tensor \mathbf{S} , predominant orientation is found by minimizing Rayleigh coefficient $\mathbf{u}^T \mathbf{S} \mathbf{u}$ through eigendecomposition of \mathbf{S} . Being symmetric and positive semi-definite \mathbf{S} yields three positive eigenvalues $\lambda_1 \leq \lambda_2 \leq \lambda_3$ and mutually orthogonal eigenvectors \mathbf{v}_1 , \mathbf{v}_2 and \mathbf{v}_3 . The eigenvector \mathbf{v}_1 corresponding to the smallest eigenvalue is an orientation leading to the smallest variation in intensities, which indicates a predominant orientation in the volume. Note that \mathbf{v}_1 is an orientation, and we usually represent it using a unit vector, but this is still not an unique representation since there are two opposite unit vectors sharing the orientation with \mathbf{v}_1 .

Eigendecomposition of a 3-by-3 real symmetric matrix can be computed efficiently using an analytic approach by Smith ² which uses an affine transformation and a trigonometric solution of a third order polynomial.

If there is no strong orientation in the volume, all eigenvalues will be similar, and dominant direction will be influenced by small local variations or the noise in the data. For this reason it is customary to analyze the ratio between eigenvalues to determine the degree of anisotropy in the neighborhood, and how (locally) line-like or plane-like the imaged structure is, see illustration 16.1. Inspired by diffusion tensor processing ³, we define values, so-called shape measures,

$$c_l = \frac{\lambda_2 - \lambda_1}{\lambda_3}, \quad c_p = \frac{\lambda_3 - \lambda_2}{\lambda_3}, \quad c_s = \frac{\lambda_1}{\lambda_3}$$

where c_l gives a measure of linearity, while c_p and c_s measure planarity and sphericity. Shape values are positive and sum to 1.

In summary, structure tensor is a 3-by-3 matrix that can be computed in each volume voxel. The computation of structure tensor requires

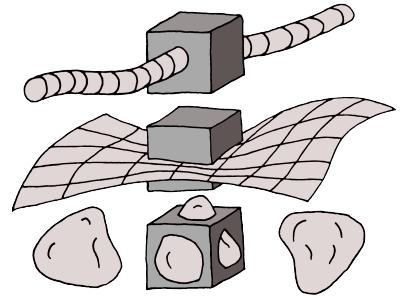


Figure 16.1: Neighborhoods and structures corresponding to linear, planar and spherical shape. On a linear structure (top) cubical neighborhood can move along the predominant direction leading to small change in intensities, while other two orthogonal directions lead to significantly larger and approximately equal change. On a planar structure (middle) two directions lead to small and approximately equal change in intensities, while third direction leads to significantly larger change. For a case with no predominant direction (bottom) the three orthogonal directions lead to roughly equal changes in intensities.

² Oliver K Smith. Eigenvalues of a symmetric 3×3 matrix. *Communications of the ACM*, 4(4):168, 1961. URL <https://dl.acm.org/citation.cfm?id=366316>

³ C-F Westin, Stephan E Maier, Hatsuho Mamata, Arya Nabavi, Ferenc A Jolesz, and Ron Kikinis. Processing and visualization for diffusion tensor MRI. *Medical image analysis*, 6(2):93–108, 2002

two parameters: noise scale σ and integration scale ρ . Being symmetric, structure tensor can be represented with 6 scalar values. The most important information extracted from structure tensor is a dominant orientation. Dominant orientation is a unit vector with equivalence relation $-\mathbf{v} \equiv \mathbf{v}$. Shape measures, also extracted from structure tensor, may also be of interest. Shape measures are three scalar c_l , c_p and c_s , summing to 1.

16.3 Visualization

Having extracted structure tensor and performed its eigendecomposition, following values are available for every volume voxel.

- Voxel intensity V , a scalar value in a certain range.
- Shape measures c_l , c_p and c_s , three scalar values summing to 1.
- Dominant orientation \mathbf{v}_1 , an orientation vector (unit vector with equivalence relation $-\mathbf{v} \equiv \mathbf{v}$).
- Other information, such as \mathbf{v}_2 , \mathbf{v}_3 , and the values λ_1 , λ_2 and λ_3 is also available, but usually of no special interest.

Visualizing this information often requires some care.

Shape measures, being three scalar values can be conveniently represented using three RGB color channels. Voxels with large c_l will appear red, large c_p will be green, and large c_s blue. This is the approach used in Figure 16.2.

Predominant orientation, is a 3D unit vector with equivalence relation $-\mathbf{v} \equiv \mathbf{v}$. A common way of visualizing orientation is to use absolute values of vector coordinates, i.e. $|v_x|$, $|v_y|$, and $|v_z|$, as three RGB color channels. Orientations roughly aligned with x direction will be red, those aligned with y green, and z blue. This has a desirable property that $-\mathbf{v}$ and \mathbf{v} map to the same color. Furthermore, when the imaged object has a certain geometry aligned with the coordinate system (for example, elongated object aligned with z axis), this coloring scheme may be favorable for the interpretation of orientations. Undesirable property of the RGB color scheme is that different orientations map to the same color, for example four orientations corresponding to main diagonals in unit cube all map to gray.

Shape measure and predominant orientation are computed for every voxel in the volume – regardless of whether the voxel is within the object or material which we investigate. When using volume rendering to visualize the extracted measures, if the value is shown in every voxel, the valuable information might be occluded. For this reason it might be beneficiary to combine visualization of extracted measures

with the intensity values, which carry information on voxels containing, or not containing, material. A visually pleasing result is obtained if voxels containing no material are shown transparent. Furthermore, predominant orientation is only relevant for voxels exhibiting high linearity, so shape measure may be combined with visualization of predominant orientation. This is the approach used in Figure 16.2.

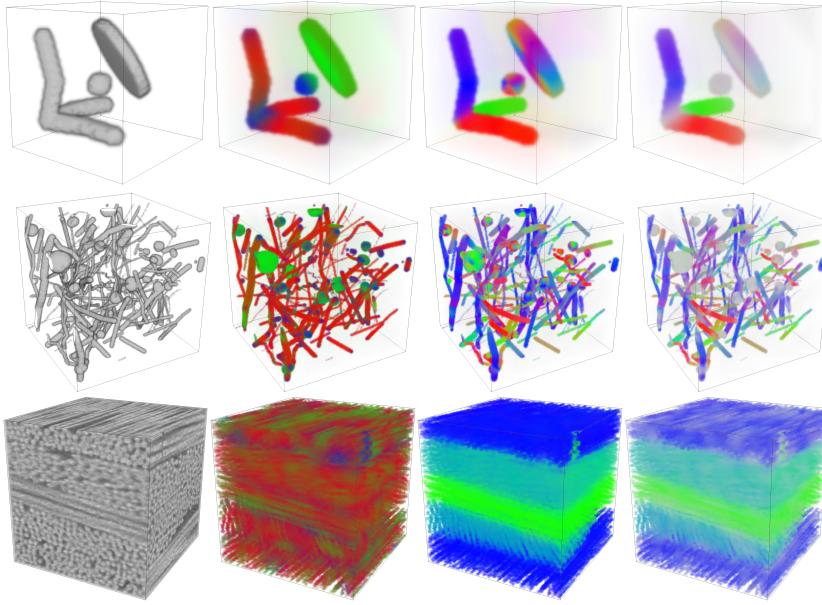


Figure 16.2: Orientation information. In each row, from left to right: 3D rendering of the volumetric data, shape measures visualized in colors, predominant orientation in the material phase, visualized in colors, predominant orientation weighed by the measure of linearity. Top two rows show the result on synthetic data, while in the bottom row we see the result of orientation analysis for composite material.

16.4 Applications

Some uses of orientation information are:

- Volumetric visualization of the shape measures and/or the predominant orientation.
- Producing histograms of orientations by binning orientation vectors. These distributions live on a half sphere, which can complicate the visualization, comparison, and fitting of distributions.
- Comparing orientation information extracted from volumetric data with, for example, the orientations obtained via modeling and/or simulation.
- Using local orientation to segment the volume into regions of constant orientation.
- Using local orientation to tune smoothness constraint in MRF segmentation.
- Using local orientation to guide fibre tracking.

17 CNN for segmentation

IMAGE SEGMENTATION is often needed as an intermediate step when quantifying structures in images, and we have previously been working with patch-based segmentation. In this exercise you should train a convolutional neural network to segment electron microscopy images of neuronal structures. This data was part of the ISBI Challenge: *Segmentation of neuronal structures in EM stacks*¹, and can be found in the file `EM_ISBI_Challenge.zip`.

The data is from a serial section Transmission Electron Microscope and depicts the ventral nerve of a Drosophila larva. An example image is shown in Figure 17.1. The task is to segment the membranes between cells. These appear mostly darker than the rest of the image, but they are often thin and smeared out, and there are other dark regions that are not membranes, but structures such as mitochondria, that should be part of the cell class. Therefore, this segmentation problem is difficult and requires either a biologist that knows the anatomy or an advanced automated image segmentation method. You should aim at building the automated segmentation method.

The data consists of 30 images with associated labels and additional 30 test images without labels. All images are 512×512 pixels. The task is to build and train a neural network that can segment this type of images. Normally, you would split your data into a training, validation and test set. You would use the training set for learning the model parameters and the validation set to ensure generalization of the model, e.g. that it does not overfit to the training set. In an ideal performance assessment, you will only use the test set once to measure the actual performance of your method. When the same test set is used multiple times, you would optimize for precisely that test set, which will bias your performance assessment.

The problem here is that you only have 30 images with labels to train, validate, and test your algorithm, if it should be done using quantitative measures. The test set does not come with the ground truth labels, since it was kept for evaluating performance of algorithms at the ISBI Challenge, and therefore the images in the test set only allows you to

¹ Ignacio Arganda-Carreras, Srinivas C Turaga, Daniel R Berger, Dan Cireşan, Alessandro Giusti, Luca M Gambardella, Jürgen Schmidhuber, Dmitry Laptev, Sarvesh Dwivedi, Joachim M Buhmann, et al. Crowdsourcing the creation of image segmentation algorithms for connectomics. *Frontiers in neuroanatomy*, 9:142, 2015

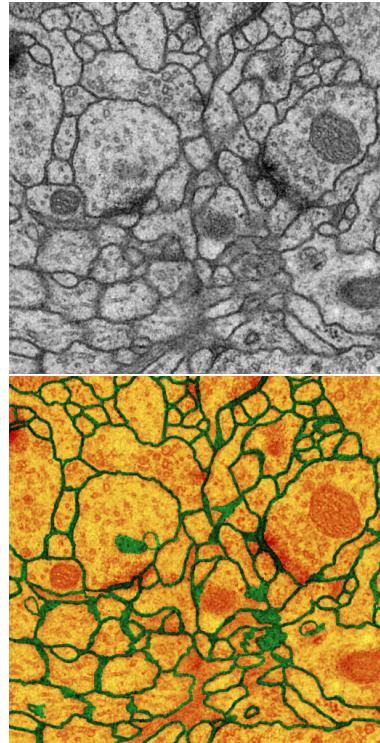


Figure 17.1: Example of EM data for segmentation. Top image is the original EM slice and the bottom image is the labels overlaid.

do a qualitative evaluation. It is your task to choose the images for training, validation, and test, and you should argue for your choice.

You can find inspiration for building your segmentation method in other peoples work. The ISBI challenge data set has e.g. been used in the U-net paper², which is a very successful model for segmentation based on deep learning. This network uses four down-sampling steps followed by four up-sampling steps, and at each resolution there more convolutions and activation steps. In addition there are so-called skip connections that connects layers at different scales. This architecture can be drawn in an u-shape, which has given the name to the method. Furthermore, the paper describes data augmentation suited for this type of data, so you might use the paper to get inspiration for your solution.

² Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. In *International Conference on Medical image computing and computer-assisted intervention*, pages 234–241. Springer, 2015

17.1 Suggestions

You are suggested to implement your model using either the Deep Learning Toolbox from MATLAB or Pytorch or similar for Python. Furthermore, to avoid spending too much time in the initial training, it can be a good idea to start with smaller images than the 512×512 pixels. You can e.g. split the images into smaller patches and work from these. Then you can at a later stage, when you are sure that your model is working as expected, increase the size of your images.

It is also a good idea to start with a simple model, where it is easy to ensure that all steps are working as they should. Then you can gradually increase the complexity.

Besides training the deep learning model and setting it up such that it works as expected, there are many choices to be made for building your model and for optimizing it. You are welcome to try out different approaches and investigate the effect of e.g. data augmentation. It is a very good idea to be systematic and show the effect of different parameter choices – either using quantitative and qualitative evaluation measures.

18 Superresolution from line scans

SCANNING ALONG A SET OF PARALLEL LINES is a common setting in medical imaging. For example, consider optical coherence tomography (OCT), well established in ophthalmology for obtaining images of the retina. Using OCT, the retina is scanned in along a line with a high transversal resolution. Collecting a number of scans along parallel lines, a larger area of the retina may be covered. Since scanning speed of the OCT systems employed in clinic is limited, and prolonged scanning is unpleasant for the patient, the distance between the parallel lines is often large compared to the transverse resolution of the scans. Therefore, the resolution of the scan is much coarser in one direction. In other words, each pixel covers a non-square area. Elongated pixels appear as stripes and influence the visual appearance of the image. The stripes can disturb the interpretation of the image and make it difficult to distinguish the anatomical structures, especially evident with blood vessels running parallel to scan lines.

To reveal additional anatomical structures, another OCT scan may be performed, along the lines orthogonal to the first scan, as a pair of images shown in Figure 18.1. Several problems emerge in connection to this. The eye might move during scanning, and the intensity might vary significantly between the scans. And most importantly, how to combine two scans covering the same area, one with high resolution in x direction, and the other in y , such that the resulting image has a satisfactory quality? To simplify the problem, we consider a set up as in Figure 18.2, where pixels of unknown values are to be estimated from the pixels of known values.

This problem is very similar to single-image superresolution, which can be obtained with great quality using neural networks. In this mini-project, you can try using neural network for upsampling line scans. The project involves setting up a framework based on the publicly available frameworks. We would expect the performance to improve significantly if prior knowledge about the appearance of the images is incorporated in the method. The training should therefore be performed on images having similar appearance as OCT image, but

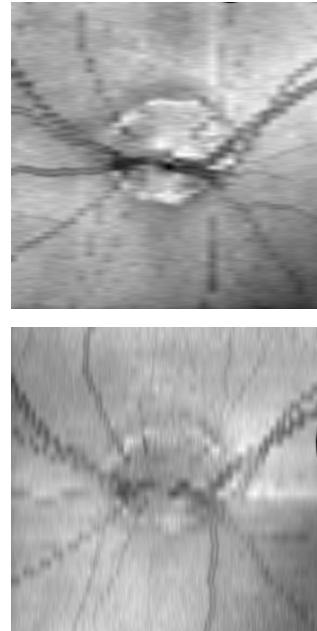


Figure 18.1: Example of images obtained using line scans in orthogonal directions. Notice that vertical lines are less clear in the first image, while horizontal lines are unclear in the second.

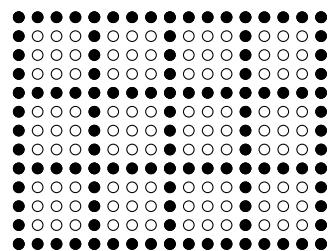


Figure 18.2: An image with pixels divided in pixels with known intensity (black), and pixels with unknown intensity (white), here shown with an upsampling rate of 4.

that can be images of vasculature obtained using other modalities. We will provide you such images.

To evaluate the performance of the upsampling, the conventional approach is to use peak signal-to-noise ratio (PSNR) metric. Furthermore, a comparison with the simple base-line upsampling scheme would be nice. For example, consider a schme where each direction is linearly upsampled, and the two contributions are combined as an average. Such an approach is show in Figure 18.3.

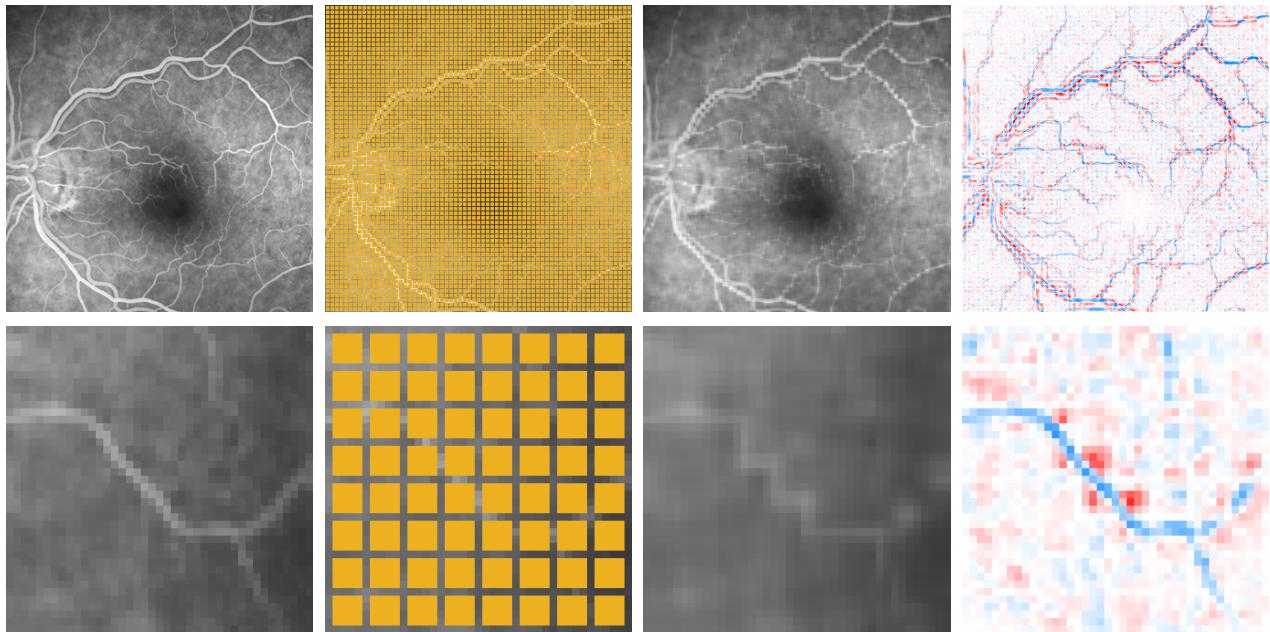


Figure 18.3: A simple linear upsampling scheme. First column: ground truth; second column: unknown pixels masked; third column: upsampling result; fourth column: the error, with color indicating the sign. Top row: image of vasculature; bottom row: zoom in on a small part of the image.